

Course Overview

CS 202: Advanced Operating Systems

Instructor Self-Intro

- Graduated from UNC-Chapel Hill in 2013
- Assistant/associate professor at UT-Dallas during 2013-2022
- Full-time visiting professor at ETH Zurich during 2020.1-2021.6
- Associate professor at UCR since July 2022!
- Autonomous embedded systems, real-time systems, robotics
- <https://intra.ece.ucr.edu/~cong>
- Check csrankings.org

Today

- Course organization and mechanics
- Introduction to OS

What is this course about?

- How has the role of the OS evolved over time?
 - How does the past inform the present?
- What are the underlying principles of OS?
- What are current and future trends in OS?
- Make it real: lab assignments to get some experience with OS development
- *Get you ready to do Systems research*

Some topics we will cover

- OS design models and how they evolved
 - Monolithic kernels, micro-kernels, ...
 - extensibility, protection, performance

Some topics we will cover

- OS design models and how they evolved
 - Monolithic kernels, micro-kernels, ...
 - extensibility, protection, performance
- Concurrency:
 - Scheduling (including provable scheduler design for timing-critical systems)
 - Synchronization
 - Multicore OS

Some topics we will cover

- OS design models and how they evolved
 - Monolithic kernels, micro-kernels, ...
 - extensibility, protection, performance
- Concurrency:
 - Scheduling (including provable scheduler design for timing-critical systems)
 - Synchronization
 - Multicore OS
- File systems:
 - Sequential, networked, distributed, internet scale

Pre-requisites

- CS 202 is a tier-1 course
 - Lectures will recap basics of OS
 - But you are expected to read textbook chapters for background
- To do well, you must understand all lecture materials and read required materials
- Architecture, networks, distributed systems courses are also a plus

Class format

- For every topic:
 - Some overview
 - Read related book chapters
 - Discuss research papers
- Research papers:
 - Critiques for some required papers
 - You are responsible for the papers and materials discussed in class
- Lecture format: 45 min (first half) + 5 min (break) + 30 min (2nd half)

Questions while reading papers

- What are the primary goals (hypothesis)?
 - 2 sentence elevator pitch
- Pros and Cons
- Why did the authors do what they did the way they did?
 - Understand motivation for design
- What were the driving forces for the paper at the time it was written?
- What parts still hold? What parts don't?
- How did they experiment to support their ideas?
- **Write the review by yourself, do NOT leverage online tools such as chatgpt!**

Reading Research Papers

- Guidelines for reading papers
 - Make sure to identify authors' goals and assumptions.
Not always directly stated.
 - Look for high-level takeaways.
 - **Simulate the whole process in your head**
 - Follow a multi-pass reading strategy
 - Pass1: Get overview
 - Pass2: Read details and make notes
 - Pass3: Re-read details to evaluate.
 - Think how techniques used are applicable today.
Identify extensions.

Projects

- 3 Lab Assignments
 - Modifications to xv6 (<https://github.com/mit-pdos/xv6-riscv>)
 - Often not directly connected to the research topics discussed
 - Primary goal is to improve systems building skills
- All labs will be done in **groups** (3 members recommended)
 - Form your group as soon as possible
 - Group with fewer members acceptable, but the scope remains the same
- Submission
 - Short report & code per group on Canvas
 - Brief summaries of individual contributions
 - Short demo
- Late policy: 10% penalty per day; up to 2 late days permitted

Course Materials

- Textbook: OS, 3 easy pieces: [http://
pages.cs.wisc.edu/~remzi/OSTEP/](http://pages.cs.wisc.edu/~remzi/OSTEP/)
 - Its free!
 - Its excellent!
- Lecture Slides/Schedule/ProjectInfo
 - [https://intra.ece.ucr.edu/~cong/teaching/UCR/AOS/
F23.html](https://intra.ece.ucr.edu/~cong/teaching/UCR/AOS/F23.html)
- Published research papers (course website)

Grading Criteria

- Class participation and paper critiques: 20 %
- Lab assignments: 40 %
- Midterm: 20 %
- Final: 20 %

- Midterm:
 - Closed book with 3 cheat sheet (A4 size, both sides), 80 minutes
- Final exam:
 - Closed book with 3 cheat sheet (A4 size, both sides)
 - Final is non-cumulative
 - 3 hours
- A set of CEP questions will be asked in both Midterm and Final exams

Questions and Feedback

- Grades will be available online at elearn/Canvas
 - <https://elearn.ucr.edu/>
 - TA will stop by our next lecture and talk about labs, paper critique submission, grades etc.
- You can provide feedback anytime

Cheating and Plagiarism

- Your code and critiques may be compared against others!
- UCR cheating policies
 - Please read and understand:
<https://conduct.ucr.edu/policies/academic-integrity-policies-and-procedures>
- What consequence can happen?
 - Grade penalty
 - Course failure and department notification
 - Academic sanctions: probation, suspension, dismissal..

Situation

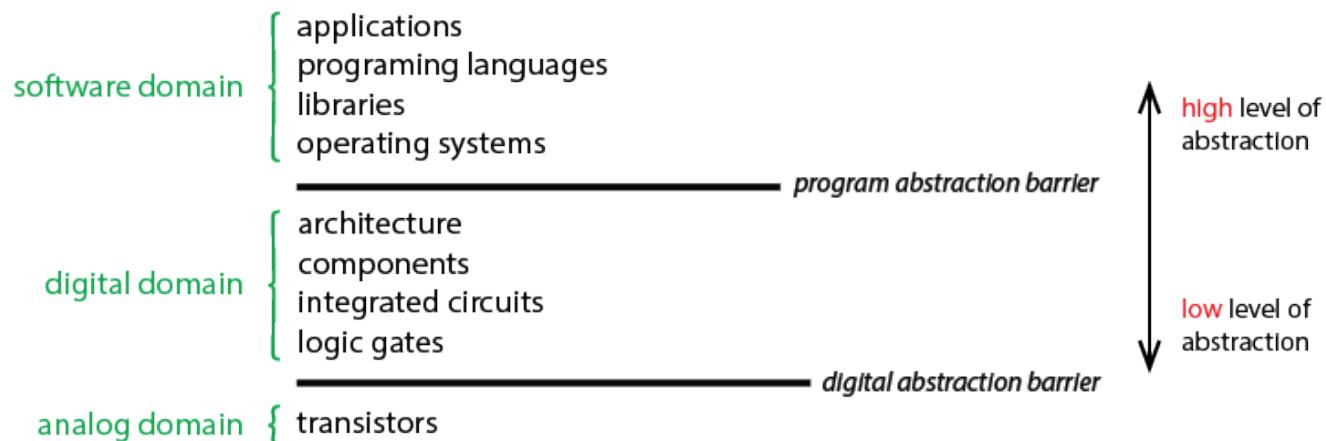
- We all have multiple applications running on our smart phones and computers
 - Written by programmers that don't know each other
 - They all just magically work – how?
- Goal today: get you ready to discuss OS structure, our first topic

Computing Systems – a hierarchy of abstractions

- Computing systems are a **series of abstractions**
 - Impossible to think about a system from electrons to application in one shot
 - What are some abstraction layers we have from transistors to applications?

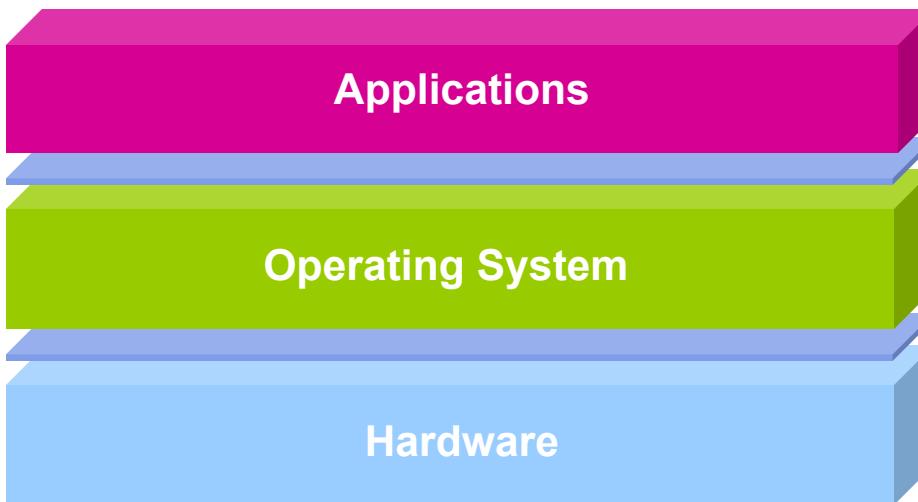
Computing Systems – a hierarchy of abstractions

- Computing systems are a **series of abstractions**
 - Impossible to think about a system from electrons to application in one shot
 - What are some abstraction layers we have from transistors to applications?



- This class: OS level abstractions

What is an OS?



- Has direct access to underlying hardware
- Hides hardware complexity
 - Offers nice abstractions to the applications through system calls
- Manages the hardware on behalf of one or more applications
- Ensures that applications are isolated and protected from each other

Getting more technical

- What is an OS?
 - A piece of software that *abstracts* and *arbitrates* a computing system

Getting more technical

- What is an OS?
 - A piece of software that *abstracts* and *arbitrates* a computing system
- A manager in a shop
 - Directs resources
 - Controls CPUs, memory, devices...
 - Enforces working policies
 - Fairness, resource limits, ...
 - Simplifies complex tasks
 - Abstracts hardware; offers system calls

OS Purposes

- Abstraction of hardware
- Multiplexing
- Isolation
- Sharing
- Security
- Performance
- A range of uses

Abstraction and Arbitration

- OS offers **abstractions** and **arbitration**
- Example of arbitration
 - Allocate memory or CPU time
 - Arbitrate access to a shared device
- Examples of abstractions
 - Files, sockets, process, thread, ...

Abstractions, mechanisms, policies

- Memory management example
- *Abstraction*: memory page
- *Mechanisms*: allocate, deallocate
- *Policies*: page replacement, LRU, LFU, ...

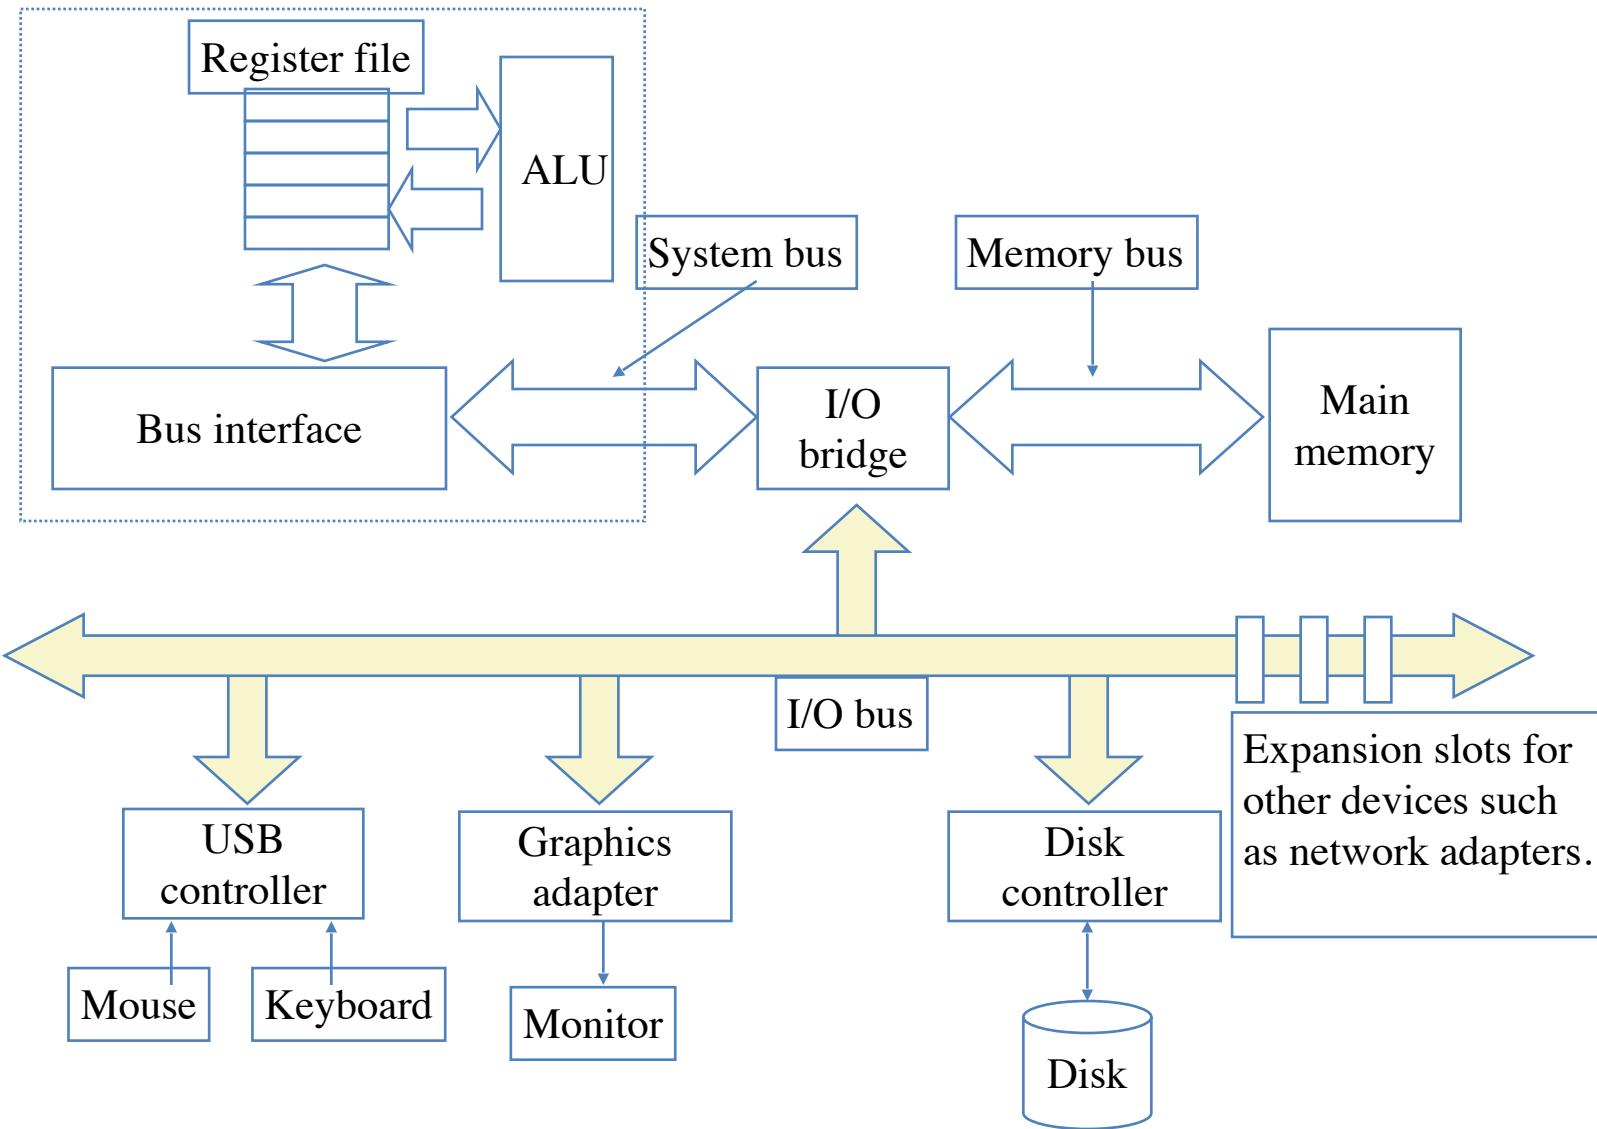
Design principles

- Separation of **mechanism** and **policy**
 - Implement flexible mechanisms to support many policies
 - Develop better policies based on the same mechanism
- Policies must optimize for use cases
 - Where will the OS be used?
 - What will the user want to execute?

Hardware and Resources

- Good understanding of the hardware is essential to understanding OS
- What hardware?
 - Smart phone/tablets?
 - Desktops?
 - Servers?
 - Computing clusters?
 - Cloud?
- How different are these?

They are not that different!



How does the OS interact with the hardware?

- OS
 - Has protected access to hardware resources
 - Arbitrates among competing requests
 - Receives and handles events from the hardware

What support does the hardware provide?

- Manipulating privileged machine state
 - Protected instructions
 - Manipulate device registers, TLB entries, etc.
 - Controlling access
- Generating and handling “events”
 - Interrupts, exceptions, system calls, etc.
 - Respond to external events
 - CPU requires software intervention to handle fault or trap
- Mechanisms to handle concurrency
 - Interrupts, atomic instructions

Catering to Applications

Catering to Applications

- Provide resource needs of an application
 - CPU, memory, device access

Catering to Applications

- Provide resource needs of an application
 - CPU, memory, device access
- When applications launch, the OS loads the program from file into memory
 - Allocates memory for code, data, heap and stack
 - Can the application ask for more resources?
 - Yes, OS receives additional requests and provides resources as needed

Catering to Applications

- Provide resource needs of an application
 - CPU, memory, device access
- When applications launch, the OS loads the program from file into memory
 - Allocates memory for code, data, heap and stack
 - Can the application ask for more resources?
 - Yes, OS receives additional requests and provides resources as needed
- OS also reacts to events in the system

CPU management

- Abstractions
 - Program: static entity
 - Process: program in execution
 - Unit of resource allocation and one thread of execution

Memory management

- Abstractions:
 - Address space for each process
 - Pages to manage memory
- OS implements these abstractions using the available hardware support
 - Paging, segmentation, TLBs, caches...

Storage/file system

- Abstraction: files and directories
 - Others possible: e.g., object store
- Implemented in a variety of different ways
 - Traditional file system: mapping of files to storage
 - Network file system
 - Distributed FS
 - Internet scale FS

Some thoughts on why OS is hard/interesting

- Unforgiving
- Tensions
 - Efficient – Abstraction
 - Powerful – Simple API
 - Flexible - Secure
 - Performance - consistency/determinism/predictability/robustness

4-way intersection demo (coordinating via an RSU)



ROS2-based Autoware.auto



XronOS-based Autoware.auto

4-way intersection demo (coordinating via an RSU)



ROS2-based Autoware.auto



XronOS-based Autoware.auto

4-way intersection demo (coordinating via an RSU)



ROS2-based Autoware.auto



XronOS-based Autoware.auto

4-

Relinquishing predictability in middleware is unacceptable
for safety-critical AS

XronOS: An event-driven middleware that provides
predictable coordination without sacrificing performance



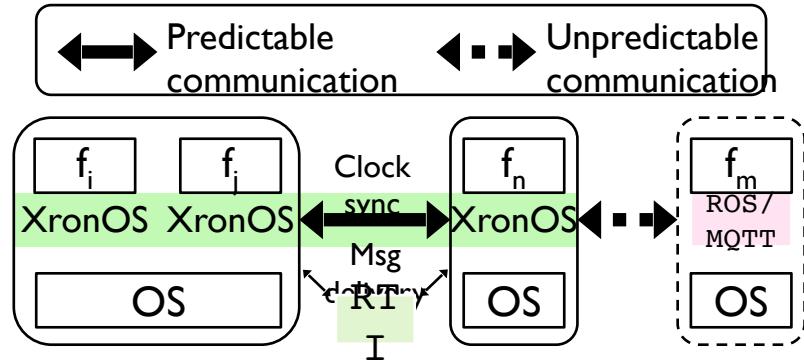
ROS2-based Autoware.auto



XronOS-based Autoware.auto

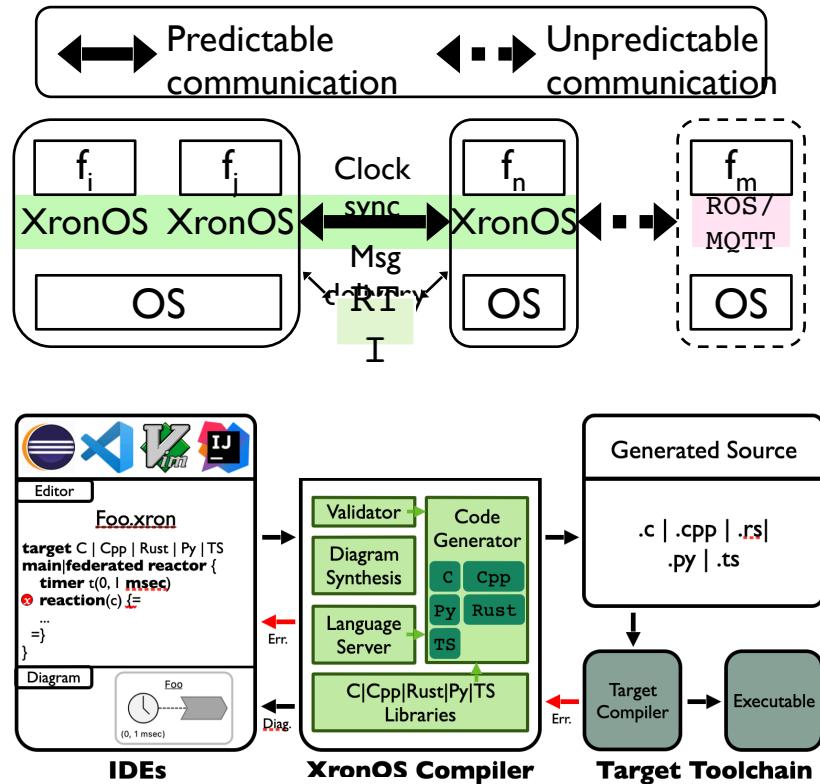
Design goals of XronOS

- ❖ **Enable predictable coordination via explicit temporal semantics based on logical time**
 - Reflect multiplicity of distinct logical & physical timelines
 - ♦ Predictable coordination via the exposure of the relationships between events across timelines
 - Allow the **re-use** of existing ROS and MQTT libraries and nodes
 - Trading off **availability v.s. consistency** in light of Brewer's CAP theorem
 - ♦ Centralized coordination (consistency over availability)
 - ♦ De-centralized coordination (availability over consistency)



Design goals of XronOS

- ❖ **Enable predictable coordination via explicit temporal semantics based on logical time**
 - Reflect multiplicity of distinct logical & physical timelines
 - ♦ Predictable coordination via the exposure of the relationships between events across timelines
 - Allow the **re-use** of existing ROS and MQTT libraries and nodes
 - Trading off **availability v.s. consistency** in light of Brewer's CAP theorem
 - ♦ Centralized coordination (consistency over availability)
 - ♦ De-centralized coordination (availability over consistency)



**Systems research is all about optimizing
tradeoffs and tensions**

Systems research is all about optimizing tradeoffs and tensions

- Think about creating a scalable, fault-tolerant, and high performance distributed file system that could meet the demands of FAAMG's rapidly growing infrastructure

Systems research is all about optimizing tradeoffs and tensions

- Think about creating a scalable, fault-tolerant, and high performance distributed file system that could meet the demands of FAAMG's rapidly growing infrastructure
- **Performance**

Systems research is all about optimizing tradeoffs and tensions

- Think about creating a scalable, fault-tolerant, and high performance distributed file system that could meet the demands of FAAMG's rapidly growing infrastructure
- **Performance** -> Distributing

Systems research is all about optimizing tradeoffs and tensions

- Think about creating a scalable, fault-tolerant, and high performance distributed file system that could meet the demands of FAAMG's rapidly growing infrastructure
- **Performance** → Distributing → Faults

Systems research is all about optimizing tradeoffs and tensions

- Think about creating a scalable, fault-tolerant, and high performance distributed file system that could meet the demands of FAAMG's rapidly growing infrastructure
- **Performance** → **Distributing** → **Faults** → **Fault Tolerance**

Systems research is all about optimizing tradeoffs and tensions

- Think about creating a scalable, fault-tolerant, and high performance distributed file system that could meet the demands of FAAMG's rapidly growing infrastructure
- **Performance** → **Distributing** → **Faults** → **Fault Tolerance**
Replicas <-

Systems research is all about optimizing tradeoffs and tensions

- Think about creating a scalable, fault-tolerant, and high performance distributed file system that could meet the demands of FAAMG's rapidly growing infrastructure
- **Performance** → **Distributing** → **Faults** → **Fault Tolerance**
Inconsistency <- **Replicas** <-

Systems research is all about optimizing tradeoffs and tensions

- Think about creating a scalable, fault-tolerant, and high performance distributed file system that could meet the demands of FAAMG's rapidly growing infrastructure
- **Performance** → **Distributing** → **Faults** → **Fault Tolerance**
Consistency <- **Inconsistency** <- **Replicas** <-

Systems research is all about optimizing tradeoffs and tensions

- Think about creating a scalable, fault-tolerant, and high performance distributed file system that could meet the demands of FAAMG's rapidly growing infrastructure
- **Performance** → Distributing → Faults → Fault Tolerance
Low <- Consistency <- Inconsistency <- Replicas <- performance

Why this class could be challenging

- From the instructor's perspective
 - The “comprehensive” way of teaching AOS yet in a QUARTER system!
 - Tier 1 course: need to consider students who have not taken undergraduate OS classes
- From students' perspective
 - Digesting the comprehensiveness in ~9-10 weeks
 - Some (parts) of the lectures could be boring for those who have taken OS classes before, yet could be rather challenging for students who have not taken any OS classes
 - However, in order to get a good enough grade (e.g., A- and above), you have to get a solid, true understanding of all materials, particularly the research papers covered

Conclusions

- Today was a quick overview of the role of an OS
- Check course schedule and *reading* materials
- Form a group for lab assignments