



# Real-time schedulers & RTOS core services

CS 202: Advanced Operating Systems





# Real-time Systems



# What are Real-Time Systems?

- A real-time system is a system that must satisfy explicit (bounded) response-time constraints or risk severe consequences
- Real-time systems have a dual notion of correctness:
  - **Logical** correctness (“do the right thing”)
  - **Temporal** correctness (“do it on time”)
- A system wherein **predictability** is more important than **performance**
- Example: A robot arm picking up objects from a conveyor belt






# Infamous Real-Time Systems






# What Really Happened on Mars?



- Unmanned spacecraft landed on Mars in 1997
  - Reset reinitializes everything and terminates the execution of the current ground commanded activities
  - Caused by the classical real-time priority inversion problem (mutex-protected shared data area)
  - Tests before launch did not consider “higher data rate than we could have ever imagined” case
  - Later solved by applying a “priority inheritance protocol” first published in RTSS 1990

# Apollo 11

- The first spacecraft landing on the moon
  - Software problem almost caused the landing to be nearly aborted
  - Engineers in charge decided to ignore the problem
  - Control system overloaded => buffer overflow => alarm signals
  - Low-priority jobs were not executed, but not critical at all



# Patriot Missile Control System



- ❖ System used to protect Saudi Arabia during Gulf War
- ❖ Task: detect flying objects, perform prediction, trajectory matches prediction => Patriot missile launched
- ❖ 1991.2.25 --- Scud missile hit city of Dhahran, classified as false alarm (no Patriot missile launched), 28 soldiers killed!
- ❖ Software bug: real-time clock accumulating a delay of 57 microseconds per minute; the battery in question has been in operation for over 100 hours => 343 milliseconds

# Lessons To Be Learned

- If we cannot guarantee something not to go wrong, it will go wrong
- For real-time systems, arguments like “it works now” or “I think it will work” are worthless
- Extensive testing (which can find many bugs) **never guarantee correctness**



# Lessons To Be Learned

- If we cannot guarantee something not to go wrong, it will go wrong
- For real-time systems, arguments like “it works now” or “I think it will work” are worthless
- Extensive testing (which can find many bugs) **never guarantee correctness**

**Correctness should be guaranteed at design time  
by formal analysis & verification with  
clearly stated assumptions and assertions**

# Priority-driven Scheduling

- Priority-driven are dominantly applied in real-time systems
  - Jobs have priorities, and whenever the scheduler selects a job to run, it selects the highest-priority ready job
- Also called work-reserving or list scheduling
  - Never leave the processor idle intentionally
  - Defined by the list of job priorities

# Outline


- We consider both earliest-deadline-first (EDF) and least-laxity-first (LLF) (often called least-slack-time-first) scheduling

# Optimality of EDF

[Liu and Layland, Horn] When preemption is allowed and jobs do not self-suspend nor contend for resources, the EDF algorithm can produce a feasible schedule of a set  $J$  of jobs with arbitrary release times and deadlines on a processor if and only if  $J$  has feasible schedules.

# Proof

- We show that any feasible schedule of  $J$  can be systematically transformed into an EDF schedule
- Suppose parts of two jobs  $J_i$  and  $J_k$  are executed out of EDF order:




- This situation can be corrected by performing a “swap”:





# Proof (Continued)

- If we inductively repeat this procedure, we can eliminate all out-of-order violations
- However, the resulting schedule may still fail to be an EDF schedule because it has idle intervals where some job is ready:




- Such idle intervals can be easily eliminated by moving some jobs forward:



# LLF Scheduling

- **Definition:** At any time  $t$ , the slack (or laxity) of a job with deadline  $d$  is equal to  $(d - t)$  minus the time required to complete the remaining portion of the job



- **LLF Scheduling:** The job with the smallest laxity has highest priority at all times

*↳ more overhead compared to EDF  
⇒ Calculating laxity*

# Optimality of LLF


Theorem 4-3: When preemption is allowed and jobs do not self-suspend nor contend for resources, the LLF algorithm can produce a feasible schedule of a set  $J$  of jobs with arbitrary release times and deadlines on a processor if and only if  $J$  has feasible schedules.

- The proof is similar to that for EDF and is left as an exercise
- Question: Which of EDF and LLF would be preferable in practice?

# Preemptive v.s. Nonpreemptive EDF


Theorem: Nonpreemptive EDF is not optimal.

- **Proof by a counterexample**: Consider a system of three jobs  $J_1$ ,  $J_2$ , and  $J_3$  such that  $(r_1, e_1, d_1) = (0, 3, 10)$ ,  $(r_2, e_2, d_2) = (2, 4, 8)$ ,  $(r_3, e_3, d_3) = (4, 6, 14)$ . Here is a feasible schedule:



# Proof (Continued)

- Here shows the nonpreemptive EDF schedule, we see a deadline miss!



- Question: Should we conclude from this result that preemptive EDF is always better than non-preemptive EDF in practice?
  - Note: The EDF optimality proof assumes there is no penalty due to preemption
  - Are there other practical issues we have ignored?

*↳ High context switch overhead*



# Real-time operating systems

# Core Services

- ❖ Most RTOSs offer APIs for at least:
  - ❖ **Timers** and **clocks** (e.g., to allow threads to “sleep”)
  - ❖ **Scheduling primitives** (priority, preemption control)
  - ❖ **Synchronization** (e.g., semaphores)
  - ❖ **Interrupt control** (e.g., to register ISRs, etc.)
  - ❖ Access to **I/O devices** (e.g., map I/O ports into a thread’s address space)
  - ❖ **Reserving memory** (e.g., to prevent memory from being paged out by the OS) *↳ Predictability*

# Time Services

- ❖ Every system has at least one hardware device that contains a **counter**
  - ❖ Counts down from some initial value and generates an interrupt when 0 is reached
- ❖ A counter can (usually) be programmed to act as either:
  - ❖ **Periodic Timer:** Generates an interrupt every  $p$  time units
  - ❖ **One-Shot Timer:** Must be re-initialized after each interrupt
- ❖ The counter can serve as a **clock** and/or a **timer**
  - ❖ **Clock:** A counter and ISR that together keep (system) time
  - ❖ **Timer:** A counter and ISR that triggers an event  $e$  at time  $t$  (relative or absolute)

# Clock Resolution

- ❖ The **resolution** of a clock is the granularity of time provided by the clock
  - ❖ The hardware resolution may be on the order of nanoseconds
  - ❖ However, the resolution seen by the application is usually in the order of a few microseconds to milliseconds
- ❖ Clock resolution is fundamentally limited by:
  - ❖ **CPU frequency** (interrupts are only processed between instructions)
  - ❖ **Max. instruction length** (on modern Intel systems, some instructions can execute for hundreds of cycles)
  - ❖ **Interrupt latency** (time-service interrupts may be delayed)  
*↳ Split interrupt handling*

# Software Clocks

- ❖ “Current time” has many definitions
- ❖ Most RTOSs keep track of at least the following clocks:
  - ❖ **Absolute (wall-clock) time** (including daylight saving time, leap seconds, etc.)
  - ❖ **Monotonic time** (without discontinuities, e.g., microseconds since system startup)
  - ❖ **Per-thread execution time** (CPU allocation)



# Software Timers

- ❖ Most OSs (including all real-time POSIX-compliant systems) allow a thread to have its own software timers
- ❖ Threads can **create**, **set**, **cancel**, and **destroy** timers
- ❖ When a timer is created by a thread, the kernel creates a timer data structure, which includes:
  - ❖ The clock associated with the timer
  - ❖ The thread that created the timer
  - ❖ The expiration of the timer in **absolute** or **relative** time, once set
  - ❖ The timer type: **one-shot** or **periodic**
  - ❖ An event handler: a routine to execute when the timer expires

# Software Timers

- ❖ **Synchronous Timers** suspend the calling thread until the timer has expired
  - ❖ For example, `timer_sleep(relative_time)` or `timer_sleep_until(absolute_time)`
- ❖ **Asynchronous Timers** count down while the calling thread executes
  - ❖ When the timer expires, the associated handler is called
  - ❖ Threads may have multiple asynchronous timers active
  - ❖ Asynchronous timers can be used as a timing monitor to log missed deadlines (see Figure 12-4 in Liu)

# Scheduler Activation

- ❖ **Event-driven** *→ Sleeping heavily design*
  - ❖ The scheduler is invoked on job release, job completion, and when a job blocks/resumes
  - ❖ Job releases are triggered by interrupts / timers
- ❖ **Quantum-driven**
  - ❖ The scheduler is invoked by a periodic timer
  - ❖ Periodic job releases may be triggered by the scheduler
- ❖ **Hybrid approaches are common**
  - ❖ In current Linux versions, the scheduler is invoked in response to events, but also by a periodic scheduler tick (usually once every millisecond)

# Scheduling Mechanisms

- ❖ Virtually all RTOSs support **Priority-driven Scheduling**
    - ❖ Many RTOSs provide **256** priority levels
      - We have already seen that 256 levels approximates an ideal system
    - ❖ General-purpose OSs usually provide **10-20** levels
      - NT has only 16 real-time priority levels
      - Exception: Linux, which provides 100 levels.
    - ❖ The priority is usually set (once) when the thread is created and stored in the TCB
    - ❖ The kernel (usually) maintains a ready queue for each priority level
    - ❖ Dispatching the highest priority ready thread requires finding the highest priority nonempty queue
      - How do we do this?
      - **What is the complexity of this?**
- 20 levels not enough for real time systems  
→ More granular control*

# Dispatching Threads

- ❖ An efficient way of doing this in software:
  - ❖ Data structures:
    - Assume a  $K$ -bit **queue-status word** and  $\Omega$  priority queues
    - Let each bit of the word represent the status of  $\Omega/K$  priority queues
    - If a bit is set, one of the  $\Omega/K$  associated queues has a ready thread
  - ❖ Finding a ready thread:
    - If the **queue-status word**  $> 0$  then a thread is ready
      - Perform a binary search on the word (comparing its value to powers of 2) to find the highest-order bit set
      - Find the highest priority non-empty queue of the  $\Omega/K$  queues associated with the highest-order bit set, and dequeue the job at the head of the queue
      - Clear the associated bit in the **queue-status word** if all  $\Omega/K$  queues are empty
  - ❖ Complexity: at most  $\Omega/K + \log_2 K$ 
    - At most 13 comparisons for a 32-bit **queue-status word** and 256 priorities



# Thread Dispatching in Practice

- ❖ Many architectures have instructions for finding the highest non-zero bit in hardware
  - ❖ On x86, the **Bit Scan Forward (BSF)** instruction yields the index of the highest non-zero bit in a 32-bit word
- ❖ Using one bit per queue, the scheduler can find the highest-priority non-empty queue with at most eight BSF instructions (assuming 256 priority levels)
  - ❖ Only four instructions on 64-bit systems

$(256/64 = 4)$

# Preemption Locks

- ❖ Almost all RTOSs allow **a thread to disable preemption by disabling interrupts**
  - ❖ Preemption locks are usually used to protect critical sections in both the kernel and the application threads
- ❖ However, some also support **disabling preemption by setting a mode control flag to disable the scheduler**
  - ❖ With this feature, hardware interrupts (such as the clock) are not disabled when context switching is disabled

# Resource Access Control

- ❖ The most common resource access control method used in real-time systems is a user-level implementation of the **nonpreemptable critical section (NPCS) protocol**:
  - ❖ Each thread disables preemption (using some form a preemption lock) just before accessing a resource
  - ❖ The thread enables preemption (releases the preemption lock) as soon as it is done with the resource
- ❖ How does this guarantee that the resource is always available when the thread accesses it?

# RTOS Design Goals

- ❖ Most RTOSs are used in embedded systems with limited resources (CPU, memory)
  - ⇒ *Microkernel*
- ❖ **Design time:** the OS should be modular and extensible
- ❖ **Run time:** often a fixed set of tasks, only little changes in workload
- ❖ OS services designed to limit worst-case latencies
- ❖ In contrast, best-effort OS (e.g., Windows, Linux)
  - ❖ “One size fits all”
  - ❖ Must handle wide variety of workloads
  - ❖ Optimized for average-case performance

# Capabilities of Commercial RTOSs

- ❖ Commercially available RTOSs:
  - ❖ LynxOS, QNX, pSOSystem, Nucleus RTOS, VRTX, and VxWorks
- ❖ Each of these systems shares the following attributes:
  - ❖ Compliant or partially compliant to the Real-Time POSIX API Standard:
    - ▶ Preemptive, fixed-priority scheduling
    - ▶ Standard synchronization primitives (mutex and message passing)
    - ▶ Each also has its own API
  - ❖ Modular and scalable:
    - ▶ The kernel is small so that it can fit in ROM in embedded systems
    - ▶ I/O, file, and networking modules can be added

# Shared Attributes (Continued)

## ❖ Fast and efficient:

- ❖ Most are microkernel systems
- ❖ Low overhead
- ❖ Small context switch time, interrupt latency, and semaphore get/release latency: **usually one to a few microseconds**
- ❖ Nonpreemptable portions of kernel functions are highly optimized, short, and as deterministic as possible
- ❖ Many have system calls that require no trap: **applications run in kernel mode**

## ❖ Flexible scheduling:

- ❖ All offer at least 32 priority levels: min required by real-time POSIX
- ❖ Most offer 128 or 256 priority levels
- ❖ FIFO or RR scheduling for equal-priority threads

# Shared Attributes (Continued)

- split interrupt in two parts*  
*⇒ Small part ⇒ "acknowledging";*  
*executed first*

  - ❖ Support split interrupt handling
- ⇒ complex part ⇒ Scheduled for later*

  - ❖ Relatively High Clock and Timer Resolution:
    - ❖ Nominal timer resolution in the nanosecond range
    - ❖ Actual timer resolution in the microsecond range
- maintaining predictability*

  - ❖ No Paging or Swapping:
    - ❖ May provide logical-to-physical memory translation but no paging
    - ❖ May not offer memory protection: **often the kernel and all tasks execute in kernel mode, sharing one common address space**
    - ❖ Level of memory protection may be settable (ranging from "none" to "private virtual memory")
- ❖ Optional Networking Support:
    - ❖ Can be configured to support TCP/IP with an optional module

*"Partial Paging"*  
 ↓  
*"Cache colouring"*  
 ↓  
*Partitioning cache lines among processes*