

# Course Summary and Final Review

CS 202: Advanced Operating Systems

# What was this course about?

- How has the role of the OS evolved over time?
- What are the underlying principles of OS?
- What are current and future trends in OS?
- Make it real: lab assignments to get some experience with OS development
- Know how to critically read academic papers
- *Ready to do Systems research?*

# Topics we have covered

Week	Date	Day	Class activity
0	09/22	Thu	Course Overview
1	09/27	Tue	OS Organization XV6 Overview
	09/29	Thu	Processes and Threads
2	10/04	Tue	Virtual Memory
	10/06	Thu	Extensible OS Design I
3	10/11	Tue	Extensible OS Design II
	10/13	Thu	Scheduling
4	10/18	Tue	Scheduler Design
	10/20	Thu	Concurrency and Synchronization
5	10/25	Tue	Synchronization (cont'd)
	10/27	Thu	<b>Midterm</b>
6	11/01	Tue	Multicore OS
	11/03	Thu	Scalability to Many Cores
7	11/08	Tue	I/O, Storage, and File Systems
	11/10	Thu	File Systems II
8	11/15	Tue	Cloud Storage
	11/17	Thu	Virtualization
9	11/22	Tue	Virtualization (cont'd)
	11/24	Thu	<b>Campus holiday; No class</b>
10	11/29	Tue	Embedded and IoT
	12/01	Thu	Review and Wrap-up

# Papers we have reviewed

Readings and Assignment
<a href="#">Ch02</a>
<a href="#">Ch04</a> , <a href="#">Ch05</a> , <a href="#">Ch06</a> , <a href="#">Ch26</a>
<a href="#">Ch13</a> , <a href="#">Ch18</a> , <a href="#">Ch19</a> , <b>Lab 1</b>
Reading: <a href="#">Spin</a> (critique)
Reading: <a href="#">Exokernel</a> (critique), Optional: <a href="#">L4</a> , <a href="#">seL4</a>
<a href="#">Ch07</a> , <a href="#">Ch08</a> , <a href="#">Ch10</a>
Reading: <a href="#">Lottery scheduling</a> (critique), <a href="#">Stride scheduling</a> , Optional: <a href="#">Scheduler activations</a> , <b>Lab 2</b>
<a href="#">Ch28</a> , <a href="#">Ch29</a> , <a href="#">Ch31</a> , <a href="#">Ch32</a>
Reading: <a href="#">RCU</a>
Reading: <a href="#">Multikernel</a> (critique), Optional: <a href="#">LegoOS</a>
Reading: <a href="#">Linux scalability</a> (critique), Optional: <a href="#">Decade of wasted cores</a>
<a href="#">Ch37</a> , <a href="#">Ch39</a> , <a href="#">Ch40</a> , <a href="#">Ch41</a> , <b>Lab 3</b>
Reading: <a href="#">LFS</a> , <a href="#">FFS</a> , <a href="#">NFS</a> , Optional: <a href="#">AFS</a>
Reading: <a href="#">GFS</a> (critique), Optional: <a href="#">HDFS</a>
Reading: <a href="#">Xen</a> (critique), <a href="#">Appendix B</a>
Optional: <a href="#">SW/HW techniques for x86</a> , <a href="#">NOVA</a>
Reading: <a href="#">Tock</a> (critique), Optional: <a href="#">TinyOS</a> , <a href="#">Contiki</a> , <a href="#">Ink</a>

# Final Exam

- 12/06/2022, Tuesday, 09:00am-10:30pm, MSE 104
- Closed book with cheat sheets
  - Up to 3 sheets of letter or A4 paper (same as midterm)  
Proctors will check number of sheets properly.
- Question types:
  - True and false, multiple choices (single or multi select), short answers  
Multi-select: Specify number of correct answers.
  - Background on OS concepts & papers discussed

# Scope

- Lecture slides from 10 to 19 (ignore “11\_Review for Midterm.pdf”)

	10/20	Thu	Concurrency and Synchronization	<a href="#">Ch28</a> , <a href="#">Ch29</a> , <a href="#">Ch31</a> , <a href="#">Ch32</a>
5	10/25	Tue	Synchronization (cont'd)	Reading: <a href="#">RCU</a>
	10/27	Thu	<b>Midterm</b>	
6	11/01	Tue	Multicore OS	Reading: <a href="#">Multikernel</a> (critique), Optional: <a href="#">LegoOS</a>
	11/03	Thu	Scalability to Many Cores	Reading: <a href="#">Linux scalability</a> (critique), Optional: <a href="#">Decade of wasted cores</a>
7	11/08	Tue	I/O, Storage, and File Systems	<a href="#">Ch37</a> , <a href="#">Ch39</a> , <a href="#">Ch40</a> , <a href="#">Ch41</a> , <b>Lab 3</b>
	11/10	Thu	File Systems II	Reading: <a href="#">LFS</a> , <a href="#">FFS</a> , <a href="#">NFS</a> , Optional: <a href="#">AFS</a>
8	11/15	Tue	Cloud Storage	Reading: <a href="#">GFS</a> (critique), Optional: <a href="#">HDFS</a>
	11/17	Thu	Virtualization	Reading: <a href="#">Xen</a> (critique), <a href="#">Appendix B</a> Optional: <a href="#">SW/HW techniques for x86</a> , <a href="#">NOVA</a>
9	11/22	Tue	Virtualization (cont'd)	
	11/24	Thu	<b>Campus holiday; No class</b>	
10	11/29	Tue	Embedded and IoT	Reading: <a href="#">Tock</a> (critique), Optional: <a href="#">TinyOS</a> , <a href="#">Contiki</a> , <a href="#">Ink</a>
	12/01	Thu	Review and Wrap-up	

# Synchronization

- Atomicity
- Mutual exclusion
- Spinlock
- Disabling interrupts
- Test-And-Set (TAS) lock
- Semaphore
- Deadlock conditions
- Lock ranking
- RCU

# True and False

- Read-Copy Update (RCU) requires garbage collection for old versions of data True
- Hardware atomic instructions are expensive since they may stall other CPUs in a multi-core system True
- Deadlocks do not occur with spinlocks False



# Multicore OS

- Scalability
- Amdahl's law
- Cache coherence
- Problem of Test-And-Set (TAS) Lock
- Multikernel
  - Message passing
  - Replicas
  - CPU drivers and monitors
  - Inter-core communication
  - Memory management
  - System knowledge base

# Scalability

- Eliminate bottlenecks in the Linux kernel
  - Multicore packet processing
    - Per-core RX/TX queues
  - Sloppy counters
    - Local reserve for reference counters
  - Lock-free comparison
  - Per-core data structures
  - Eliminating false sharing
    - Use separate cache lines

# Short Answer

- If 80% of a program is parallelizable, what is the maximum speedup achievable with an infinite number of processors? (assume ideal condition)

$1/(1-p) + p/s$  using Amdahl's law

If  $s \rightarrow \text{infinity}$ , speedup approaches  $1/(1 - p)$ .

$$1/(1-0.8) = 5$$

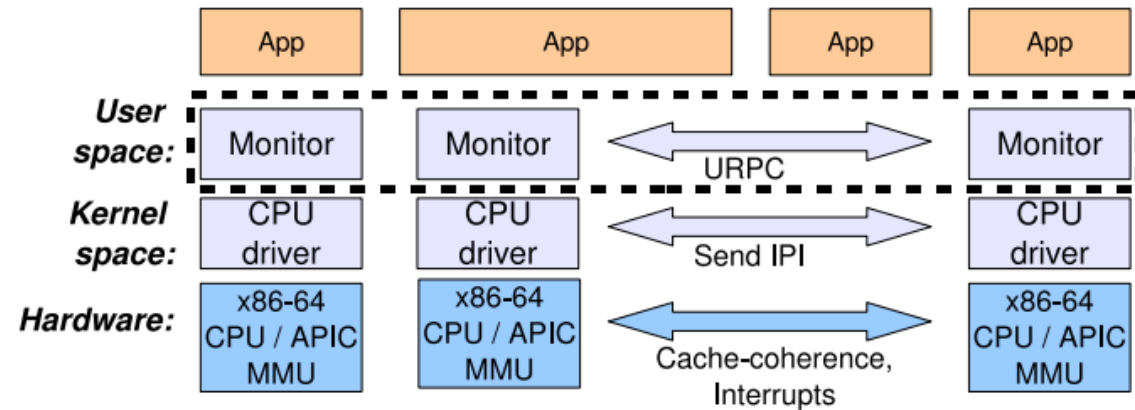
# Multiple choices

- Typically, it is hard to achieve  $N$  times of speedup by adding  $N$  processors. Which of the following is good for better scalability?
  - A. Multiple tasks updates the shared data protected by RCU simultaneously
  - B. Tasks compete for memory bus E - is good, F
  - C. Too few tasks are running
  - D. Shared data updated often are located on the same cache line as those read often
  - E. The kernel uses per-core data structure
  - F. Protect all shared kernel states using a single mutex lock

Ans. E - is good

D - false sharing is bad  
F - bad, better to have fine-grained locks.

# Short Answer



- “In multikernel, the CPU drivers are completely  
 [ event-driven, ], [ nonpreemptable ], and [ single-threaded ].”

Choose the correct words to fill in the blanks.

Word pool:

event-driven, time-triggered, priority-driven,  
 preemptible, nonpreemptable,  
 single-threaded, multi-threaded, multi-processed  
 hardware-neutral, portable, extensible

# Short Answer

- “In multikernel, the CPU drivers are completely [                    ], [                    ], and [                    ].”  
Choose the correct words to fill in the blanks.

Word pool:

event-driven, time-triggered, priority-driven,  
preemptible, nonpreemptible,  
single-threaded, multi-threaded, multi-processed  
hardware-neutral, portable, extensible

# I/O and Storage Devices

- I/O layers
  - Polling vs. Interrupts
- Disks and SSD
  - Disk I/O time = Seek Time + Rotation Time + Transfer Time
  - SSD Advantages
  - Write Amplification
  - Wear Leveling
- Virtual File System
- Inode
- UNIX File System (UFS)

# FFS and LFS

- Fast File System (FFS)
  - Cylinder groups
  - Larger block sizes
  - Sub-blocks
- Log-structure File System (LFS)
  - Sequential writes
  - inode map
  - Checkpoint Region
- Pros and cons



# True and False

- Virtual file system (VFS) provides a uniform interface to user programs regardless of the actual file system being used True
- Larger block size of Fast file system (FFS) help improve “latency” False, only helps throughput
- In Log-structure File System (LFS), there may exist multiple versions of the same file

True

# Multiple choices

- Which one is correct about polling vs. interrupts?
  - A. Interrupts consume more CPU time than polling for long I/O operations
  - B. Polling improves CPU utilization if the device needs service from the CPU occasionally
  - C. Interrupts allow asynchronous I/O in programs
  - D. The number of instructions for handling an event after polling is higher than that for handling the same event after receiving an interrupt

Ans. C

# Multiple choices

- Which of the following is **NOT** the feature of the Fast File System (FFS)?
  - A. Cylinder groups to improve average access time
  - B. Larger block size to improve bandwidth
  - C. Larger block size to support larger files
  - D. Replicated superblocks in cylinder groups
  - E. Pre-allocate blocks to improve write performance

E is wrong

# Network File System (NFS)

- Stateful vs. stateless protocols
- NFS
  - Stateless
  - Single centralized server
  - File lookup
  - Caching and consistency
  - Idempotent requests
  - File locking
  - Time synchronization

# Google File System (GFS)

- Single master server
- Chunk servers
- Replicas
- User-space API
- Namespace
- Lease & Mutation order
- Record append

# True and False

- In the design of network-based file systems, stateful protocols make crash/disconnect recovery easier

False

- Google File System (GFS) is designed for expensive high-endurance disk drives that rarely fails

False

# Multiple choices

- To open a file `/foo/bar/cs202/test.c` stored in the server, how many lookup messages will be made by the client?
  - A. 0
  - B. 1
  - C. 2
  - D. 3
  - E. 4

E. 4 (foo + bar + cs202 + test.c)

# Virtualization

- Trap-and-emulate
- x86 virtualization challenges
- Dynamic binary translation
- Shadow paging
- Para-virtualization
  - Xen
  - Dom0
  - CPU/memory/IO virtualization
- Hardware extension



# Multiple choices

- The trap-and-emulate approach faces multiple challenges in virtualizing classical x86 architectures (without hardware extensions). Which one is correct?
  - A. The guest OS does not know that it's running in an unprivileged mode
  - B. A privileged instruction in the guest OS does not trigger a trap
  - C. x86 does not provide a mechanism to set write-protected pages for shadow paging
  - D. x86's multi-level paging prevents the use of shadow page tables

B

# Short Answer

- Can shadow paging be faster in address translation than nested/extended paging? Give your answer with a short explanation

Yes

# Embedded OS

- TinyOS
  - Concurrency: Event-driven architecture; no preemption
  - Modularity: scheduler + graph of components
  - Compiled into one executable
  - Static memory allocation
- Contiki
  - Preemptive multithreading on top of event-driven kernel
- Tock
  - Rust - Isolate drivers
  - MPU - Isolate applications
  - Grants - enable dynamic memory without crashing the kernel

# True and False

- TinyOS implements priority-based scheduling for tasks

False

- Tock supports dynamic memory allocation for processes

True

- Tock supports virtual memory

False