

# Course Summary and Final Review

CS 202: Advanced Operating Systems

# What was this course about?

- How has the role of the OS evolved over time?
- What are the underlying principles of OS?
- What are current and future trends in OS?
- Make it real: lab assignments to get some experience with OS development
- *Ready to do Systems research?*

# Topics we have covered

Week	Date	Day	Class activity
0	09/23	Thu	Course Overview
1	09/28	Tue	OS Organization XV6 Overview
	09/30	Thu	Processes and Threads
2	10/05	Tue	Virtual Memory
	10/07	Thu	Extensible OS Design I
3	10/12	Tue	Extensible OS Design II
	10/14	Thu	Scheduling
4	10/19	Tue	Scheduler Design
	10/21	Thu	Synchronization
5	10/26	Tue	Concurrency and Memory Consistency
	10/28	Thu	<b>Midterm</b>
6	11/02	Tue	Multicore OS
	11/04	Thu	Scalability to Many Cores
7	11/09	Tue	I/O and File Systems
	11/11	Thu	<b>Campus holiday; No class</b>
8	11/16	Tue	File Systems II
	11/18	Thu	Distributed File Systems
9	11/23	Tue	Cloud Storage
	11/25	Thu	<b>Campus holiday; No class</b>
10	11/30	Tue	Virtualization
	12/02	Thu	OS for Embedded and IoT

# Papers we have reviewed

Readings and Assignment
<a href="#">Ch02</a>
<a href="#">Ch04</a> , <a href="#">Ch05</a> , <a href="#">Ch06</a> , <a href="#">Ch26</a>
<a href="#">Ch13</a> , <a href="#">Ch18</a> , <a href="#">Ch19</a> , Lab 1
Reading: <a href="#">Spin</a> (critique)
Reading: <a href="#">Exokernel</a> (critique) Optional: <a href="#">L4</a>
<a href="#">Ch07</a> , <a href="#">Ch08</a> , <a href="#">Ch10</a>
Reading: <a href="#">Lottery Scheduling</a> (critique) <a href="#">Stride Scheduling</a> , <a href="#">Scheduler Activations</a> , Lab 2
<a href="#">Ch28</a> , <a href="#">Ch29</a> , <a href="#">Ch31</a> , <a href="#">Ch32</a>
Reading: <a href="#">RCU</a>
Reading: <a href="#">Barrelfish</a> (critique), Optional: <a href="#">LegoOS</a>
Reading: <a href="#">Linux scalability</a> (critique), Optional: <a href="#">Decade of wasted cores</a>
<a href="#">Ch37</a> , <a href="#">Ch39</a> , <a href="#">Ch40</a> , <a href="#">Ch41</a> , Lab 3
Reading: <a href="#">LFS</a> , <a href="#">FFS</a>
Reading: <a href="#">NFS</a> , Optional: <a href="#">AFS</a>
Reading: <a href="#">GFS</a> (critique), Optional: <a href="#">HDFS</a>
Reading: <a href="#">Xen</a> (critique), <a href="#">Appendix B</a>
Optional: <a href="#">SW/HW techniques for x86</a>
Reading: <a href="#">Tock</a> (critique) <a href="#">TinyOS</a> , Optional: <a href="#">Contiki</a> , <a href="#">InK</a>

# Final Exam

- 12/04/2021, Saturday, 9 am to 11 am, at Sproul 1102
- Closed book with cheat sheets
  - Up to ~~3~~ 5 sheets of letter or A4 paper
- Question types:
  - True and false, single choice, ~~multiple choices~~, short answers, ~~long answers~~
  - Background on OS concepts & papers discussed
  - No “None of the above” type of choices

# Scope

- Lecture slides from 11 to 20 (ignore “12\_Review for Midterm.pdf”)

10/21	Thu	Synchronization	<a href="#">Ch28</a> , <a href="#">Ch29</a> , <a href="#">Ch31</a> , <a href="#">Ch32</a>
10/26	Tue	Concurrency and Memory Consistency	Reading: <a href="#">RCU</a>
10/28	Thu	<b>Midterm</b>	
11/02	Tue	Multicore OS	Reading: <a href="#">Barrelfish</a> (critique), Optional: <a href="#">LegoOS</a>
11/04	Thu	Scalability to Many Cores	Reading: <a href="#">Linux scalability</a> (critique), Optional: <a href="#">Decade of wasted cores</a>
11/09	Tue	I/O and File Systems	<a href="#">Ch37</a> , <a href="#">Ch39</a> , <a href="#">Ch40</a> , <a href="#">Ch41</a> , Lab 3
11/11	Thu	<b>Campus holiday; No class</b>	
11/16	Tue	File Systems II	Reading: <a href="#">LFS</a> , <a href="#">FFS</a>
11/18	Thu	Distributed File Systems	Reading: <a href="#">NFS</a> , Optional: <a href="#">AFS</a>
11/23	Tue	Cloud Storage	Reading: <a href="#">GFS</a> (critique), Optional: <a href="#">HDFS</a>
11/25	Thu	<b>Campus holiday; No class</b>	
11/30	Tue	Virtualization	Reading: <a href="#">Xen</a> (critique), <a href="#">Appendix B</a> Optional: <a href="#">SW/HW techniques for x86</a>
12/02	Thu	OS for Embedded and IoT	Reading: <a href="#">Tock</a> (critique), <a href="#">TinyOS</a> , Optional: <a href="#">Contiki</a> , <a href="#">InK</a>

# Synchronization

- Atomicity
- Mutual exclusion
- Spinlock
- Disabling interrupts
- Test-And-Set (TAS) lock
- Semaphore
- Deadlock conditions
- Lock ranking
- RCU

# True and False

- Read-Copy Update (RCU) requires garbage collection for old versions of data

TRUE

- Hardware atomic instructions are expensive since they may stall other CPUs in a multi-core system

TRUE

- Deadlocks do not occur with spinlocks

FALSE



# Multicore OS

- Scalability
- Amdahl's law
- Cache coherence
- Problem of Test-And-Set (TAS) Lock
- Multikernel
  - Message passing
  - Replicas
  - CPU drivers and monitors
  - Inter-core communication
  - Memory management
  - System knowledge base

# Scalability

- Eliminate bottlenecks in the Linux kernel
  - Multicore packet processing
    - Per-core RX/TX queues
  - Sloppy counters
    - Local reserve for reference counters
  - Lock-free comparison
    - Generation counter
  - Per-core data structures
  - Eliminating false sharing
    - Use separate cache lines

# Short Answer

- If 80% of a program is parallelizable, what is the maximum speedup achievable with an infinite number of processors?

$$SPEEDUP = \frac{1}{(1 - p) + \frac{p}{s}}$$

If  $s \rightarrow \infty$ , speedup approaches  $1/(1-p)$

$$1/(1-0.8) = 5$$

# Single Choice

- Typically, it is hard to achieve “X” times of speedup by adding “X” processors. Which of the following is **NOT** the correct reasons for imperfect scalability?
  - A. Multiple tasks updates the shared data protected by RCU simultaneously
  - B. Tasks compete for memory bus
  - C. Too few tasks are running
  - D. Shared data updated often are located on the same cache line as those read often
  - E. The kernel uses per-core data structure
  - F. Shared kernel states are protected by mutex locks

E: good for scalability

# Short Answer

- In multikernel, the CPU drivers are completely event-driven, [                      ], and [                      ]. Choose the correct words to fill in the blanks.

Word pool: time-triggered, message-based, preemptible, nonpreemptible, consistent, global, single-threaded, multi-threaded, hardware-neutral,

single-threaded, nonpreemptible

# I/O and Storage Devices

- I/O layers
  - Polling vs. Interrupts
- Disks and SSD
  - Disk I/O time = Seek Time + Rotation Time + Transfer Time
  - SSD Advantages
  - Write Amplification
  - Wear Leveling
- Virtual File System
- Inode
- UNIX File System (UFS)

# FFS and LFS

- FFS
  - Cylinder groups
  - Larger block sizes
  - Sub-blocks
- LFS
  - Sequential writes
  - inode map
  - Checkpoint Region
- Pros and cons

# True and False

- Virtual file system (VFS) provides a uniform interface to user programs regardless of the actual file system being used  
**TRUE**
- Larger block size of Fast file system (FFS) help improve latency  
**FALSE**
- In Log-structure File System (LFS), there may exist multiple versions of the same file  
**TRUE**



# Single Choice

- Which one is correct about polling vs. interrupts?
  - A. Interrupts consume more CPU time than polling for long I/O operations
  - B. Polling improves CPU utilization if the device needs service from the CPU occasionally
  - C. Interrupts allow asynchronous I/O in programs
  - D. The number of instructions for handling an event after polling is higher than that for handling the same event after receiving an interrupt

C

# Single Choice

- Which of the following is **NOT** the feature of the Fast File System (FFS)?
  - A. Cylinder groups to improve average access time
  - B. Larger block size to improve bandwidth
  - C. Larger block size to support larger files
  - D. Replicated superblocks in cylinder groups
  - E. Pre-allocate blocks to improve write performance

E

# Network File System (NFS)

- Stateful vs. stateless protocols
- NFS
  - Stateless
  - Single centralized server
  - File lookup
  - Caching and consistency
  - Idempotent requests
  - File locking
  - Time synchronization

# Google File System (GFS)

- Single master server
- Chunk servers
- Replicas
- User-space API
- Namespace
- Lease & Mutation order
- Record append

# True and False

- In the design of network file systems, stateful protocols make crash/disconnect recovery easier

FALSE

- Google File System (GFS) is designed for expensive high-endurance disk drives that rarely fails

FALSE

# Single Choice

- To open a file `/foo/bar/cs202/test.c` stored in the server, how many lookup messages will be made by the client?
  - A. 1
  - B. 2
  - C. 3
  - D. 4
  - E. 5

E: 4 lookup messages

# Virtualization

- Trap-and-emulate
- x86 virtualization challenges
- Dynamic binary translation
- Shadow paging
- Para-virtualization
  - Xen
  - Dom0
  - CPU/memory/IO virtualization
- Hardware extension

# Single Choice

- The trap-and-emulate approach faces multiple challenges in virtualizing classical x86 architectures (without hardware extensions). Which one is correct?
  - A. The guest OS does not know that it's not running in a privileged mode
  - B. A privileged instruction in the guest OS does not trigger a trap
  - C. x86 does not provide a mechanism to set write-protected pages for shadow paging
  - D. x86's hierarchical page table structure prevents the use of shadow page tables

**B**



# Single Choice

- The trap-and-emulate approach faces multiple challenges in virtualizing classical x86 architectures (without hardware extensions). Which one is correct?
  - A. The guest OS does not know that it's not running in a privileged mode
  - B. A privileged instruction in the guest OS does not trigger a trap
  - C. x86 does not provide a mechanism to set write-protected pages for shadow paging
  - D. x86's hierarchical page table structure prevents the use of shadow page tables

**B**

# Short Answer

- Can shadow paging be faster in address translation than nested/extended paging? Give your answer with a short explanation

# Embedded OS

- TinyOS
  - Concurrency: **Event-driven** architecture; no preemption
  - Modularity: scheduler + graph of components
  - Compiled into one executable
  - Static memory allocation
- Tock
  - Rust - Isolate drivers
  - MPU - Isolate applications
  - Grants - enable dynamic memory without crashing the kernel

# True and False

- TinyOS implements priority-based scheduling for tasks

FALSE

- Tock supports dynamic memory allocation for processes

TRUE

- Tock supports virtual memory

FALSE