

# CS 202: Advanced Operating Systems, Spring 2022

## Midterm Exam

89

Name & UCR Login ID: ~~XXXXXXXXXX~~

### I. True and False (2 points each) 28

1. (☒ / ☐ F) OS has exclusive access to hardware.  $\tau$
2. (☐ T / ☒ F) Processes in the user space can execute halt instructions directly.  $\epsilon$
3. (☐ T / ☒ F) Some interrupt handlers can be implemented in the user space.  $\epsilon$
4. (☒ / ☐ F) To issue a system call in xv6, a user-level process passes the function pointer of the corresponding kernel function to the kernel.
5. (☐ T / ☒ F) xv6 is considered a library OS.  $\epsilon$
6. (☐ T / ☒ F) xv6 implements priority-based scheduling by default.  $\epsilon$
7. ☒ (☐ T / ☒ F) Only one thread or process can be in the RUNNING state on a single CPU.  $\tau$
8. ☒ (☒ / ☐ F) Context switching between the threads of the same process does **not** need to save and restore machine states such as register values.  
need to save and restore Stack and register value
9. (☒ / ☐ F) In UNIX-like systems, fork() is a system call to create a new process.
10. (☐ T / ☒ F) All threads in the same process share the same program counter.
11. (☒ / ☐ F) User-level threads maintain thread control blocks in the user space.  
User-level threads, the thread control block (TCB) is maintained in user space and not in kernel space
12. (☒ / ☐ F) With virtual memory, instructions executed by the CPU use virtual addresses to locate data in physical memory.
13. (☒ / ☐ F) MMU is a must requirement for modern virtual memory subsystems.
14. (☒ / ☐ F) Paging-based virtual memory helps mitigate external fragmentation.  
The solution to external fragmentation is compaction and paging.
15. (☐ T / ☒ F) Preemptive scheduling is always better than non-preemptive scheduling.
16. (☐ T / ☒ F) User-level threads **cannot** implement preemptive scheduling.  
preemptive scheduling by using a timer interrupt mechanism or by yielding control to another thread voluntarily.

## II. Single choice (3 points each) 12

17. Which of the following is correct about virtual memory with paging?
- A. Solves the internal fragmentation problem paging solves external fragmentation and best-fit solves internal
  - B. Makes sharing of data between different processes harder
  - ☒ C. Offers protection domains each process is isolated in its own virtual address space
  - D. Improves cache hit ratio
  - ☒ E. None of the above

18. Consider a system using 32-bit address space, 1MB ( $=2^{20}$  bytes) page size. Page table entry (PTE) is 4 bytes each. If single-level paging is used, what is the size of the page table for each process?

$$2^{32} \times 2^{-20} \times 2^2$$

$12+2=14$

- A. 4096 bytes ( $=2^{12}$ )
  - ☒ B. 16384 bytes ( $=2^{14}$ )
  - C. 262144 bytes ( $=2^{18}$ )
  - D. 1048576 bytes ( $=2^{20}$ )
  - E. 16777216 bytes ( $=2^{24}$ )
19. Which of the following is correct about the Translation Lookaside Buffer (TLB)?
- A. Delays large copies of memory as long as possible in fork()
  - B. Effective for write memory requests but not for reads
  - C. Performance is worse than caching PTEs in L1 cache
  - ☒ D. TLB typically has a high hit ratio due to locality
  - E. TLB often is the largest cache in the system
20. Which of the following is correct about OS extensibility and protection?
- A. DOS-like OSs provide good protection
  - B. DOS-like OSs have more border crossings than monolithic kernels
  - C. Monolithic kernels have more border crossings than microkernels
  - ☒ D. Monolithic kernels are bad at extensibility
  - E. Microkernels are bad at extensibility

21. Which of the following is **NOT** correct about priority-based scheduling?

- ☒ A. Need to know CPU burst time in advance
- B. Priority assignment schemes are often ad hoc
- C. Hard to achieve fair-share scheduling ?
- D. Subject to the priority inversion problem
- E. Effective in favoring important tasks over others

Justification of D This is a situation where a lower priority process holds a resource that a higher priority process needs, causing the high priority process to be blocked and not able to run

22. Which of the following is correct about multiprocessor scheduling?
- ☒ A. Partitioned scheduling (aka. <sup>mult</sup>single queue multiprocessor scheduling) causes higher task migration overhead than global scheduling (aka. <sup>global</sup>mult queue multiprocessor sched.)
  - ☐ B. Partitioned scheduling is better at reclaiming unused processor time ✗
  - ☐ C. Global scheduling tends to achieve better cache affinity ✗
  - ☐ D. Push migration and pull migration cannot be used together in the same OS ✗
  - ☒ E. None of the above

### III. Single choice/short answers (4 points each) 20.

23. Which of the following apply to the SPIN operating system?
- ☐ A. No longer relies on virtual memory to enforce protection
  - ☐ B. Extensions are written in Modula-3
  - ☐ C. Capabilities are implemented directly through the use of pointers
  - ☐ D. Garbage collection may occur at runtime
  - ☒ E. All apply to SPIN
24. Which of the following apply to the Exokernel?
- ☐ A. Exokernel runs library OSs in the kernel space
  - ☒ B. Downloaded code can reduce border crossings
  - ☐ C. Visible resource revocation gives better efficiency as revocations happen more frequently
  - ☐ D. No intervention from Exokernel is needed when traps, faults, or interrupts occur
  - ☐ E. All apply to Exokernel
25. Which of the following is **NOT** L4 microkernel's claim?
- ☐ A. L4 Microkernel refutes the argument that "Microkernels are inherently slow".
  - ☐ B. The high overhead of previous microkernels (e.g., Mach) is an implementation issue
  - ☒ C. As a microkernel is made more portable, its efficiency is likely to be improved. —
  - ☐ D. Major performance issues like switching overhead can be greatly reduced if the implementation considers architecture-specific features
  - ☐ E. All are correct.
26. Which of the following is correct about the lottery scheduler?
- ☐ A. Deterministically fair scheduling ✗
  - ☐ B. Good at achieving short-term fairness
  - ☐ C. Scheduler overhead is high
  - ☒ D. Provides mechanisms to address the priority inversion problem
  - ☐ E. Ticket inflation leads to starvation



27. Which of the following is **NOT** correct about the Linux CFS scheduler?

- A. Considers target latency to ensure response time
- B. Considers minimum granularity to mitigate overhead
- C. Makes a scheduling decision based on virtual time instead of physical time
- ☒ D. Uses a random function to achieve probabilistic-fair scheduling
- E. All above are incorrect

(29)

IV. Short answers (5 points each) - Don't write long; one or two sentences should be enough.

28. Why the `fork()` system call is said to return "twice"? Explain briefly.

*fork returns "twice". It returns child's PID to parent and 0 to the child.*

29. What is the head-of-line blocking problem under FCFS scheduling? (CEP question)

*Head of the line blocking - is when long process can impede short process.*

*Eg:- CPU bound process followed by I/O bound processes*


30. What is the head-of-line blocking problem of FCFS scheduling? (CEP question)

Same as 29

31. Briefly compare main design differences between SPIN and Exokernel. One sentence for each should suffice. (CEP question)

*SPIN - adds application specific functionality in kernel.*

*Exokernel - make the barrier as low as possible.*




32. Consider a system running the following three processes:

Process	Burst Time	Arrival Time
P1	5	0
P2	4	1
P3	3	2


Compute the average turnaround time under (i) Non-preemptive Shortest Job First (SJF) and (ii) Preemptive SJF (PSJF) scheduling. Note: turnaround time = Completion time - Arrival time.

(i) Non-preemptive SJF




$$\frac{(5-0) + (9-1) + (12-2)}{3} = \frac{5+8+10}{3} = \frac{23}{3}$$

(ii) Preemptive



$$\frac{(12-0) + (8-1) + (5-2)}{3} = \frac{12+7+3}{3} = \frac{22}{3}$$



$$\frac{(5-0) + (8-2) + (12-1)}{3} = \frac{5+6+11}{3} = \frac{22}{3}$$

33. Consider a system with three processes:

Process	Arrival Time	Tickets
P1	0	5
P2	0	3
P3	0	2

Fill in the following table to show the scheduling decisions by Stride Scheduling. Assume 30 as the large constant (i.e., stride = 30 / number of tickets)

	Selected process	Pass values after selection		
		P1	P2	P3
Init	N/A	6	10	15
Time 1	P1	12	10	15
Time 2	P2	12	20	15
Time 3	P1	18	20	15
Time 4	P3	18	20	30
Time 5	P1	24	20	30
Time 6	P2	24	30	30

P1 → P2 → P1 → P3 → P1 → P2

[ End of document. Good luck! ]