

Parallel Algorithms

Yihan Sun

Your feedback is welcome!

- **Today is the last lecture with new knowledge, we are almost done!**
 - Most of you are doing very well so far, congrats!
- **Your feedback is very helpful for improving the course (and hence help future UCR CS students)**
 - Provide your feedback on iEval
 - Finish our survey: <https://forms.gle/m6Z698VremJkvCCt7>
 - We want to know how you like the course/topics/OH/assignments
- **If you finish both by Dec 2 (the last lecture), you get 2 bonus candies**
- **Both are anonymous, and to do so you attach the screenshots after you submit**

How are candies counted and used?

- **Bonus for programming homework: up to 10**
- **Bonus for written homework: up to 7**
- **Midterm exam & midterm Challenge: up to 6**
- **UCRPC: up to 5**
- **Class participation: up to 9**
 - Class participation (Nov 16 and today): up to 2
 - Discussion participation (week 4): up to 1
 - In-lecture participation (Q&A): up to 3
 - Online (campuswire) participation (Q&A): up to 3
- **Providing feedback: 2 candies**

How are candies counted and used?

- **Total: likely to be up to 20**
- **Converting to your class participation**
 - If you get 1-3 candies, you are likely to get 1-3 points to your final grade
 - If you get 4-20 candies, we will convert it to 4-10 points, based on some piecewise function
- **Also, those with top 10 most candies and/or top 5 in each session will receive our prizes**
 - Will be given to you on Dec 2's lecture, and you need to come to the classroom and pick up your prizes

- **Lastly: Written HW 5 is due on Nov 30 (Tuesday), one day earlier than usual**
- We need to finalize your grade by the end of the next week

Last Lecture

- **Dynamic multithreading and Scheduler:**


- Dynamic multithreading: only specify parallelism / dependency for tasks
- Scheduler: Help you map your parallel tasks to processors

- **Fork-join**

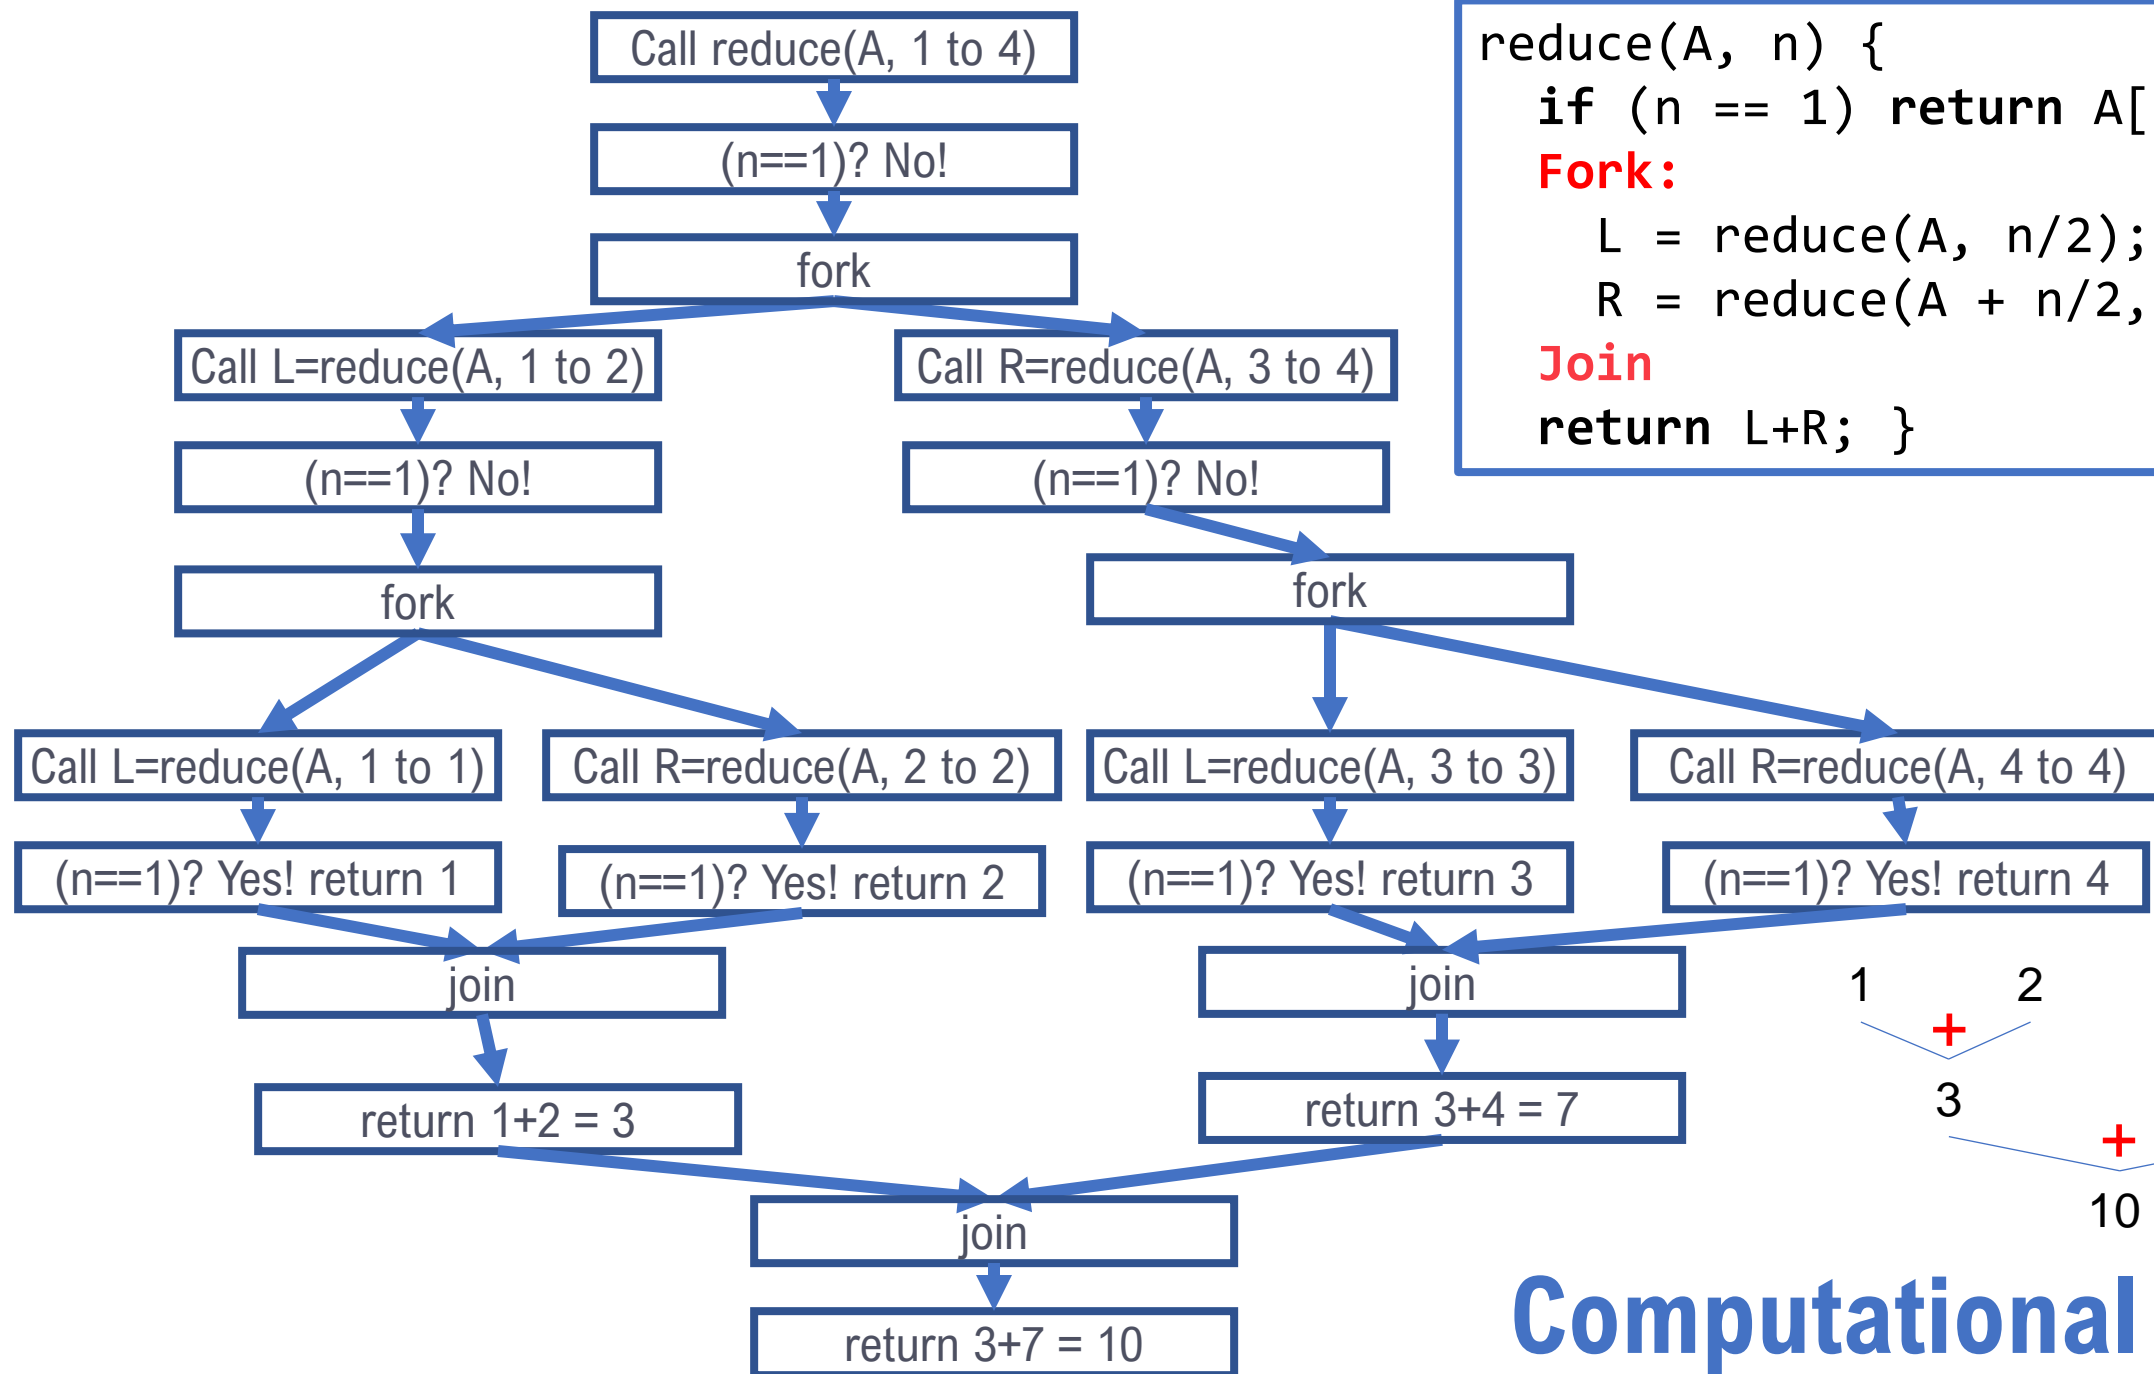
- Fork: two tasks run in parallel
- Join: after all forked threads finish, synchronize them

- **Computational DAG**


- Draw all operations in a DAG
- $A \rightarrow B$ means that operation B must be done after A



```
reduce(A, n) {
    if (n == 1) return A[0];
    In parallel:
        L = reduce(A, n/2);
        R = reduce(A + n/2, n-n/2);
    return L+R;
}
```



```
reduce(A, n) {  
  if (n == 1) return A[0];  
  Fork:  
    L = reduce(A, n/2);  
    R = reduce(A + n/2, n-n/2);  
  Join  
  return L+R; }
```




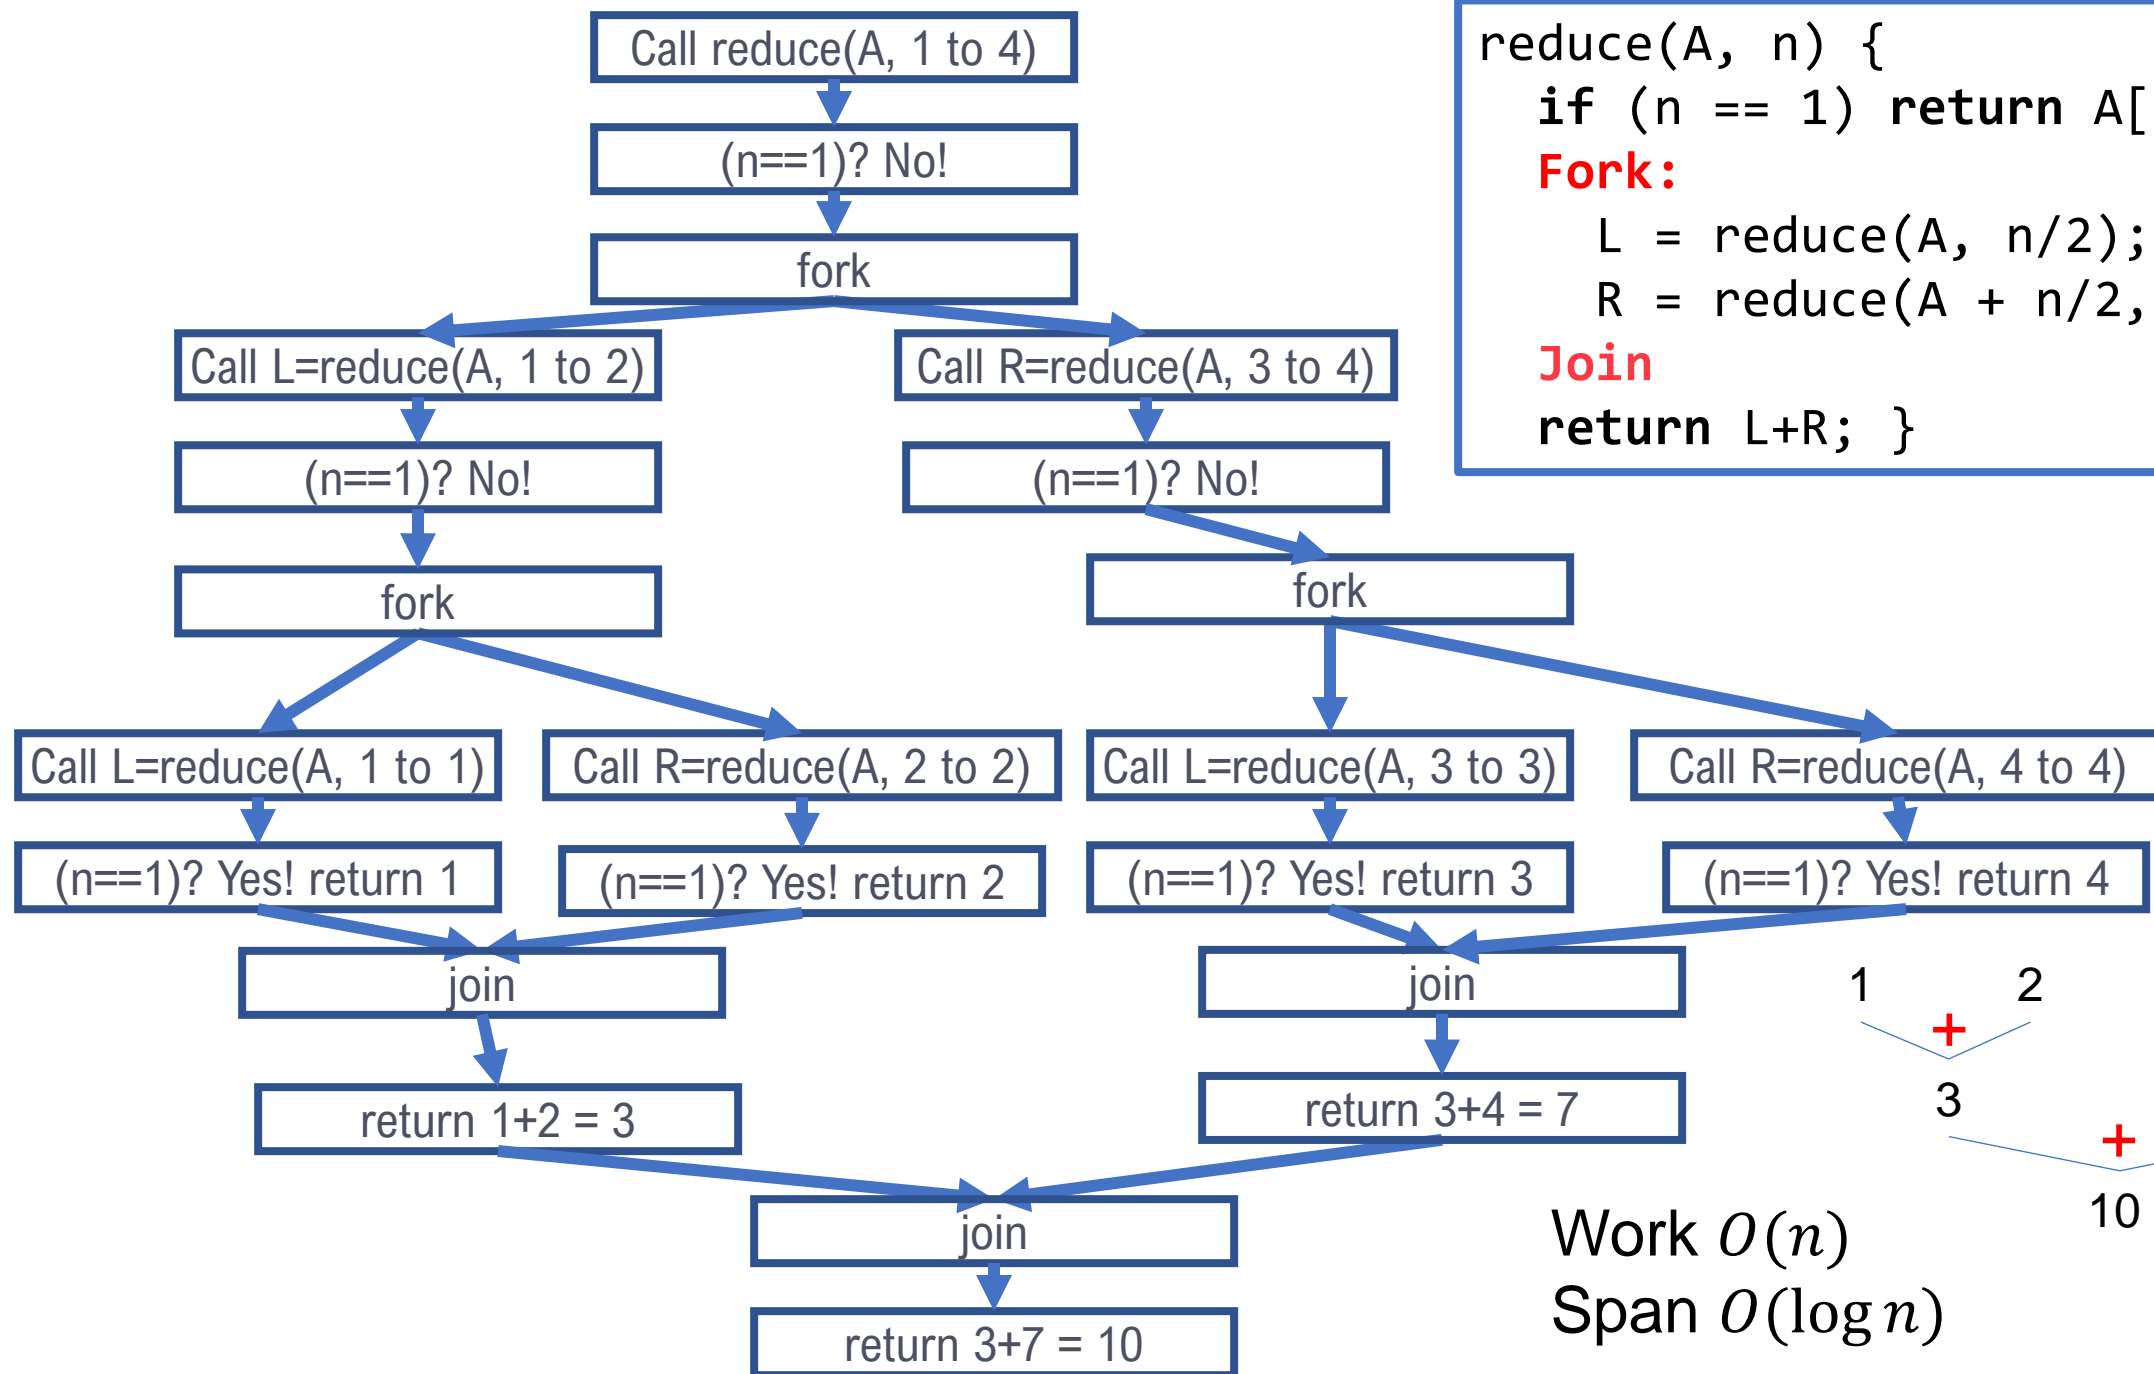
Computational DAG

Last Lecture

- **Work-span**


- Work: total number of operations, sequential complexity
- Span (depth): the longest chain in the computational DAG





```

reduce(A, n) {
  if (n == 1) return A[0];
  Fork:
    L = reduce(A, n/2);
    R = reduce(A + n/2, n-n/2);
  Join
  return L+R; }
  
```



Work $O(n)$
Span $O(\log n)$

Computational DAG


```
reduce(A, n) {  
    if (n == 1) return A[0];  
    In parallel:  
        L = reduce(A, n/2);  
        R = reduce(A + n/2, n-n/2);  
    return L+R;  
}
```

$$W(n) = 2W\left(\frac{n}{2}\right) + \Theta(1)$$
$$S(n) = \Theta(1) + S\left(\frac{n}{2}\right)$$

Work $O(n)$
Span $O(\log n)$

Another way to implement reduce

- $W(n) = \Theta(n) + W\left(\frac{n}{2}\right)$
- $\Rightarrow W(n) = \Theta(n)$



- $S(n) = \Theta(\log n) + S\left(\frac{n}{2}\right)$
- $\Rightarrow S(n) = \Theta(\log^2 n)$

```
reduce(A, n) {  
  if (n == 1) return A[0];  
  if (n is odd) n=n+1;  
  parallel_for i=1 to n/2  
    B[i]=A[2i]+A[2i+1];  
  return reduce(B, n/2); }
```

Needs $\log n$ rounds
to spawn tasks!


How do work and span relate to the real execution and running time?

Scheduling a parallel algorithm

- A DAG with work W and span S can be executed using p processors in time $O\left(\frac{W}{p} + S\right)$
- Both W and S matter!
- For small p , W is more important
- For large p , S is more important

Prefix Sum (Scan)

Prefix sum



The most widely-used building block in parallel algorithm design

Two algorithms to implement a reduce

Can we use the same idea in *reduce* to compute prefix sum?

```
reduce(A, n) {  
    if (n == 1) return A[0];  
    In parallel:  
        L = reduce(A, n/2);  
        R = reduce(A + n/2, n-n/2);  
    return L+R;  
}
```

Divide-and-conquer:

Dealing with the left and right halves recursively

```
reduce(A, n) {  
    if (n == 1) return A[0];  
    if (n is odd) n=n+1;  
    parallel_for i=1 to n/2  
        B[i]=A[2i]+A[2i+1];  
    return reduce(B, n/2); }
```

Decrease problem size:

Shrink the original size into a half

Decrease problem size


Can we use the same idea in *reduce* to compute prefix sum?

```
reduce(A, n) {  
  if (n == 1) return A[0];  
  if (n is odd) n=n+1;  
  parallel_for i=1 to n/2  
    B[i]=A[2i]+A[2i+1];  
  return reduce(B, n/2); }
```

Decrease problem size:

Shrink the original size into a half

Prefix Sum



```
reduce(A, n) {  
  if (n == 1) return A[0];  
  if (n is odd) n=n+1;  
  parallel_for i=1 to n/2  
    B[i]=A[2i]+A[2i+1];  
  return reduce(B, n/2); }
```

Prefix sum of A

A	1	2	3	4	5	6	7	8
---	---	---	---	---	---	---	---	---

1	3	6	10	15	21	28	36
---	---	---	----	----	----	----	----


Prefix sum of B

3	10	21	36
---	----	----	----

B	3	7	11	15
---	---	---	----	----

- Shrink the problem size into $\frac{n}{2}$, possibly in parallel
- Solve the same problem on $\frac{n}{2}$
- Convert the result of the subproblem to the final answer, possibly in parallel


Prefix sum



```
Function Out = PrefixSum(In[1..n]) {
  if (n==1) Out[1] = In[1];
  parallel_for (i=1 to n/2)
    B[i] = In[2i-1]+In[2i]

  C = PrefixSum(B[1..n/2]);

  Out[1] = In[1];
  parallel_for (i=2 to n) {
    if (i%2) Out[i] = C[i/2];
    else Out[i] = C[i/2] + In[i]; } }
```



$O(n)$ work
 $O(\log^2 n)$ span

Another way to solve prefix sum


- **Based on the divide-and-conquer reduce algorithm, we can also design a prefix sum algorithm**
 - $O(n)$ work, $O(\log n)$ span
 - Slightly more complicated

Filtering / packing

Parallel filtering / packing

- Given an array A of elements and a predicate function f , output an array B with elements in A that satisfy f

$$f(x) = \begin{cases} \text{true} & \text{if } x \text{ is odd} \\ \text{false} & \text{if } x \text{ is even} \end{cases}$$



Parallel filtering / packing

- Sequentially, we just read the array from left to right and put those satisfying f into an input array
- How can we know the length of B in parallel?
 - Count the number of red elements – parallel reduce
 - $O(n)$ work and $O(\log n)$ span


$A =$

4	2	9	3	6	5	7	11	10	8
0	0	1	1	0	1	1	1	0	0

Parallel filtering / packing

- How can we know where should 9 go?
 - 9 is the first red element, 3 is the second, ...

```
Filter(A, n, B, f) {  
    new array flag[n], ps[n];  
    parallel_for (i = 1 to n) {  
        flag[i] = f(A[i]);  
    }  
    ps = prefix_sum(flag, n);  
    parallel_for(i=1 to n) {  
        if (f(A[i]))  
            B[ps[i]] = A[i];  
    }  
}
```




Parallel Filtering/packing

- $O(n)$ **work**,
- $O(\log^2 n)$ **span** (if we use the scan algorithm above)
- **Can be** $O(\log n)$ **span** with a better scan algorithm

Parallel Quicksort

Sequential quicksort algorithm

- Find a random pivot p in the array (e.g., the middle one)
- Put all elements in A that are $< p$ on the left, and all elements in A that are $\geq p$ on the right* (“**partition**”)




• How to do partition?

*: Usually we need to find $>$, $<$ and $=$. Here for simplicity we assume keys are distinct so that doesn't make much difference

Sequential sorting algorithms: quicksort

- How to move elements around? (using 6 as a pivot)



```
Partition(A, n, x) {  
    i = 0; j = n-1;  
    while (i < j) {  
        while (A[i] < x) i++;  
        while (A[j] > x) j++;  
        if (i < j) swap A[i] and A[j];  
        i++; j--;  
    }  
}
```

- $O(n)$ time for one round

Sequential quicksort

- Use a pivot and partition the array into two parts
- Sort each of them recursively

```
qsort(A, n) {  
    t = partition(A, A[random()]);  
    qsort(A, t);  
    qsort(A+t, n-t);  
}
```

Parallel quicksort

- Use a pivot and partition the array into two parts
- Sort each of them recursively, **in parallel**

```
qsort(A, n) {  
    t = partition(A, A[random()]);  
    In parallel:  
    qsort(A, t);  
    qsort(A+t, n-t);  
}
```

Parallel quick sort

- The partitioning algorithm costs $O(n)$ time. So even if the problem is always perfectly partitioned


- $W(n) = 2W\left(\frac{n}{2}\right) + O(n)$
- $S(n) = S\left(\frac{n}{2}\right) + O(n)$
- $S(n) = O(n)?$

```
qsort(A, n) {  
    t = partition(A, A[random()]);  
    In parallel:  
        qsort(A, t);  
        qsort(A+t, n-t);  
}
```

- Have to partition in parallel!

Application of filter: partition in quicksort

- For an array A , move elements in A smaller than k to the left and those larger than k to the right



- Left: all elements < 6
- Right: all elements ≥ 6

Using filter for partition

(Looking at the left part as an example)

using 6 as a pivot

A 6 2 9 4 1 3 5 8 7 0

flag 0 1 0 1 1 1 1 0 0 1

A X 2 X 4 1 3 5 X X 0

Prefix sum
of flag 0 1 1 2 3 4 5 5 5 6

pack 2 4 1 3 5 0 6 9 8 7

```
Partition(A, n, k, B) {  
    new array flag[n], ps[n];  
    parallel_for (i = 1 to n) {  
        flag[i] = (A[i]<k);  
    }  
    ps = scan(flag, n);  
    parallel_for(i=1 to n) {  
        if (A[i]<k)  
            B[ps[i]] = A[i];  
    }  
    //symmetric for the right half  
}
```

Predicator: if $A[i] < \text{pivot}$

Parallel quicksort

```
qsort(A, n) {  
    t = parallel_partition(A, A[random()]);  
    In parallel:  
        qsort(A, t);  
        qsort(A+t, n-t);  
}
```

- **Work**

- Exactly the same as sequential version
- $O(n \log n)$ in expectation

- **Span**

- $O(\log n) \times (\text{\#rounds of recursions}) \approx O(\log^2 n)$
 - $O(\log^2 n)$ whp.

- **Actually, quicksort is not the most commonly used sorting algorithm in parallel**

- It is no longer in-place: lose the advantage

Summary

- **Parallel algorithms: identify / eliminate dependencies between operations**
- **“Time complexity” evaluation**
 - Work and span
 - Work: the total number of operations performed
 - Span: the longest dependency chain
- **Useful primitives:**
 - Reduce: the sum of all elements in an array
 - Scan: prefix sum
 - Filter/pack: extract elements with certain property
 - Can be used to implement parallel partition and quicksort

Summary

Useful ideas:

- **Divide-and-conquer: work on two/multiple subproblems in parallel**
 - Reduce
 - Quicksort
 - (Also parallel merge sort, but we need a parallel merging algorithm)
- **Decrease and conquer**
 - Decrease the size of the problem in parallel (polylog span), usually to $\frac{n}{2}$
 - Solve the smaller problem
 - Finish in $\log n$ rounds. So total span will still be polylog