

Analyzing algorithms

Yihan Sun

Course announcement

- **Entrance Exam due this weekend**
 - Try to start working on it soon if you haven't
 - Come to the office hours if you need help
- **Course policy test: due next Tuesday**
 - 1 point to your final grade, and required
 - Resubmit-able multiple choices problems
- **Regarding course-related logistics**
 - Contact Zijin for non-homework-related questions
 - Contact Xiangyun for homework-related questions

Collaboration: the “whiteboard” policy

- You are welcome to chat with each other (also welcome to come to OHs), but you come with nothing and leave with nothing
- When you type your answers / code, it must be done on your own. It must be close-book. It must be typed by you, word by word.
- Any violation may result in severe outcome. Usually -100% current/all homework assignment score, fail the course, report to the university, the university may make further decisions
- Must cite if any idea is from other sources, including people, books, websites, AI, etc.


Analyzing algorithms

Yihan Sun

Analyzing algorithms

- Predict how your algorithm performs in practice

- **Criteria:**

- Running time  Time complexity
- Space usage
- Cache I/O
- Disk I/O
- Network bandwidth
- Power consumption
- Lines of codes

What are good algorithms?

- Tale about Gauss



One day Gauss's teacher asked his class to add together all the numbers from 1 to 100, assuming that this task would occupy them for quite a while. He was shocked when young Gauss, after a few seconds thought, wrote down the answer 5050. (source: <https://nrich.maths.org/2478>)

Other students:


```
sum = 0;  
for (int i = 1; i <= n; i++)  
    sum += i;
```

$O(n)$ time complexity

Gauss:

$\text{sum} = (1+n)*n/2;$

$O(1)$ time complexity



$$(1+10) + (2+9) + (3+8) + (4+7) + (5+6) = ?$$

Computational Model

What is time complexity?

```
s = 0
for (i = 1 to n) {
    s = s + A[i]
}
```

$O(n)$ time complexity

```
sum(A, n) {
    if (n == 1) return A[0];
    L = sum(A, n/2);
    R = sum(A + n/2, n-n/2);
    return L+R;
}
```

$O(n)$ time complexity


- Why the time complexity is $O(\mathbf{n})$? What are we counting?

What is “time complexity”?

- Count the number of “instructions” in the algorithm
- **Random Access Machine (RAM) model**
 - We have an arbitrarily large memory
 - We can
 - do arithmetic calculation
 - Read/write to a memory location given the address
 - Every operation takes **unit time**

Random-Access Machine (RAM)

- Unit cost for:
 - Any instruction on w -bit words (how large do we need for w ?)
 - Read/write a single memory location from an infinite memory
- The cost measure: **time complexity**



What is a **computational model**?

- What resource do we have?
- What operations are we allowed to do?
- What is the cost of each operation?

What is “time complexity”?

- To estimate the time needed for an algorithm
- To define the “cost” of an algorithm, we first need to define what “costs” time
- **Computational model**

Algorithm A

Cost bound $f(A)$

Computational
Model

Cost measure f

What is a **computational model**?

- What resource do we have?
- What operations are we allowed to do?
- What is the cost of each operation?

What is time complexity?

```
s = 0
for (i = 1 to n) {
    s = s + A[i]
}
```

$O(n)$ time complexity

```
sum(A, n) {
    if (n == 1) return A[0];
    L = sum(A, n/2);
    R = sum(A + n/2, n-n/2);
    return L+R;
}
```

$O(n)$ time complexity

- Why the time complexity is $O(n)$? What are we counting?
 - The number of operations
- Why the time complexity is $O(n)$? Why we omit the constants and lower-order terms? Why we use “asymptotic analysis”?

Anything else we need?

- Count the number of “instructions” in the algorithm
- **Random Access Machine (RAM) model**
 - We have an arbitrarily large memory
 - We can
 - do arithmetic calculation
 - Read/write to a memory location given the address
 - Every operation takes **unit time**

```
sum = 0;  
for (int i = 1; i <= n; i++)  
    sum += i;
```

$3n+2$ operations?

```
sum = (1+n)*n/2;
```

3 operations?

What is “time complexity”?

- **Random Access Machine (RAM) model**
 - Every operation, including memory access, arithmetic operations, etc., takes **unit time**
- To make our life easier, we only analyze **order of the cost**, and omit
 - Lower-order terms
 - Constants
- We care about **how faster a function grows** for large n

Asymptotic Analysis

Asymptotic notation

```
sum = 0;  
for (int i = 1; i <= n; i++)  
    sum += i;
```

$3n+2$ operations?

$O(n)$ operations

```
sum = (1+n)*n/2;
```

3 operations?

$O(1)$ operations

- OK, then what does big-O mean?

Asymptotic notations

- Big-O: asymptotically **smaller than or equal to** \leq
larger than or equal to \geq
smaller than $<$
larger than $>$
equal to $=$
- n is $O(n)$
- $3n + 2$ is $O(n)$
- $\log n + \sqrt{n} + 4$ is also $O(n)$
- What happens if we want to say other cases?

Analogy to real numbers

Functions	Real numbers
$f(n) = O(g(n))$	$a \leq b$
$f(n) = \Omega(g(n))$	$a \geq b$
$f(n) = \Theta(g(n))$	$a = b$
$f(n) = o(g(n))$	$a < b$
$f(n) = \omega(g(n))$	$a > b$

Popular Classes of Functions

Example

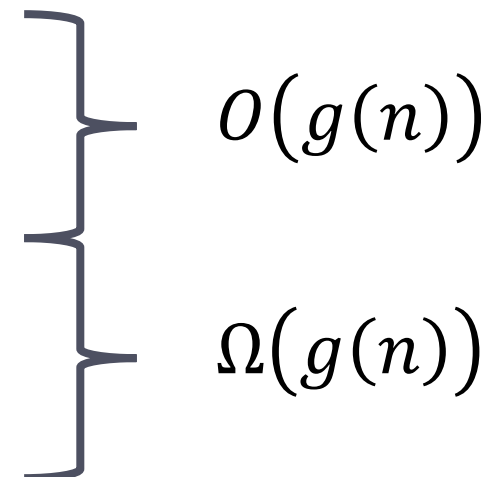
• Constant:	$f(n) = \Theta(1)$	1, 10, 100000000
• Logarithmic:	$f(n) = \Theta(\log(n))$	$\log n, \log n + 3 \log \log n$
• Poly-logarithmic:	$f(n) = O(\log^k n)$	$\log^2 n, \log^9 n + 8$
• Sublinear:	$f(n) = o(n)$	$\log n, \sqrt{n}, n^{1/5}$
• Linear:	$f(n) = \Theta(n)$	$n, 5n + \log n$
• Super-linear:	$f(n) = \omega(n)$	$n^3, n \log n$
• Quadratic:	$f(n) = \Theta(n^2)$	$n^2, 3n^2 + n$
• Polynomial:	$f(n) = O(n^k)$	$n^3 + 2n^2 + 4, 4n^5$
• Exponential:	$f(n) = \Theta(k^n)$	2^n

($k \geq 0$ is a constant)

Compare two functions

- $\lim_{n \rightarrow \infty} \frac{f(n)}{g(n)}$

- 0: $f(n) = o(g(n))$
- **Constant** $c > 0$: $f(n) = \Theta(g(n))$
- ∞ : $f(n) = \omega(g(n))$



$O(g(n))$

$\Omega(g(n))$

Commonly-used functions

For large enough n , we have (asymptotically):

$$\begin{array}{ccccccc} c > 0 & & c_1 > 1 & 0 < c_2 < 1 & & c_3 > 1 & c_4 > 1 \\ \bullet \ c < \log \log n < \log n < \log^{c_1} n < n^{c_2} < n < n \log n < n^{c_3} < c_4^n < n^n \end{array}$$

Analyze running time

- **Random Access Machine (RAM) model**
 - Every operation, including memory access, arithmetic operations, etc., takes **unit time**
- We only care about the **order of the cost** for simplicity, and omit
 - Constants
 - Lower-order terms
- We usually consider **worst-case** running time for **general input**
 - In some cases, we also analyze bounds with probabilistic guarantees (e.g., **average** running time for randomized algorithms)

If you are not familiar with these notations

- **Read CLRS Chapter 3: “Growth of Functions”**
 - Definitions of the asymptotic notation
 - How to compare the growth of two functions
 - What are the classic “classes” of functions

How does the asymptotic notation relate to the computational model (algorithm analysis)?

Algorithm A

Cost bound $f(A)$



**Computational
Model**



Cost measure f

Questions (true or false):

- The goal of defining computational models is to allow for asymptotic analysis
- We have to use asymptotic notations when analyzing algorithms
 - It is almost true for time complexity, since it is too sloppy to give a distinct weight to each operation

Is RAM model perfect?

$$n = 10^9$$

`for (int i = 0; i < n; i++) A[i] = A[i] + 1;`

0.141072s

`for (int i = 0; i < n; i++) A[i] = (long long)(i)*4323 % n + 1;`

0.580809s

`for (int i = 0; i < n; i++) A[i] = A[(long long)(i)*4323 % n] + 1;`

3.25008s

How long would the other for-loop take?

A.0.14s

D.1.0s

B.0.2s


E.3.0s

C.0.6s

F.5.0s



Not all CPU operations are created equal



Distance which light travels while the operation is performed





Image from ithare.com:

<http://ithare.com/infographics-operation-costs-in-cpu-clock-cycles/>


I/O model

- Access memory is expensive
- But we have cache in our machine!
- We count the cost only if it is a “cache miss”
 - Accessing data in the cache is free
 - # of transfers between cache and memory




Asymmetric RAM model

- Write to the memory is more expensive than read



Parallel computational models

- Multiple “threads/processes” can do computation together, and share a memory
- Computation starts with one thread
- Each thread is like a regular RAM
- Each thread can also “create” another thread to run in parallel
- What should we count to evaluate the running time in this case?



So.. Why are we still using RAM model?

- In the sequential setting, it actually provide a way to analyze algorithms nicely (e.g., the comparison of quicksort and selection sort)
- Many other important models are based on that (e.g., a parallel model would assume multiple threads/processors, each behaving like a RAM)

Analysis

- Time complexity and RAM model
- Other models exist!
- Analyzing algorithms => time complexity, other cost (space, I/O, etc.), how efficient an algorithm is
- Analyzing PROBLEMS => lower bound of a problem, how hard an algorithm is?
 - Will be covered next week

Lower bound analysis

Analyzing “problems” – lower bound analysis

- Given a problem, what is the smallest cost we need to pay?
- Given an array, what is the “best complexity” you can get to sort it?
 - We know algorithms of running time $O(n \log n)$
 - But is it the “best”? Does there exist better algorithms? If not, why?
- Given a sorted array and an element x , what is the “best complexity” can you get to find the rank of x ?
 - We know binary search needs $O(\log n)$ time
 - But is it the “best”? Does there exist better algorithms? If not, why?

Example about lower bound proof: Finding the max of an array

- **Input:** an array of distinct elements
- **Model:** we only care about comparisons: each comparison is a unit cost
- How many comparisons at least do you need to find the **max of the array**?
- Let's first consider some algorithms for this.
- **Idea 1:** always compare to the champion

```
ans = a[1];  
for (i = 2 to n) {  
    if (a[i] > max) max = a[i];  
}
```

$n - 1$ comparisons


Example about lower bound proof: Finding the max of an array

- Idea 1: always compare to the champion

```
ans = a[1];  
for (i = 2 to n) {  
    if (a[i] > max) max = a[i];  
}
```

$n - 1$ comparisons

- Idea 2: single-elimination tournament?



$n - 1$ comparisons

Finding the max of an array

- Can we use fewer than $n - 1$ comparisons to find the max of an array?
- Principle:
 - If an element is **not compared to others**, we cannot guarantee we find the max
 - If an element **loses any comparison**, it cannot be the max
 - If an element **never lose a comparison**, it's possible to be the max
- Proof:
 - Call an element “**free**” if it hasn't been compared to anyone (need to verify!)
 - Call an element “**bad**” if it loses at least one comparison (rule it out!)
 - Call an element “**top**” if it never loses any comparison (a candidate!)
 - Initial status: 0 top, n free, 0 bad
 - Final status: 1 top, 0 free, $n - 1$ bad
 - Only when we reached this final status, we find the max

Finding the max of an array – lower bound proof

- Preliminary

- “top”: never loses any comparison
- “bad”: lost at least one comparison
- “free”: never compared to others
- Initial status: 0 top, n free, 0 bad
- Final status: 1 top, no free, $n - 1$ bad

- What happens when we compare two elements?

	#top	#bad	#free	
top vs. top	-1	+1	=	
top vs. bad	=	=	=	(top wins)
	-1	+1	=	(bad wins)
top vs. free	=	+1	-1	(top wins)
	-1+1	+1	-1	(free wins)
bad vs. bad	=	=	=	
bad vs. free	=	+1	-1	(bad wins)
	+1	=	-1	(free wins)
free vs. free	+1	+1	-2	

Each comparison increases at most 1 bad!

=> We need at least $n - 1$ comparisons to get the final status ($n - 1$ bad)

Finding the max of an array – lower bound proof

- To get an optimal algorithm, we have to **guarantee to increase #bad by 1** in every comparison
- Only compare
 - top to top, or
 - free to free, or
 - top to free
 - (guarantee to increase #bad by 1)
- Algorithm:
 - If there are more than two **free** or **top** elements, **compare** them
 - Until there is only one top

	#top	#bad	#free	
top vs. top	-1	+1	=	
top vs. bad	=	=	=	(top wins)
	-1	+1	=	(bad wins)
top vs. free	=	+1	-1	(top wins)
	-1+1	+1	-1	(free wins)
bad vs. bad	=	=	=	
bad vs. free	=	+1	-1	(bad wins)
	+1	=	-1	(free wins)
free vs. free	+1	+1	-2	

Each comparison increases at most 1 bad!

=> We need at least $n - 1$ comparisons to get the final status ($n - 1$ bad)


Example about lower bound proof: Finding the max of an array

- Idea 1: always compare to the champion Every comparison is **top to free**!

```
ans = a[1];  
for (i = 2 to n) {  
    if (a[i] > max) max = a[i];  
}
```

$n - 1$ comparisons

- Idea 2: single-elimination tournament?



First level: free to free

Others: top to top

$n - 1$ comparisons

Upper bound vs. lower bound

- An **upper bound** $f(n)$ of the cost of a problem means there exists an algorithm that takes at most $f(n)$ steps on any input of size n
 - So, given this problem, we can just run this algorithm to get an answer
 - $f(n)$ is guaranteed to be sufficient: we don't need more than $f(n)$ costs
- A **lower bound** of $g(n)$ means for any algorithm, there exists an input for which all algorithm takes at least $g(n)$ steps on that input
 - Whatever algorithm you use, you cannot get better than $g(n)$!!
- When upper bound meets lower bound...
 - An algorithm has cost $f(n)$, and the best you can do is $g(n) = f(n)$
 - That's an **optimal** algorithm!

Upper bound vs. lower bound

- Finding max of an array

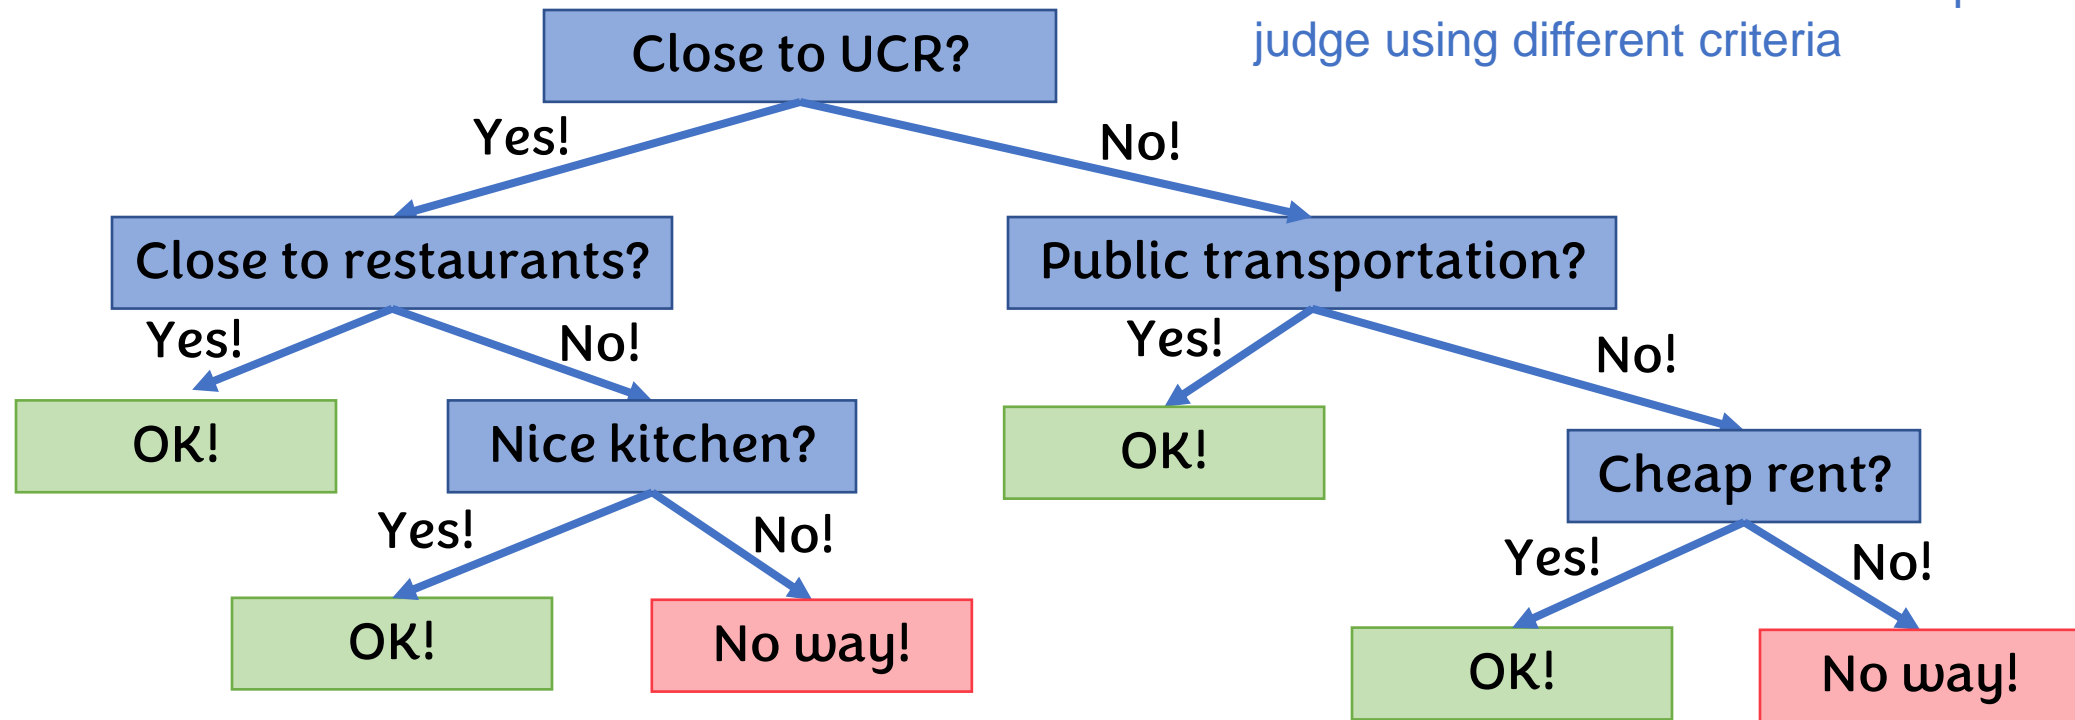
- Lower bound: $n - 1$ comparisons (we just proved that)
 - If you use fewer than $n - 1$ comparisons, you cannot find the answer
- Upper bound: the compare-to-champion algorithm does exactly $n - 1$ comparisons
 - If you want to find the answer, you don't need more than $n - 1$ comparison
- Upper bound meets lower bound!
- Comparison-to-champion is an **optimal algorithm (w.r.t. #comparisons)** to find the max

Decision trees and lower bound proof

Decision trees: partition the decision space and conquer

- You want to rent a house in Riverside, you finally find a candidate, and you need to decide...

Similar to divide-and-conquer:
based on the answer to the first question, we judge using different criteria



See how AI uses a decision-tree-like idea to guess what you are thinking: <http://20q.net>

What question will you ask?

elephant

TV

bee

clock

chair

Joshua tree

rose

computer

- **Is it a plant?**
 - Yes: Joshua tree, rose. No: TV, clock, chair, elephant, bee, computer
- **Does it use electricity?**
 - Yes: TV, clock, computer. No: elephant, bee, chair, Joshua tree, rose
- **Can you find it in this classroom?**
 - Yes: chair, computer. No: elephant, TV, bee, clock, Joshua tree, rose
- **Does it have legs?**
 - Yes: elephant, chair, bee. No: TV, clock, Joshua tree, rose, computer
- **Is it alive?**
 - Yes: elephant, bee, Joshua tree, rose. No: TV, clock, chair, computer

Lower bound proof using a “decision tree”

Sorting lower bound

- What is the minimum number of operations we need to sort n elements?
- This question is too vague to be answered
 - We know $O(n \log n)$ sorting algorithms: they are “almost” the best sorting algorithms we know of...
 - But can we use $O(n)$ to do this? Well, if your elements are integers in $[1, n]$...
 - `for (i = 1 to n) count[A[i]]++;`
 - `for (x = 1 to n) {`
 - `for (j = 1 to count[x])`
 - `Print x;`
 - `}`
 - How should we formalize the question??

We need a formal approach

- Look at computational models which specify exactly which operations may be performed on the input, and what they cost
 - E.g., performing a comparison, or swapping a pair of elements
- An upper bound of $f(n)$ means the algorithm takes at most $f(n)$ steps on any input of size n
- A lower bound of $g(n)$ means for any algorithm there exists an input for which the algorithm takes at least $g(n)$ steps on that input

We need a formal approach

- Look at computational models which specify exactly which operations may be performed on the input, and what they cost
 - E.g., performing a comparison, or swapping a pair of elements
- For sorting algorithms, we usually use the “**comparison model**”
 - We only count the number of comparisons used in the algorithm
 - (similar to the finding-max problem, we also used comparison model)
- Why we study comparison-based sort?
 - General: no constraint on input type (integer, real, string, positive or negative, pair, complicated struct, key range, hashable or not)... as long as comparable!

Sorting in the Comparison Model

- In the **comparison model**, we have n items in some initial order
- An algorithm may compare two items (asking: is $a_i > a_j$?) at a cost of 1
 - Moving the items is free
- No other operations allowed, such as XORing, hashing, etc.
- Sorting: given an array $a = [a_1, \dots, a_n]$, output a permutation π so that $[a_{\pi(1)}, \dots, a_{\pi(n)}]$ in which the elements are in increasing order
- A sorting algorithm based on comparisons is called **comparison sort**

Lower bound for sorting

- **Of course 1 is a lower bound...**
 - You cannot guarantee to sort the entire array using 1 comparison!
- **Of course $n-1$ is a lower bound...**
 - We just proved that just finding the maximum value needs $n-1$ comparisons
- **But... Can we show “better” lower bounds?**
- **We are usually interested in tight lower bounds (the tighter, the better)**
- **For sorting, we can actually show that $\Omega(n \log n)$ is a lower bound**

Sorting lower bound in the Comparison Model

- **Theorem:** Any deterministic comparison sort algorithm must perform at least $\Omega(n \log n)$ comparisons to sort n elements in the worst case
 - More precisely, for any sorting algorithm A with size $n \geq 2$, the #comparisons needed is $\log_2(n!)$
 - i.e., for any sorting algorithm, there exists an input I of size n so that A makes $\geq \log_2 n! = \Omega(n \log n)$ comparisons to sort I
- Proof is information-theoretic

Lower bound proof outline

- In total, there are $n!$ possible inputs (permutations) for an array of size n
- We need to identify the case for the input
- What can we do via a comparison?

All permutations of ranks

a_1	a_2	a_3	a_4
1	2	3	4
1	2	4	3
1	3	2	4
1	3	4	2
1	4	2	3
1	4	3	2
2	1	3	4
.....			
4	3	2	1

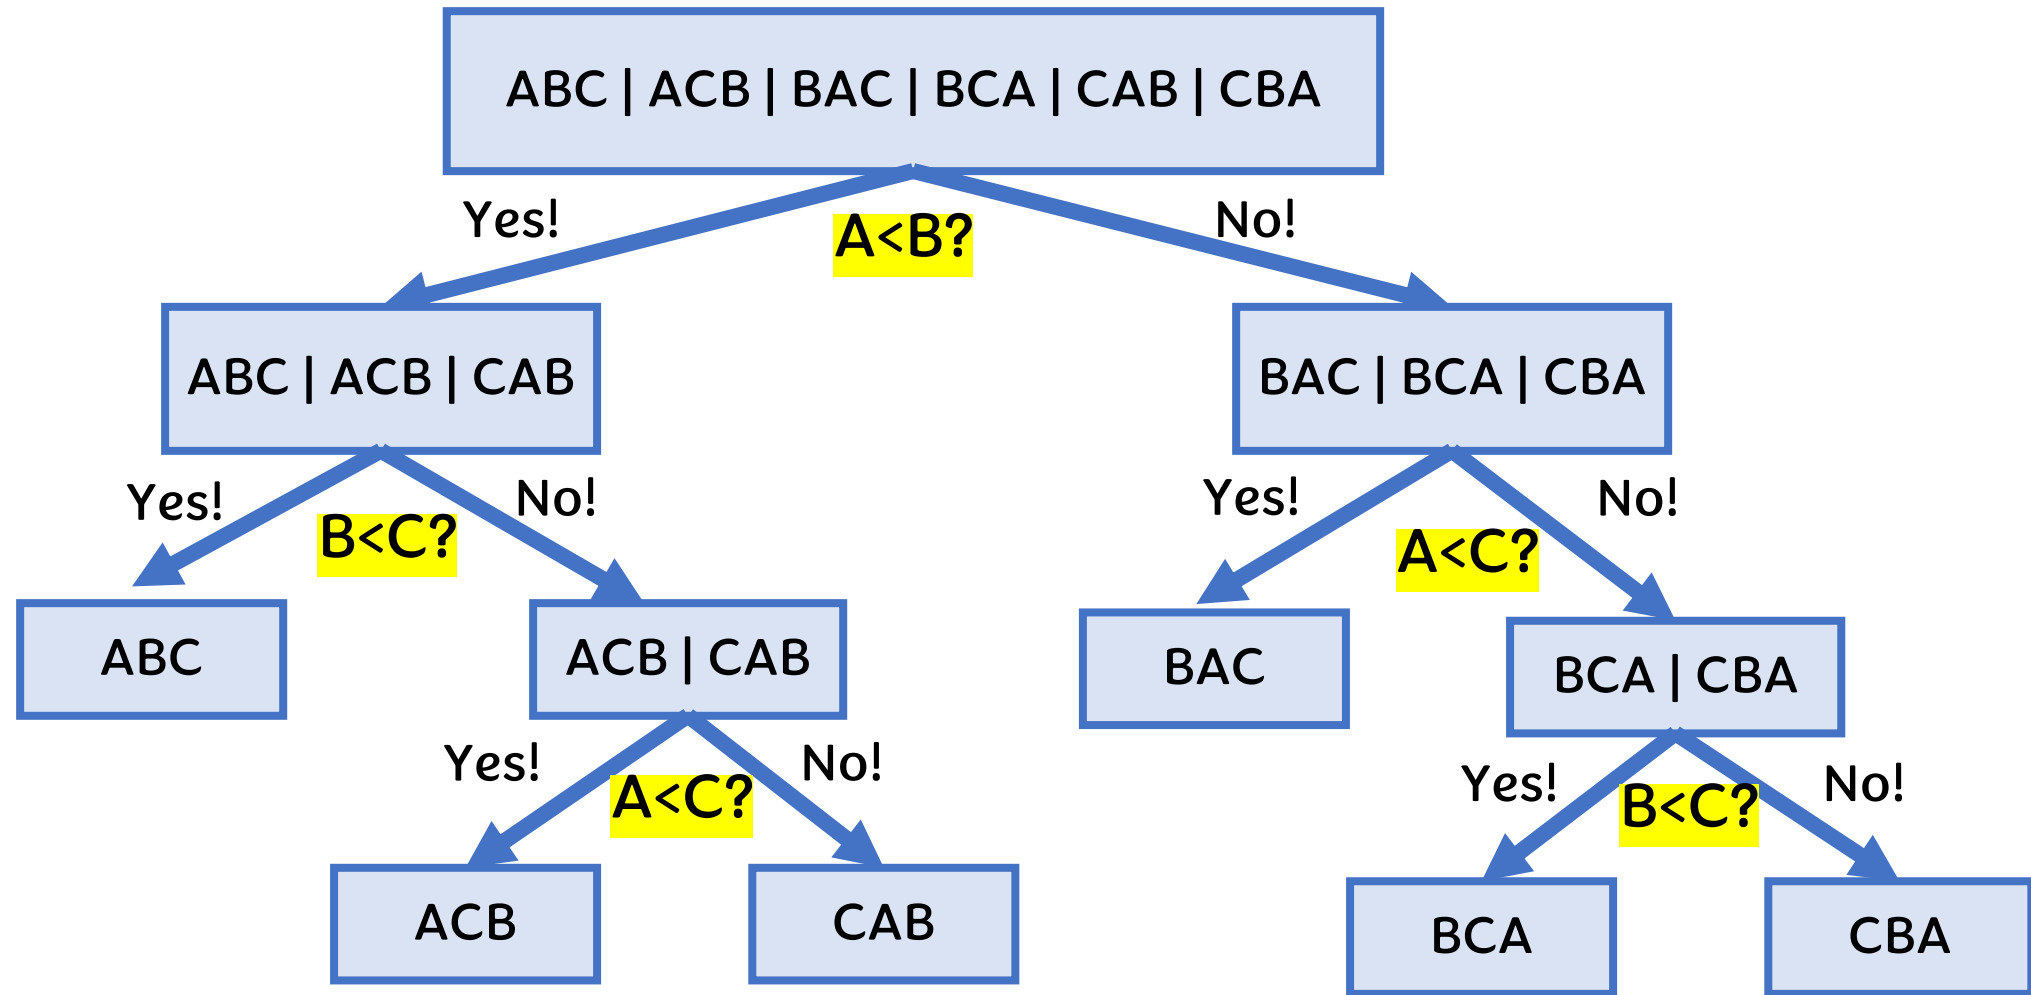
Lower bound proof outline

- In total, there are $n!$ possible inputs (permutations) for an array of size n
- We need to identify the case for the input
- What can we do via a comparison?
 - We can rule out some of the input cases
 - We have to repeat this computation until one case is left

All permutations of ranks

a_1	a_2	a_3	a_4
1	2	3	4
1	2	4	3
1	3	2	4
1	3	4	2
1	4	2	3
1	4	3	2
2	1	3	4
.....			
4	3	2	1

$$a_2 < a_3$$

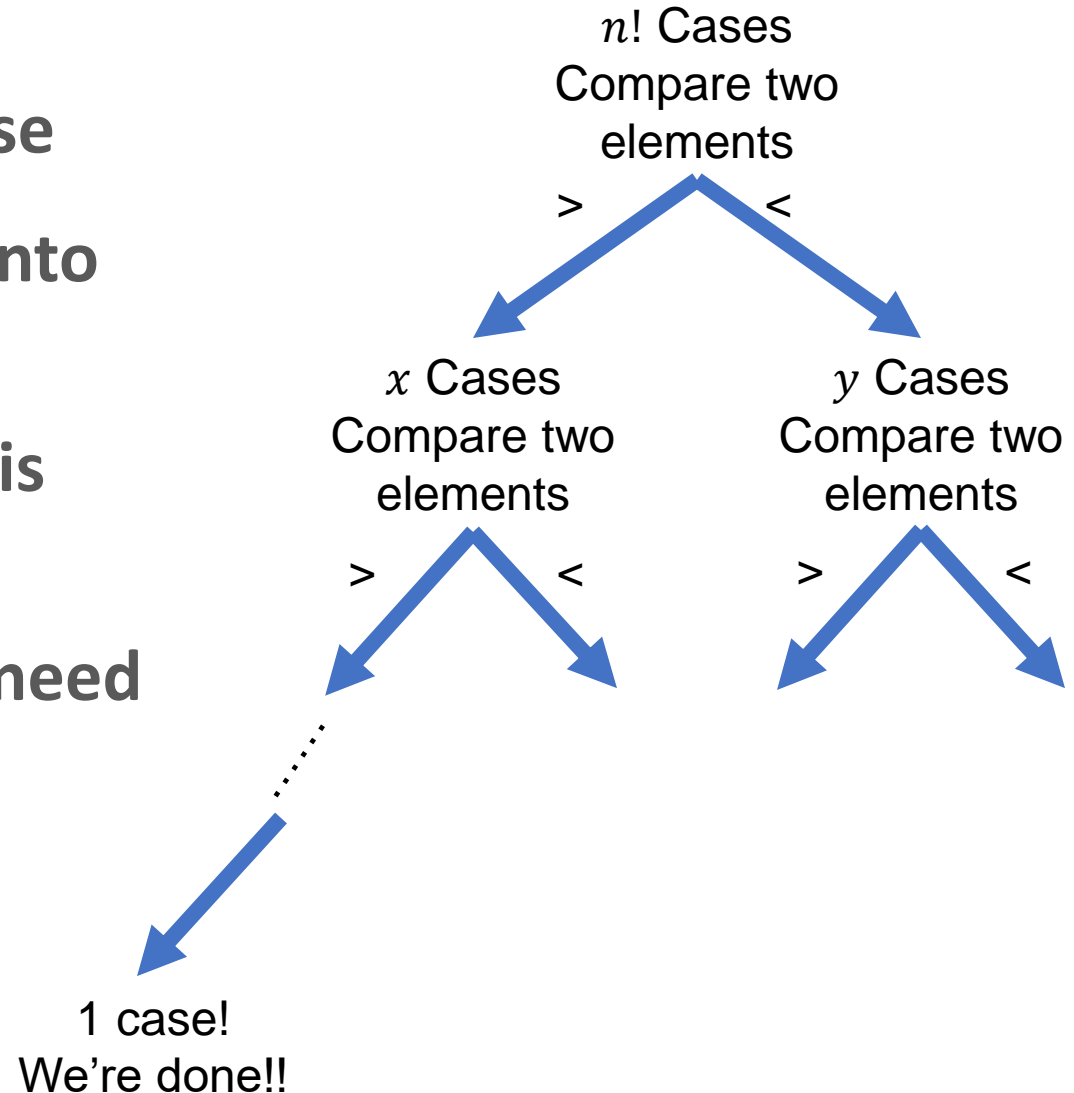


To find the correct solution for any input, we need to distinguish all possible $n!$ Input cases!

Lower bound for sorting algorithms

- We need to identify the input case
- Every comparison split all cases into two parts
- We need to have $n!$ leaves for this decision tree
- What's the number of levels we need for the deepest branch?

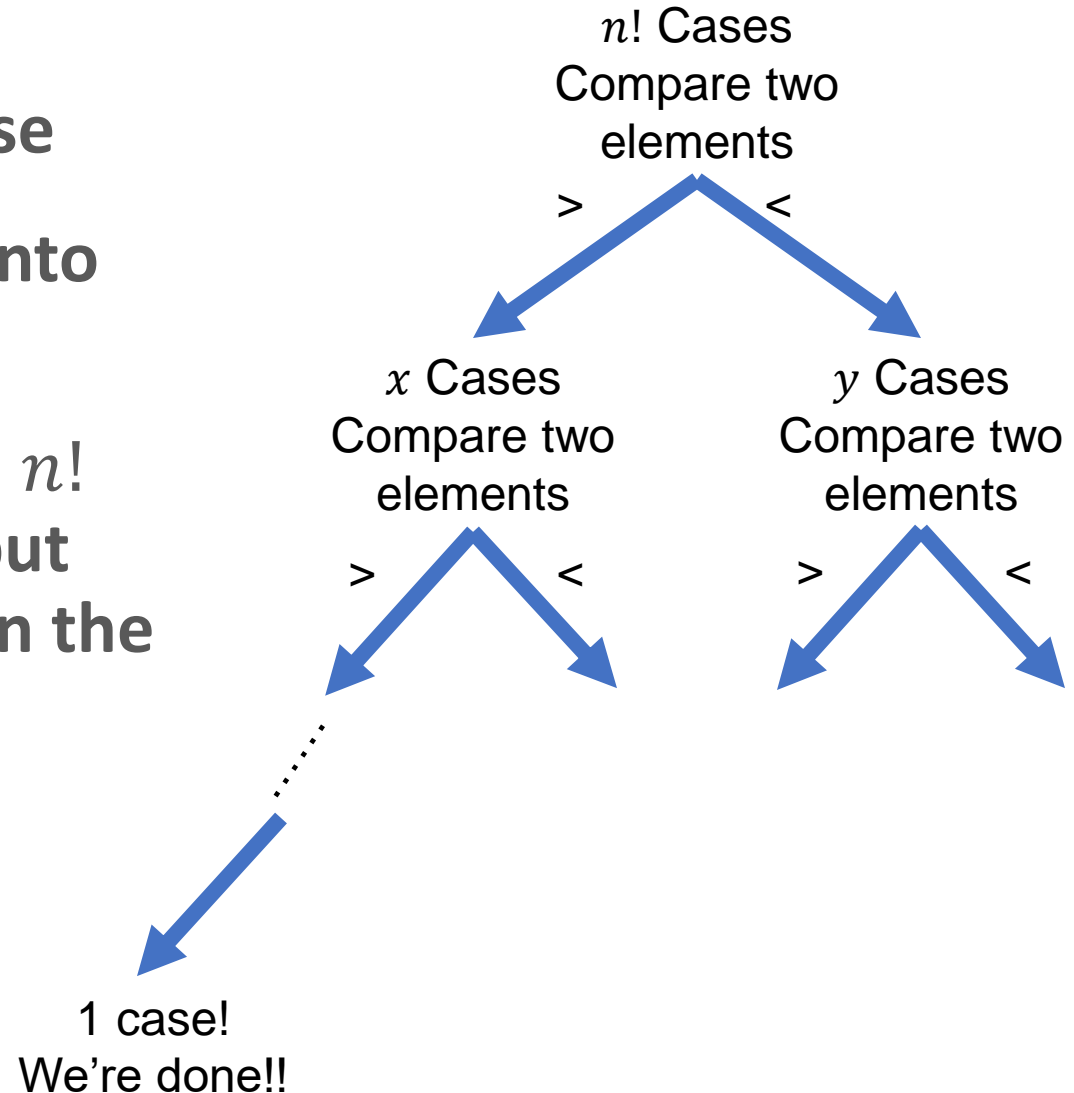
$$\log_2 n! = \Theta(n \log n)$$



Lower bound for sorting algorithms

- We need to identify the input case
- Every comparison split all cases into two parts
- Information-theoretic: need $\log_2 n!$ bits of information about the input before we can correctly decide on the output

$$\log_2 n! = \Theta(n \log n)$$



Why $\log_2 n! = \Theta(n \log n)$?

- $\log_2 n! = \log_2 n + \log_2(n-1) + \dots + \log_2 1 < n \log_2 n = O(n \log n)$
- $\log_2 n! > \log_2 n + \log_2(n-1) + \dots + \log_2 \frac{n}{2} > \frac{n}{2} \log_2 \frac{n}{2} = \Omega(n \log n)$

- We can also use the Stirling's formula:

$$n! = \sqrt{2\pi n} \left(\frac{n}{e}\right)^n + O(n)$$

- So $\log_2 n! = n \log_2 n - n \log_2 e + O(\log_2 n) = \Theta(n \log n)$

Summary for sorting lower bound

- A lower bound of $g(n)$ means for any algorithm there exists an input for which the algorithm takes at least $g(n)$ steps on that input
 - If that matches with the upper bound (time complexity), it means this algorithm is **optimal**, and the upper bound is **tight**
- For sorting algorithm, we use the comparison model that assumes an algorithm compares two items with cost 1, and all other operations are free
- We can show that to distinguish $n!$ possible inputs, we need at least $\log_2 n! = \Omega(n \log n)$ comparisons, indicating that quicksort and mergesort are optimal comparison-sort algorithms

Sorting algorithms

- **Comparison-based sorts:**


- Bubble sort: compare adjacent elements
- Selection sort: compare to find the smallest, 2nd smallest, ...
- Quicksort: compare to pivots to partition
- Merge sort: compare in merge
- Bogosort (permutation sort, stupid sort or slowsort, based on infinite monkey theorem)

- **Non-comparison-based sorts:**

- Counting sort: only positive integers in small range
- Bucket sort: have to know key-range (create buckets)
- Radix sort: only integers (sorting based on the bits)
- Sleep sort: hmmm...

Bogosort:

```
while not is_sorted(data) {  
    shuffle(data);  
}  
return data;
```



Sleep sort:


```
printNumber(n) {  
    sleep n seconds;  
    print n; }  
  
// start n threads  
parallel_for (i=1 to n)  
    printNumber(A[i]);  
wait for all threads to finish
```

What sorting algorithm should I use?


- **Merge sort and quicksort both are $O(n \log n)$**
 - In expectation for quicksort
- **Which one is faster?**
 - In practice, quicksort usually shows better performance
 - (that's why it's called quicksort. `std::sort` in STL also uses quicksort)
 - Why? Because quicksort is in-place (no extra space used), while merge requires extra space to hold the merging result temporarily (and then write back)
 - Also is more cache-friendly
- **Time complexity help you roughly predict the performance**
- **Many other practical considerations may affect performance**
- **We can develop more theoretical models, get more experience in coding, ... to better understand the performance**

Pointer machine model

- You cannot random access the memory based on address (i.e., not assume contiguous memory)
- You can access the memory only use pointers



You can access a memory location by a pointer




You **cannot** find the i -th element in an array by using $a[i]$

- Other models: cell-probe model, counter-machine model, etc.

I/O model

- Measures the number of reads and writes of an algorithm



- Examples of lower-bound proofs on the I/O model:
Improved Parallel Cache-Oblivious Algorithms for Dynamic Programming and Linear Algebra

Upper bound vs. lower bound

- An **upper bound** $f(n)$ of the cost of a problem means there exists an algorithm that takes at most $f(n)$ steps on any input of size n
 - So, given this problem, we can just run this algorithm to get an answer
 - $f(n)$ is **guaranteed to be sufficient**: we don't need more than $f(n)$ costs
- A **lower bound** of $g(n)$ means for any algorithm, there exists an input for which all algorithm takes at least $g(n)$ steps on that input
 - Whatever algorithm you use, you cannot get better than $g(n)$!!
- **When upper bound meets lower bound...**
 - An algorithm has cost $f(n)$, and the best you can do is $g(n) = f(n)$
 - That's an **optimal** algorithm!

Summary for lower bounds

- The minimum “cost” to solve a PROBLEM using ANY algorithms
- Lower bounds are for problems, not algorithms
- Commonly used models for analyzing lower bounds:
comparison model, pointer machine model, I/O model, cell-probe model, counter-machine model
- Analyzing or at least knowing the lower bound is helpful
 - Showing the NP-hard / NP-completeness can be considered as a “special” lower bound

Lower bound proof

- Further reading

- <https://courses.cs.vt.edu/~cs4104/shaffer/Spring2007/bounds.pdf>