


Fundamentals of Machine Learning

OPTIMIZATION


Amit K Roy-Chowdhury

Acknowledgments: Adapted from slides at <https://probml.github.io/pml-book/teaching1.html> by Prof. Saw Shier Nee



Outline

- Convex Function
- Gradient Descent
- Newton's Method
- Stochastic Gradient Descent
- Constrained Optimization




Optimization

The core problem in machine learning is parameter estimation (aka model fitting).


This requires solving an optimization problem, where we try to find the values for a set of variables, $\theta \in \Theta$, that minimize a scalar-valued loss function or cost function, $\mathcal{L}(\theta)$.

$$\theta^* \in \operatorname{argmin}_{\theta \in \Theta} \mathcal{L}(\theta)$$



Global / Local Optimization

$$\theta^* \in \operatorname{argmin}_{\theta \in \Theta} \mathcal{L}(\theta)$$





Definitions: Gradient, Hessian, Jacobian

$$\nabla f(p) = \begin{bmatrix} \frac{\partial f}{\partial x_1}(p) \\ \vdots \\ \frac{\partial f}{\partial x_n}(p) \end{bmatrix}, \quad \mathbf{H}_f = \begin{bmatrix} \frac{\partial^2 f}{\partial x_1^2} & \frac{\partial^2 f}{\partial x_1 \partial x_2} & \cdots & \frac{\partial^2 f}{\partial x_1 \partial x_n} \\ \frac{\partial^2 f}{\partial x_2 \partial x_1} & \frac{\partial^2 f}{\partial x_2^2} & \cdots & \frac{\partial^2 f}{\partial x_2 \partial x_n} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial^2 f}{\partial x_n \partial x_1} & \frac{\partial^2 f}{\partial x_n \partial x_2} & \cdots & \frac{\partial^2 f}{\partial x_n^2} \end{bmatrix}.$$

$$\mathbf{H}(f(\mathbf{x})) = \mathbf{J}(\nabla f(\mathbf{x})).$$

$$\mathbf{f} : \mathbf{R}^n \rightarrow \mathbf{R}^m$$

$$\mathbf{J} = \begin{bmatrix} \frac{\partial \mathbf{f}}{\partial x_1} & \cdots & \frac{\partial \mathbf{f}}{\partial x_n} \end{bmatrix} = \begin{bmatrix} \nabla^T f_1 \\ \vdots \\ \nabla^T f_m \end{bmatrix} = \begin{bmatrix} \frac{\partial f_1}{\partial x_1} & \cdots & \frac{\partial f_1}{\partial x_n} \\ \vdots & \ddots & \vdots \\ \frac{\partial f_m}{\partial x_1} & \cdots & \frac{\partial f_m}{\partial x_n} \end{bmatrix}$$




Global / Local minima


Let $g(\theta) = \nabla \mathcal{L}(\theta)$ be the gradient vector,

$H(\theta) = \nabla^2 \mathcal{L}(\theta)$ be the Hessian matrix.


- Necessary condition: If θ^* is a local minimum, then we must have $g^* = 0$ (i.e., θ^* must be a **stationary point**), and H^* must be positive semi-definite.
- Sufficient condition: If $g^* = 0$ and H^* is positive definite, then θ^* is a local optimum.



$x^2 + y^2$
(definite)



x^2
(semidefinite)



$x^2 - y^2$
(indefinite)



Constrained / Unconstrained Optimization

Inequality
constraints

Equality
constraints

We define the **feasible set** as the subset of the parameter space that satisfies the constraints:

$$\mathcal{C} = \{\theta : g_j(\theta) \leq 0 : j \in \mathcal{I}, h_k(\theta) = 0 : k \in \mathcal{E}\} \subseteq \mathbb{R}^D$$


Our **constrained optimization** problem now becomes

$$\theta^* \in \underset{\theta \in \mathcal{C}}{\operatorname{argmin}} \mathcal{L}(\theta)$$


If $\mathcal{C} = \mathbb{R}^D$, it is called **unconstrained optimization**.

If too many constraints, empty feasible sets


A common strategy to solve constrained problem is add penalty to the loss function such as using Lagrangian multiplier.



Convex / Concave Optimization




Convex



Not Convex

We say \mathcal{S} is a **convex set** if, for any $x, x' \in \mathcal{S}$, we have



$$\lambda x + (1 - \lambda)x' \in \mathcal{S}, \quad \forall \lambda \in [0, 1]$$




Convex Function

We say f is a **convex function** if its **epigraph** (the set of points above the function, illustrated in Figure 8.4a) defines a convex set. Equivalently, a function $f(x)$ is called convex if it is defined on a convex set and if, for any $x, y \in \mathcal{S}$, and for any $0 \leq \lambda \leq 1$, we have


$$f(\lambda x + (1 - \lambda)y) \leq \lambda f(x) + (1 - \lambda)f(y) \quad (8.7)$$






Convex Function


$$f(\lambda x + (1 - \lambda)y) \leq \lambda f(x) + (1 - \lambda)f(y)$$



Convex





Concave




Convex Function

Neither convex nor concave




Types of Convex Function


Strictly convex



Strictly concave



Convex but not strictly



Neither convex nor concave - saddle point



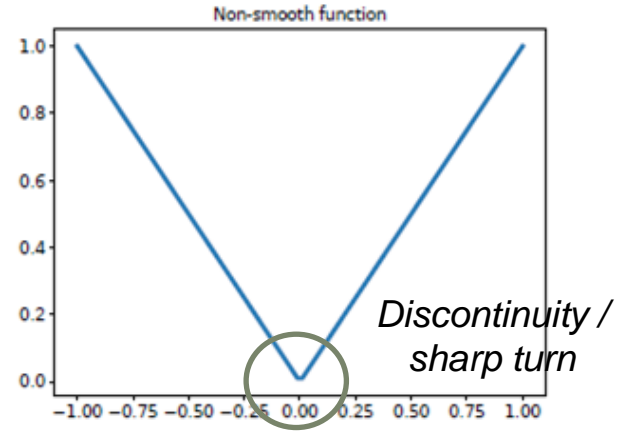


Figure 8.6: The quadratic form $f(x) = x^T A x$ in 2d. (a) A is positive definite, so f is convex. (b) A is negative definite, so f is concave. (c) A is positive semidefinite but singular, so f is convex, but not strictly. Notice the valley of constant height in the middle. (d) A is indefinite, so f is neither convex nor concave. The stationary point in the middle of the surface is a saddle point. From Figure 5 of [She94].

Smooth and Non-Smooth Optimization




(a)



(b)

$$|f(x_1) - f(x_2)| \leq L|x_1 - x_2|$$

Lipschitz constant – quantify the degree of smoothness




First Order Method

- Gradient descent
- Step size/ learning rate
- Convergence rate
- Momentum Method

Taylor series:

$$f(x + \Delta x) = f(x) + f'(x)\Delta x + \frac{1}{2!}f''(x)\Delta x^2 + \frac{1}{3!}f'''(x)\Delta x^3 + \dots$$



Gradient Descent

We need $\mathcal{L}(\theta + \eta d) < \mathcal{L}(\theta)$

Gradient at current iterate: $g_t \triangleq \nabla \mathcal{L}(\theta)|_{\theta_t} = \nabla \mathcal{L}(\theta_t) = g(\theta_t)$

Descent direction: $d^T g_t = \|d\| \|g_t\| \cos(\theta) < 0$

pick $d_t = -g_t$

First Order Method

- Gradient descent
- Step size/ learning rate
- Convergence rate
- Momentum Method

steepest descent will have global convergence iff the step size satisfies

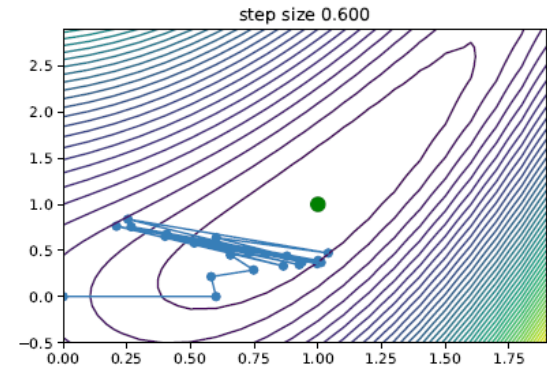

$$\rho < \frac{2}{\lambda_{\max}(\mathbf{A})}$$

$$\mathcal{L}(\theta) = \frac{1}{2}\theta^T \mathbf{A}\theta + b^T \theta + c \text{ with } \mathbf{A} \succeq 0.$$

λ_{\max} = max eigenvalue

$$\theta_{t+1} = \theta_t + \rho_t d_t$$

Updated parameter Step size Descent direction



First Order Method

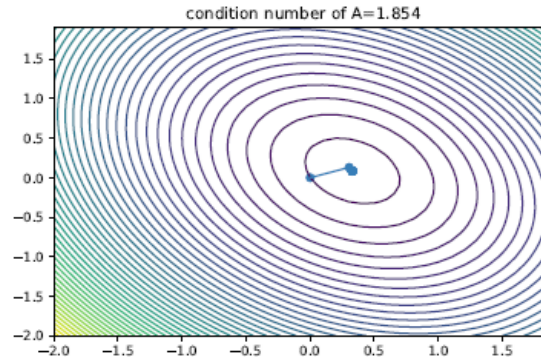

- Gradient descent
- Step size/ learning rate
- Convergence rate, μ
- Momentum Method


$$\mathcal{L}(\theta) = \frac{1}{2}\theta^T A \theta + b^T \theta + c \text{ with } A \succeq 0.$$

$$\mu = \left(\frac{\lambda_{\max} - \lambda_{\min}}{\lambda_{\max} + \lambda_{\min}} \right)^2$$

$$\mu = \left(\frac{\kappa - 1}{\kappa + 1} \right)^2, \text{ where } \kappa = \frac{\lambda_{\max}}{\lambda_{\min}} \text{ is the condition number of } A.$$


The condition number measures how “skewed” the space is, in the sense of being far from a symmetrical “bowl”





First Order Method

- Gradient descent
- Step size/ learning rate
- Convergence rate
- Momentum Method



Thinking like a ball rolling downward. At flat surface, it rolls down slowly. At sharp region, roll down faster.

momentum

$$\begin{aligned} \vec{m}_t &= \beta \vec{m}_{t-1} + \vec{g}_{t-1} \\ \theta_t &= \theta_{t-1} - \rho_t \vec{m}_t \end{aligned}$$

(exponentially weighted moving average of past gradients)

Normally $\beta=0.9$, if $\beta = 0$ - gradient descent




Adaptive Moment Estimation (Adam)

$$m_t = \beta_1 m_{t-1} + (1 - \beta_1) g_t$$
$$s_t = \beta_2 s_{t-1} + (1 - \beta_2) g_t^2$$

$$\beta_1 = 0.9, \beta_2 = 0.999 \text{ and } \epsilon = 10^{-6}$$

$$\eta_t = 0.001$$


$$\Delta \theta_t = -\eta_t \frac{1}{\sqrt{s_t} + \epsilon} m_t$$




First Order Method – Line Search

Consider we want to minimize this loss function, $f(\theta)$.


What if we use line search (iterative method to find θ^*)






First Order Method – Line Search

Start with random number, θ_t







First Order Method – Line Search

Start with random number, θ_t

Compute the gradient at x_k

$$f'(\theta_t) = \frac{f(\theta) - f(\theta_t)}{\theta - \theta_t}$$






First Order Method – Line Search


Start with random number, θ_k

Compute the gradient at θ_k

$$f'(\theta_t) = \frac{f(\theta) - f(\theta_t)}{\theta - \theta_t}$$

If $f'(\theta_t)$ is negative, move θ_t to the right





First Order Method – Line Search


Start with random number, θ_k


Compute the gradient at x_k

$$f'(\theta_t) = \frac{f(\theta) - f(\theta_t)}{\theta - \theta_t}$$

If $f'(\theta_t)$ is negative, move θ_t to the right

If $f'(\theta_t)$ is positive, move θ_t to the left





First Order Method – Line Search

Start with random number, θ

Compute the gradient at x_k

$$f'(\theta_t) = \frac{f(\theta) - f(\theta_t)}{\theta - \theta_t}$$


If $f'(\theta_t)$ is negative, move x_k to the right


If $f'(\theta_t)$ is positive, move x_k to the left

$$\theta_{t+1} = \theta_t - \alpha f'(\theta_t)$$

Large step size, α will overshoot

Small step size, α will be very slow





Second Order Method – Newton Method

Approximate with non linear graph

Compute the second derivative at θ_k


$$f''(\theta_k) = \frac{f'(\theta) - f'(\theta_k)}{\theta - \theta_k}$$

$$\theta_{t+1} = \theta_k - \alpha f'(\theta_k)$$

For Newton's method, the update formula is

$$\theta_{t+1} = \theta_k - \alpha \frac{1}{f''(\theta_k)} f'(\theta_k)$$

- Faster convergence





Second Order Method – Newton Method

Higher dimension

$$\theta_{t+1} = \theta_t - \rho_t \mathbf{H}_t^{-1} g_t$$

where

$$\mathbf{H}_t \triangleq \nabla^2 \mathcal{L}(\theta)|_{\theta_t} = \nabla^2 \mathcal{L}(\theta_t) = \mathbf{H}(\theta_t)$$

\mathbf{H} = Hessian matrix
 ρ = step size
 g_t = gradient


$$\mathbf{H}_f = \begin{bmatrix} \frac{\partial^2 f}{\partial x_1^2} & \frac{\partial^2 f}{\partial x_1 \partial x_2} & \dots & \frac{\partial^2 f}{\partial x_1 \partial x_n} \\ \frac{\partial^2 f}{\partial x_2 \partial x_1} & \frac{\partial^2 f}{\partial x_2^2} & \dots & \frac{\partial^2 f}{\partial x_2 \partial x_n} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial^2 f}{\partial x_n \partial x_1} & \frac{\partial^2 f}{\partial x_n \partial x_2} & \dots & \frac{\partial^2 f}{\partial x_n^2} \end{bmatrix},$$

Consider a quadratic approximation

$$\mathcal{L}_{\text{quad}}(\theta) = \mathcal{L}(\theta_t) + g_t^\top (\theta - \theta_t) + \frac{1}{2} (\theta - \theta_t)^\top \mathbf{H}_t (\theta - \theta_t)$$

The minimum of $\mathcal{L}_{\text{quad}}$ is at

$$\theta = \theta_t - \mathbf{H}_t^{-1} g_t$$



Gradient Descent vs Newton's method

We need $\mathcal{L}(\theta + \eta d) < \mathcal{L}(\theta)$

Gradient at current iterate: $g_t \triangleq \nabla \mathcal{L}(\theta)|_{\theta_t} = \nabla \mathcal{L}(\theta_t) = g(\theta_t)$

Gradient Descent: $d^T g_t = \|d\| \|g_t\| \cos(\theta) < 0$
pick $d_t = -g_t$

Consider only first two terms of Taylor series

Newton's method: $d_t = -H_t^{-1} g_t$

Consider only first three term of Taylor series



Stochastic Gradient Descent

Stochastic Optimization: $\mathcal{L}(\theta) = \mathbb{E}_{q(z)} [\mathcal{L}(\theta, z)]$

$$\theta_{t+1} = \theta_t - \eta_t \nabla \mathcal{L}(\theta_t, z_t) = \theta_t - \eta_t g_t$$

(distribution of random variable is independent of parameter we are optimizing over)

Consider loss function:

$$\mathcal{L}(\theta_t) = \frac{1}{N} \sum_{n=1}^N \ell(y_n, f(x_n; \theta_t)) = \frac{1}{N} \sum_{n=1}^N \mathcal{L}_n(\theta_t)$$


$$g_t = \frac{1}{N} \sum_{n=1}^N \nabla_{\theta} \mathcal{L}_n(\theta_t) = \frac{1}{N} \sum_{n=1}^N \nabla_{\theta} \ell(y_n, f(x_n; \theta_t))$$

This requires summing over all N training examples, and thus can be slow if N is large.

Minibatch:
$$g_t \approx \frac{1}{|\mathcal{B}_t|} \sum_{n \in \mathcal{B}_t} \nabla_{\theta} \mathcal{L}_n(\theta_t) = \frac{1}{|\mathcal{B}_t|} \sum_{n \in \mathcal{B}_t} \nabla_{\theta} \ell(y_n, f(x_n; \theta_t))$$

Minibatch sampling, training epochs

where \mathcal{B}_t is a set of randomly chosen examples to use at iteration t .



Constrained Optimization

$$\theta^* = \arg \min_{\theta \in \mathcal{C}} \mathcal{L}(\theta)$$


where the feasible set, or constraint set, is

$$\mathcal{C} = \{\theta \in \mathbb{R}^D : h_i(\theta) = 0, i \in \mathcal{E}, g_j(\theta) \leq 0, j \in \mathcal{I}\}$$

Taylor series:
$$h(\theta + \epsilon) \approx h(\theta) + \epsilon^T \nabla h(\theta)$$


Since both θ and $\theta + \epsilon$ are on the constraint surface, we must have $h(\theta) = h(\theta + \epsilon)$ and hence $\epsilon^T \nabla h(\theta) \approx 0$. Since ϵ is parallel to the constraint surface, $\nabla h(\theta)$ must be perpendicular to it.

$\nabla \mathcal{L}(\theta)$ is also orthogonal to the constraint surface $\nabla \mathcal{L}(\theta^*) = \lambda^* \nabla h(\theta^*)$



Lagrangian:
$$L(\theta, \lambda) \triangleq \mathcal{L}(\theta) + \lambda h(\theta)$$


$$\nabla_{\theta, \lambda} L(\theta, \lambda) = \mathbf{0} \iff \lambda \nabla_{\theta} h(\theta) = \nabla \mathcal{L}(\theta), h(\theta) = 0$$



Constrained Optimization

Lagrangian: $L(\theta, \lambda) \triangleq \mathcal{L}(\theta) + \lambda h(\theta)$

$$\nabla_{\theta, \lambda} L(\theta, \lambda) = \mathbf{0} \iff \lambda \nabla_{\theta} h(\theta) = \nabla \mathcal{L}(\theta), h(\theta) = 0$$



$$L(\theta_1, \theta_2, \lambda) = \theta_1^2 + \theta_2^2 + \lambda(\theta_1 + \theta_2 - 1)$$

$$\frac{\partial}{\partial \theta_1} L(\theta_1, \theta_2, \lambda) = 2\theta_1 + \lambda = 0$$

$$2\theta_1 = -\lambda = 2\theta_2, \text{ so } \theta_1 = \theta_2.$$

$$\frac{\partial}{\partial \theta_2} L(\theta_1, \theta_2, \lambda) = 2\theta_2 + \lambda = 0$$

$$2\theta_1 = 1$$

$$\frac{\partial}{\partial \lambda} L(\theta_1, \theta_2, \lambda) = \theta_1 + \theta_2 - 1 = 0$$

$$\theta^* = (0.5, 0.5)$$