

CUDA Parallelism Model

UNIVERSITY OF CALIFORNIA, RIVERSIDE

Example: Vector Addition Kernel (<u>Device Code</u>)



```
// Compute vector sum C = A + B
// Each thread performs one pair-wise addition
  global
void vecAddKernel(float* A, float* B, float* C, int n)
  int i = threadIdx.x + blockDim.x * blockIdx.x;
  if(i<n)
     C[i] = A[i] + B[i];
```

Example: Vector Addition Kernel Launch (Host Code)



```
void vecAdd(float* h_A, float* h_B, float* h_C, int n)
{
   // d_A, d_B, d_C allocations and copies omitted
   // Run ceil(n/256.0) blocks of 256 threads each
   vecAddKernel<<<ceil(n/256.0),256>>> (d_A, d_B, d_C, n);
}
```

More on Kernel Launch (Host Code)



```
void vecAdd(float* h_A, float* h_B, float* h_C, int n)
{
    dim3 DimGrid((n-1)/256 + 1, 1, 1);
    dim3 DimBlock(256, 1, 1);
    vecAddKernel<<<DimGrid,DimBlock>>>(d_A, d_B, d_C, n);
}
```

Kernel execution in a nutshell



```
global
  host
                                              void vecAddKernel(float *A,
void vecAdd(...)
                                                   float *B, float *C, int n)
  dim3 DimGrid(ceil(n/256.0),1,1);
                                                 int i = blockIdx.x * blockDim.x
  dim3 DimBlock (256,1,1);
                                                            + threadIdx.x;
  vecAddKernel<<<DimGrid,DimBlock>>>
   (d A, d B, d C, n);
                                                 if(i < n) C[i] = A[i] + B[i];
                                      Grid
                                       GPU
                                              SMk
                                 SM0
                                       RAM
```

More on CUDA Function Declarations



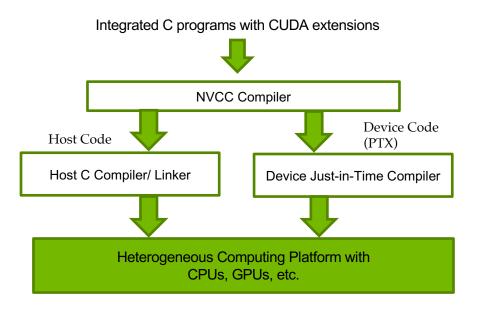
	Executed on the:	Only callable from the:
device float DeviceFunc()	device	device
global void KernelFunc()	device	host
host float HostFunc()	host	host

Device = GPU Host = CPU

- __global___ defines a kernel function
 - Each "__" consists of two underscore characters
 - A kernel function must return void
- __device__ and __host__ can be used together
- host is optional if used alone

Compiling A CUDA Program



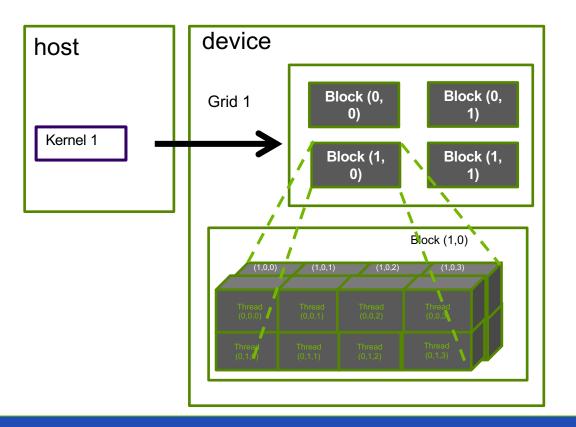




MULTI-DIMENSIONAL KERNEL CONFIGURATION

A Multi-Dimensional Grid Example



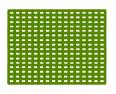


2D Grid

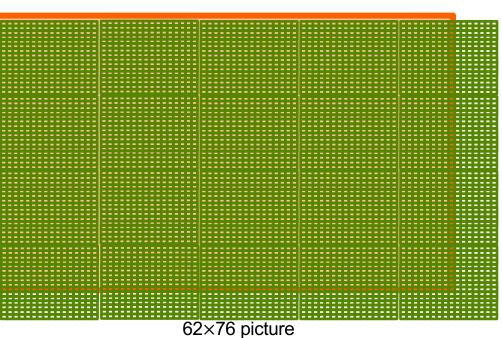
3D Blocks

Processing a Picture with a 2D Grid

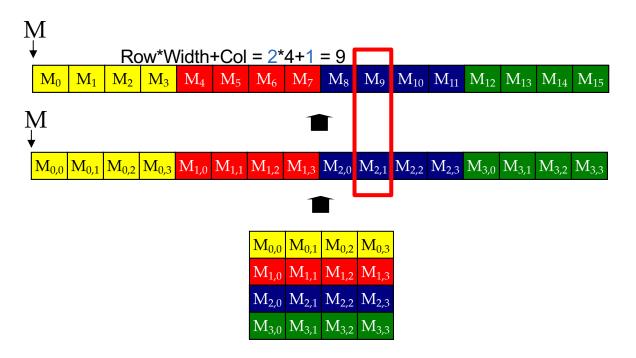




16×16 blocks



Review: Row-Major Memory Layout in C/C++ UCR



Example Kernel Code of a PictureKernel



```
global void PictureKernel(float* d Pin, float* d Pout,
                        int height, int width)
// Calculate the row # of the d Pin and d Pout element
int Row = blockIdx.y*blockDim.y + threadIdx.y;
// Calculate the column # of the d Pin and d Pout element
int Col = blockIdx.x*blockDim.x + threadIdx.x;
// each thread computes one element of d Pout if in range
if ((Row < height) && (Col < width)) {</pre>
  d Pout[Row*width+Col] = 2.0*d Pin[Row*width+Col];
```

Scale every pixel value by 2.0

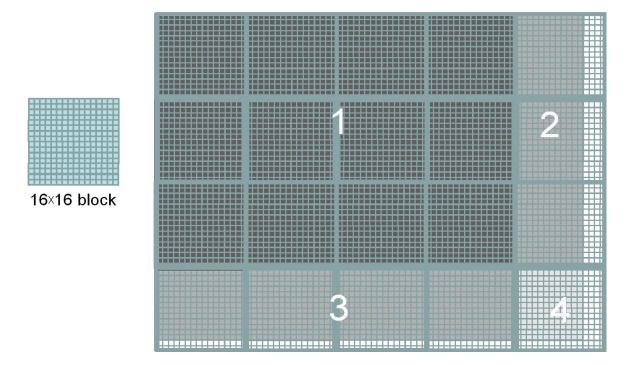
Host Code for Launching PictureKernel



```
// assume that the picture is m n,
// m pixels in y dimension and n pixels in x dimension
// input d_Pin has been allocated on and copied to device
// output d_Pout has been allocated on device
...
dim3 DimGrid((n-1)/16 + 1, (m-1)/16+1, 1);
dim3 DimBlock(16, 16, 1);
PictureKernel<<<DimGrid, DimBlock>>>(d_Pin, d_Pout, m, n);
...
```

Covering a 62×76 Picture with 16×16 Blocks





Note: Not all threads in a Block will follow the same control flow path.

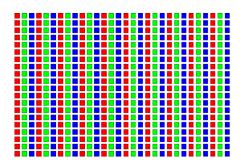


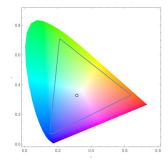
COLOR-TO-GRAYSCALE IMAGE PROCESSING EXAMPLE

RGB Color Image Representation



- Each pixel in an image is an RGB value
- The format of an image's row is (r g b) (r g b) ... (r g b)
- RGB ranges are not distributed uniformly
- Many different color spaces, here we show the constants to convert to AdbobeRGB color space
 - The vertical axis (y value) and horizontal axis (x value) show the fraction of the pixel intensity that should be allocated to G and B. The remaining fraction (1-y-x) of the pixel intensity that should be assigned to R
 - The triangle contains all the representable colors in this color space





RGB to Grayscale Conversion



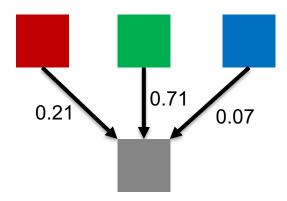


A grayscale digital image is an image in which the value of each pixel carries only intensity information.

Color Calculating Formula



- For each pixel (r g b) at (I, J) do: grayPixel[I,J] = 0.21*r + 0.71*g + 0.07*b
- This is just a dot product <[r,g,b],[0.21,0.71,0.07]> with the constants being specific to input RGB space



RGB to Grayscale Conversion Code



```
#define CHANNELS 3 // we have 3 channels corresponding to RGB
// The input image is encoded as unsigned characters [0, 255]
__alobal__ void colorConvert(unsigned char * grayImage,
                                     unsigned char * rgbImage,
                             int width, int height) {
 int x = threadIdx.x + blockIdx.x * blockDim.x;
 int y = threadIdx.y + blockIdx.y * blockDim.y;
 if (x < width && y < height) {
    // get 1D coordinate for the grayscale image
    int gray0ffset = y*width + x;
    // one can think of the RGB image having
    // CHANNEL times columns than the gray scale image
    int rgbOffset = grayOffset*CHANNELS;
    unsigned char r = rgbImage[rgbOffset ]; // red value for pixel
    unsigned char g = rgbImage[rgb0ffset + 1]; // green value for pixel
    unsigned char b = rabImage[rabOffset + 2]: // blue value for pixel
    // perform the rescaling and store it
    // We multiply by floating point constants
    qrayImage[qrayOffset] = 0.21f*r + 0.71f*q + 0.07f*b;
```

Every thread mapped to output pixel



THREAD BLOCK SCHEDULING

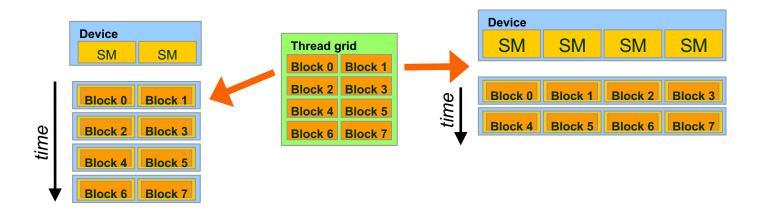
Objective



- To learn how a CUDA kernel utilizes hardware execution resources
 - Assigning thread blocks to execution resources
 - Capacity constrains of execution resources (hardware)
 - "Zero-overhead" thread scheduling

Transparent Scalability



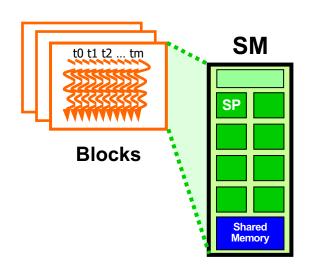


- Each block can execute in any order relative to others.
- Hardware is free to assign blocks to any processor at any time
 - A kernel scales to any number of parallel processors

Example: Executing Thread Blocks



- Threads are assigned to Streaming
 Multiprocessors (SM) in block granularity
 - Example with Fermi architecture:
 - Up to 8 blocks to each SM as resource allows
 - Fermi SM can take up to 1536 threads
 - Could be 256 (threads/block) * 6 blocks
 - Or 512 (threads/block) * 3 blocks, etc.
- SM maintains thread/block idx #s
- SM manages/schedules thread execution





- For 2D algorithms using multiple blocks, should I use 8X8, 16X16 or 32X32 blocks for Fermi?
 - Limitations: 1536 threads/SM & 8 blocks / SM

Block size	# Threads / Block	# Blocks <i>that</i> can fit in SM's threads	Max blocks / SM	# Blocks in SM	# of threads utilized in SM
8x8					/ 1536
16x16					/ 1536
32x32					/ 1536



- For 2D algorithms using multiple blocks, should I use 8X8, 16X16 or 32X32 blocks for Fermi?
 - Limitations: 1536 threads/SM & 8 blocks / SM

Block size	# Threads / Block	# Blocks that can fit in SM's threads	Max blocks / SM	# Blocks in SM	# of threads utilized in SM
8x8	64	24	8	8	512 / 1536
16x16					/ 1536
32x32					/ 1536

Limited by number of blocks per SM



- For 2D algorithms using multiple blocks, should I use 8X8, 16X16 or 32X32 blocks for Fermi?
 - Limitations: 1536 threads/SM & 8 blocks / SM

Block size	# Threads / Block	# Blocks that can fit in SM's threads	Max blocks / SM	# Blocks in SM	# of threads utilized in SM
8x8	64	24	8	8	512 / 1536
16x16	256	6	8	6	1536 / 1536
32x32					/ 1536

Fully utilize all hardware threads!



- For 2D algorithms using multiple blocks, should I use 8X8, 16X16 or 32X32 blocks for Fermi?
 - Limitations: 1536 threads/SM & 8 blocks / SM

Block size	# Threads / Block	# Blocks that can fit in SM's threads	Max blocks / SM	# Blocks in SM	# of threads utilized in SM
8x8	64	24	8	8	512 / 1536
16x16	256	6	8	6	1536 / 1536
32x32	1024	1	8	1	1024 / 1536

Limited by number of threads per SM

What are your GPU's hardware thread

UCR

limitations?

```
~ - ssh dwong@bender.engr.ucr.edu
bender /home/eeprof/dwong $ nvidia-smi -a
 : Tue Oct 4 23:30:33 2022
Timestamp
                                        : 515.48.07
Driver Version
CUDA Version
                                        : 11.7
Attached GPUs
GPU 00000000:3B:00.0
                                        NVIDIA GeForce RTX 2070 SUPER
    Product Name
   Product Brand
                                        : Geroree
   Product Architecture
                                         : Turina
   Display Mode
                                         : Disabled
   Display Active
                                        : Disabled
                                        : Disabled
   Persistence Mode
   MIG Mode
                                        : N/A
       Current
       Pending
                                        : N/A
   Accounting Mode
                                        : Disabled
                                        : 4000
   Accounting Mode Buffer Size
   Driver Model
       Current
                                        : N/A
       Pendina
                                         : N/A
   Serial Number
   GPU UUID
                                         : GPU-7b7fd507-003c-ebf8-0941-1cfd
```

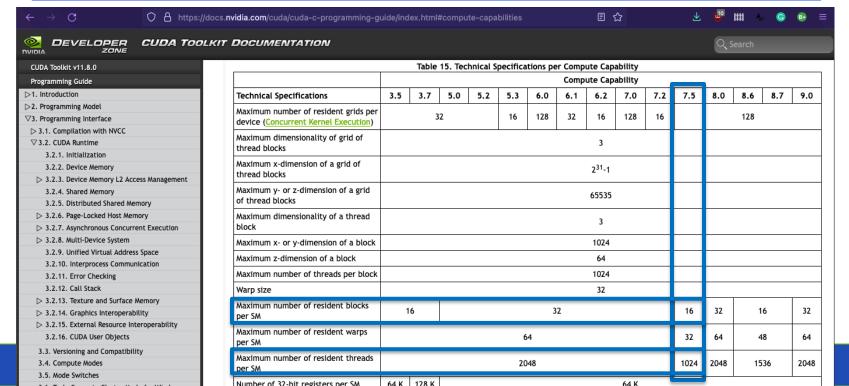
C O A ht	tps://developer. nvidia.com /cuda-gpus
Geforce RTX 2080	7.5
Geforce RTX 2070	7.5
Geforce RTX 2060	7.5
NVIDIA TITAN V	7.0
NVIDIA TITAN Xp	6.1
NVIDIA TITAN X	6.1
GeForce GTX 1080 Ti	6.1
GeForce GTX 1080	6.1
GeForce GTX 1070 Ti	6.1
GeForce GTX 1070	6.1
GeForce GTX 1060	6.1

- Each Nvidia GPU has a "Compute capability"
 - Our RTX 2070 has a compute capability of 7.5

What are your GPU's hardware thread limitations?



https://docs.nvidia.com/cuda/cuda-c-programming-guide/index.html#compute-capabilities

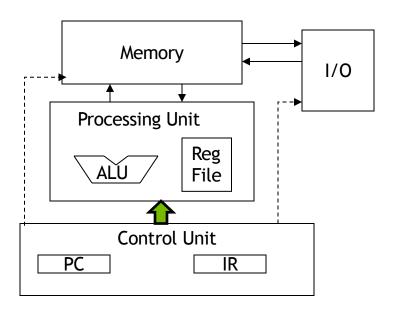




THREAD BLOCK SCHEDULING - WARPS

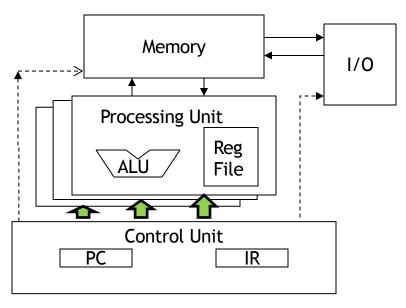
The Von-Neumann Model





The Von-Neumann Model with SIMD units



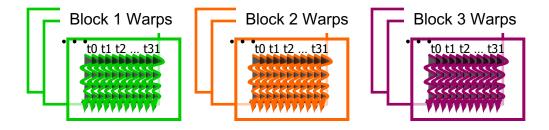


Single Instruction Multiple Data (SIMD)

Warps as Scheduling Units



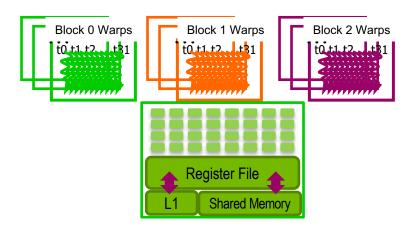
- Each Block is executed as 32-thread Warps
 - An implementation decision, not part of the CUDA programming model
 - Warps are scheduling units within SMs
 - Threads in a warp execute in SIMD (More details later)
 - Future GPUs may have different number of threads in each warp



Warp Example



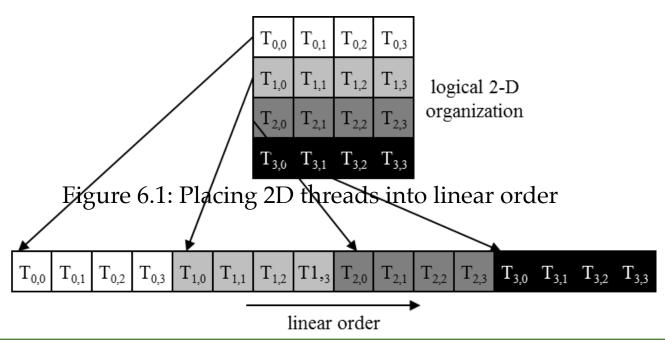
- If 3 blocks are assigned to an SM and each block has 256 threads, how many Warps are there in an SM?
 - Each Block is divided into 256/32 = 8 Warps
 - There are 8 * 3 = 24 Warps



Warps in Multi-dimensional Thread Blocks



- The thread blocks are first linearized into 1D in row major order
 - In x-dimension first, y-dimension next, and z-dimension last



Blocks are partitioned after linearization



- Linearized thread blocks are partitioned
 - Thread indices within a warp are consecutive and increasing
 - Warp 0 starts with Thread 0, Warp 1 starts with Thread 32, etc.
- Partitioning scheme is consistent across devices
 - Thus you can use this knowledge in control flow
 - However, the exact size of warps may change from generation to generation
- DO NOT rely on any ordering within or between warps
 - If there are any dependencies between threads, you must __syncthreads() to get correct results (more later).



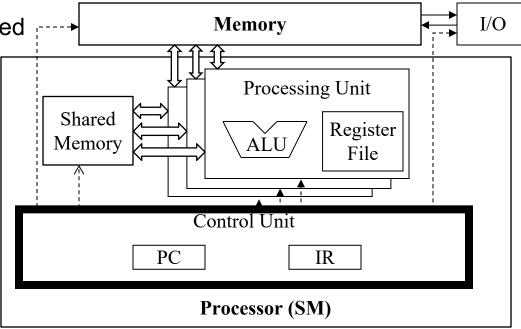
WARP DIVERGENCE

SMs are SIMD Processors



 Control unit for instruction fetch, decode, and control is shared among multiple processing units

Control overhead is minimized



SIMD Execution Among Threads in a Warp



- All threads in a warp must execute the same instruction at any point in time
 - Also called SIMT execution model
- This works efficiently if all threads follow the same control flow path
 - All if-then-else statements make the same decision
 - All loops iterate the same number of times

Control Divergence



- Control divergence occurs when threads in a warp take different control flow paths by making different control decisions
 - Some take the then-path and others take the else-path of an if-statement
 - Some threads take different number of loop iterations than others

- The execution of threads taking different paths are serialized in current GPUs
 - The control paths taken by the threads in a warp are traversed one at a time until there is no more.
 - During the execution of each path, all threads taking that path will be executed in parallel
 - The number of different paths can be large when considering nested control flow statements

Control Divergence Examples



- Divergence can arise when branch or loop condition is a function of thread indices
- Example kernel statement with divergence:
 - if (threadIdx.x > 2) { }
 - This creates two different control paths for threads in a block
 - Decision granularity < warp size; threads 0, 1 and 2 follow different path than the rest of the threads in the first warp
- Example without divergence:
 - If (blockIdx.x > 2) { }
 - Decision granularity is a multiple of blocks size; all threads in any given warp follow the same path





Device Code

```
// Compute vector sum C = A + B
// Each thread performs one pair-wise addition

__global__
void vecAddKernel(float* A, float* B, float* C, int n)
{
   int i = threadIdx.x + blockDim.x * blockIdx.x;
   if(i<n) C[i] = A[i] + B[i];
}</pre>
```