

UCR



Memory Performance and Memory Coalescing

UNIVERSITY OF CALIFORNIA, RIVERSIDE

MEMORY ACCESS PERFORMANCE

Objective

- To learn that memory bandwidth is a first-order performance factor in a massively parallel processor
 - DRAM bursts, banks, and channels
 - All concepts are also applicable to modern multicore processors

Global Memory (DRAM) Bandwidth

Ideal

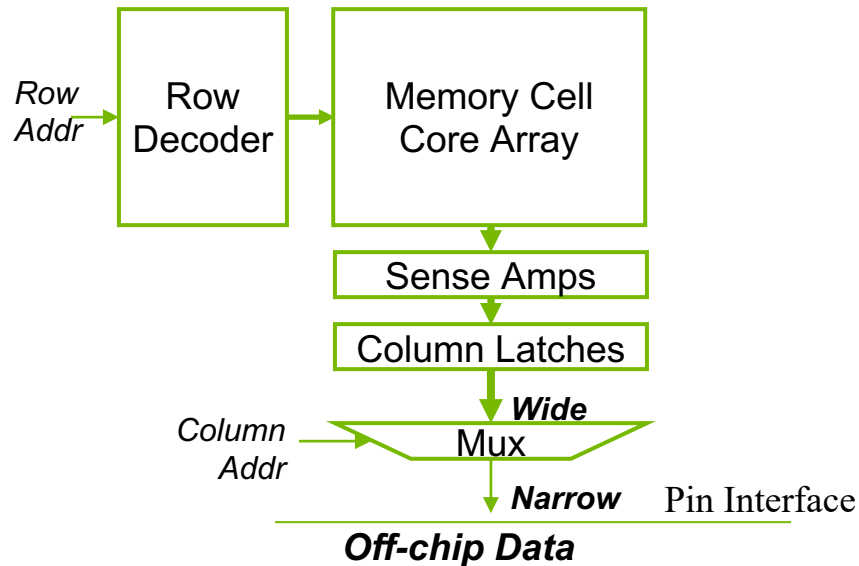


Reality

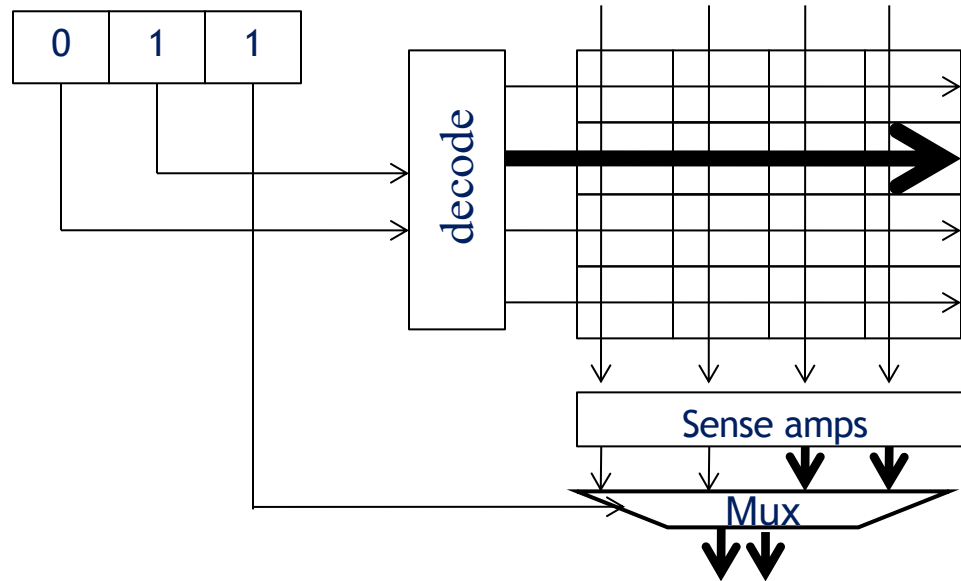


DRAM Core Array Organization

- Each DRAM core array has about 16M bits
- Each bit is stored in a tiny capacitor made of one transistor




A very small (8x2-bit) DRAM Core Array



DRAM Core Arrays are Slow

- Reading from a cell in the core array is a very slow process
 - DDR: Core speed = $\frac{1}{2}$ interface speed
 - DDR2/GDDR3: Core speed = $\frac{1}{4}$ interface speed
 - DDR3/GDDR4: Core speed = $\frac{1}{8}$ interface speed
 - ... likely to be worse in the future



DRAM Bursting

- For DDR{2,3} SDRAM cores clocked at $1/N$ speed of the interface:
- Load ($N \times$ interface width) of DRAM bits from the same row at once to an internal buffer, then transfer in N steps at interface speed
 - DDR3/GDDR4: buffer width = $8X$ interface width


DRAM Bank Operation

Access Address:
(Row 0, Column 0)
(Row 0, Column 1)
(Row 0, Column 85)
(Row 1, Column 0)

Row address 0

Row decoder

Columns



Rows

Row 1


Row Buffer ~~CONFLICT !~~

Column address 85

Column mux


Data

DRAM Bursting Timing Example



Modern DRAM systems are designed to always be accessed in burst mode. Burst bytes are transferred to the processor but discarded when accesses are not to sequential locations.

Multiple DRAM Banks



DRAM Bursting with Banking



Single-Bank burst timing, dead time on interface




Multi-Bank burst timing, reduced dead time

GPU off-chip memory subsystem

- NVIDIA GTX280 GPU:
 - Peak global memory bandwidth = 141.7GB/s
- Global memory (GDDR3) interface @ 1.1GHz
 - (Core speed @ 276Mhz)
 - For a typical 64-bit interface, we can sustain only about 17.6 GB/s (Recall DDR - 2 transfers per clock)
 - We need a lot more bandwidth (141.7 GB/s) – thus 8 memory channels
- NVIDIA Volta:
 - 16GB HBM2 memory subsystem delivers 900 GB/sec peak memory bandwidth
- NVIDIA Ampere A100:
 - 40GB HBM2 memory subsystem delivers 1555 GB/sec peak memory bandwidth


MEMORY COALESCING

DRAM Burst – A System View




- Each address space is partitioned into burst sections
 - Whenever a location is accessed, all other locations in the same section are also delivered to the processor
- Basic example: a 16-byte address space, 4-byte burst sections
 - In practice, we have at least 4GB address space, burst section sizes of 128-bytes or more

Memory Coalescing



- When *all threads of a warp* execute a load instruction, if all accessed locations fall into the *same burst section*, *only one DRAM request* will be made and the access is fully coalesced.

Un-coalesced Accesses



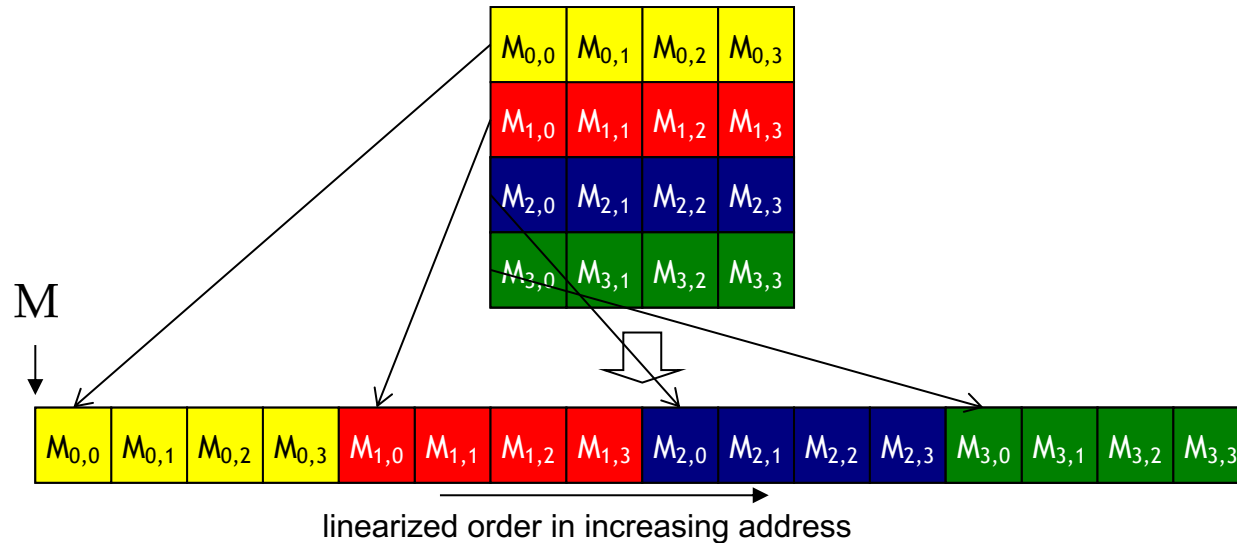
- When the accessed locations spread across burst section boundaries:
 - Coalescing fails
 - Multiple DRAM requests are made
 - The access is not fully coalesced.

- Some of the bytes accessed and transferred are not used by the threads

How to judge if an access is coalesced?

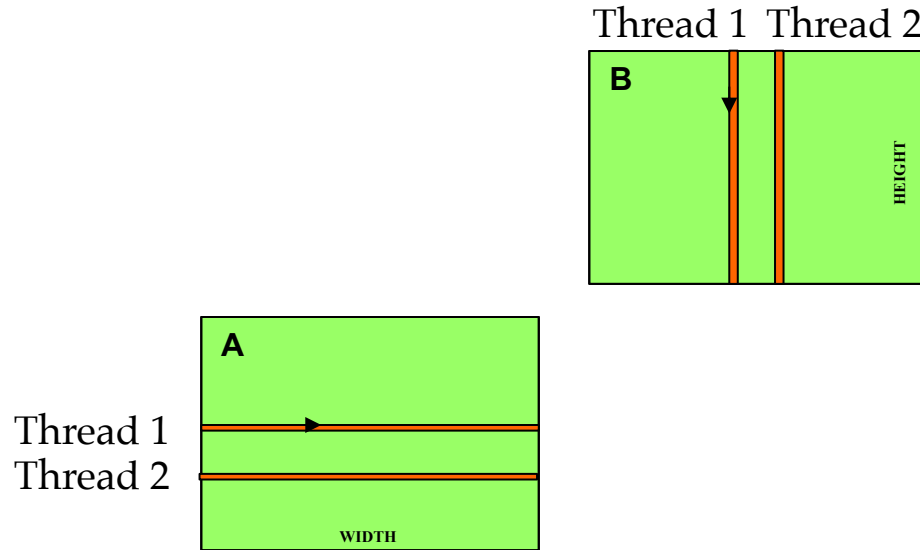
- Accesses **in a warp** are to **consecutive locations** if the index in an array access is in the form of
 - $A[(\text{expression with terms independent of threadIdx.x}) + \text{threadIdx.x}];$

A 2D C Array in Linear Memory Space

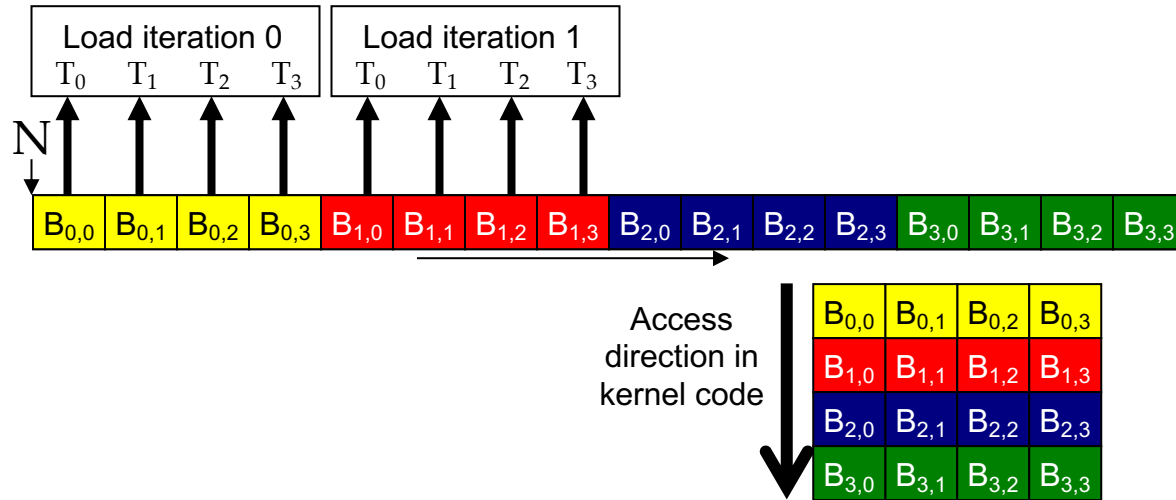


Two Access Patterns of Basic Matrix Multiplication

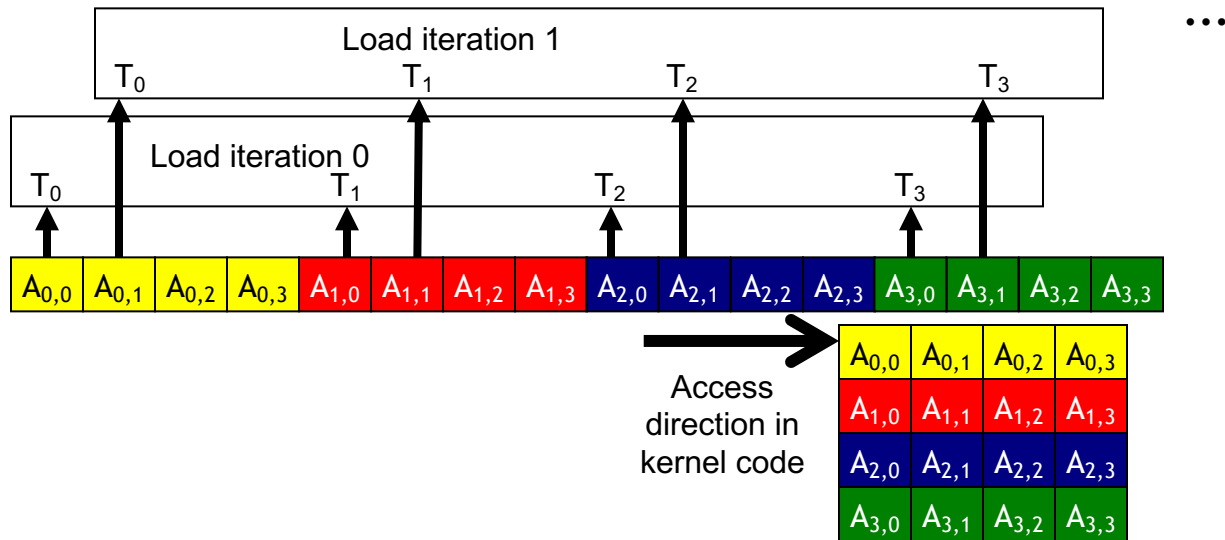
Let's assume the following thread mapping
(not representative of our actual implementation)



B accesses are coalesced



A Accesses are Not Coalesced



Tiling also optimize for coalescing

Have each thread load an A element and a B element at the same relative position as its C element.

```
int tx = threadIdx.x
int ty = threadIdx.y
Accessing tile 0 2D indexing:
    A[Row][tx]
    B[ty][Col]
```

