

Fibonacci Heaps



Priority Queue Operations


- Maintain a set Q (initially empty)
- *Insert*(x, k): add x to Q with value $d(x) = k$
 - Slight overkill for Dijkstra
 - Start by inserting s with value 0, all other nodes with value ∞
- *Deletemin*(): Find an $x \in Q$ minimizing $d(x)$. Delete x from Q and return it.
- *Decreasekey*(x, k): $d(x) \leftarrow \min\{d(x), k\}$

Priority Queue Implementations

- Next: **Binomial Trees**
 - Insert = Decreasekey = Deletemin = $O(\log n)$ time
- After that: **Fibonacci Trees**
 - Insert = Decreasekey = $O(1)$, Deletemin = $O(\log n)$


Heap-ordered trees

- Represent Q by a **forest** of rooted trees
- One element per node
- *Heap ordered*: $d(x) \geq d(\text{parent}(x))$
 - \Rightarrow element with **min. key at a root**
 - \Rightarrow two heap ordered trees can be **linked in $O(1)$ time**




Heap-ordered trees

- Represent Q by a **forest** of rooted trees
- One element per node
- *Heap ordered*: $d(x) \geq d(\text{parent}(x))$
 - \Rightarrow element with **min. key at a root**
 - \Rightarrow two heap ordered trees can be **linked in $O(1)$ time**




Binomial Trees

- A **binomial tree** is a tree that obeys a special structure
- A binomial tree B_i of “rank 0” must be a **single node**.
- A binomial tree B_{i+1} of “rank $i + 1$ ” is:
 - A B_i , with another B_i added as the rightmost child of the root.



Binomial Trees

- Equivalent definition:
 - B_{i+1} = one root node with children B_0, B_1, \dots, B_i



- Size of binomial trees: $|B_0| = 1, |B_i| = 2|B_{i-1}| \Rightarrow |B_i| = 2^i$

Binomial Trees


- B_i = binomial tree with rank i
 - B_0 = one node
 - B_{i+1} = make one B_i the rightmost child of another B_i



Rank 1

Binomial Trees


- B_i = binomial tree with rank i
 - B_0 = one node
 - B_{i+1} = make one B_i the rightmost child of another B_i



Rank 2

Binomial Trees


- B_i = binomial tree with rank i
 - B_0 = one node
 - B_{i+1} = make one B_i the rightmost child of another B_i



Rank 3

Binomial Trees


- B_i = binomial tree with rank i
 - B_0 = one node
 - B_{i+1} = make one B_i the rightmost child of another B_i



Rank 4


Binomial Trees

- B_i = binomial tree with rank i
 - B_0 = one node
 - B_{i+1} = make one B_i the rightmost child of another B_i



Binomial Trees

- B_i = binomial tree with rank i
 - B_0 = one node
 - B_{i+1} = make one B_i the rightmost child of another B_i




Binomial *Heaps*

- Structure of a binomial heap:
 - (1) A list of heap-ordered binomial trees
 - (2) At most one tree of each rank
- *Insert*(X, k):
 - Create a new node X (a B_0 tree)
 - Add X to the list of trees (1)
 - **Now there might be two rank-0 trees**
 - Do some work to restore property (2)

Insertion in a Binomial Heap

Insert(x, k): create new B_0 tree (one node) for x


Now more than one B_0 tree



Insertion in a Binomial Heap

Insert(x, k):


Link B_0 trees. Now there's too many B_1 trees.



Insertion in a Binomial Heap

Insert(x, k):


Link B_1 trees. Now too many B_2 trees.



Insertion in a Binomial Heap

Insert(x, k):


Link B_2 trees.



Insertion in a Binomial Heap

Insert(x, k):

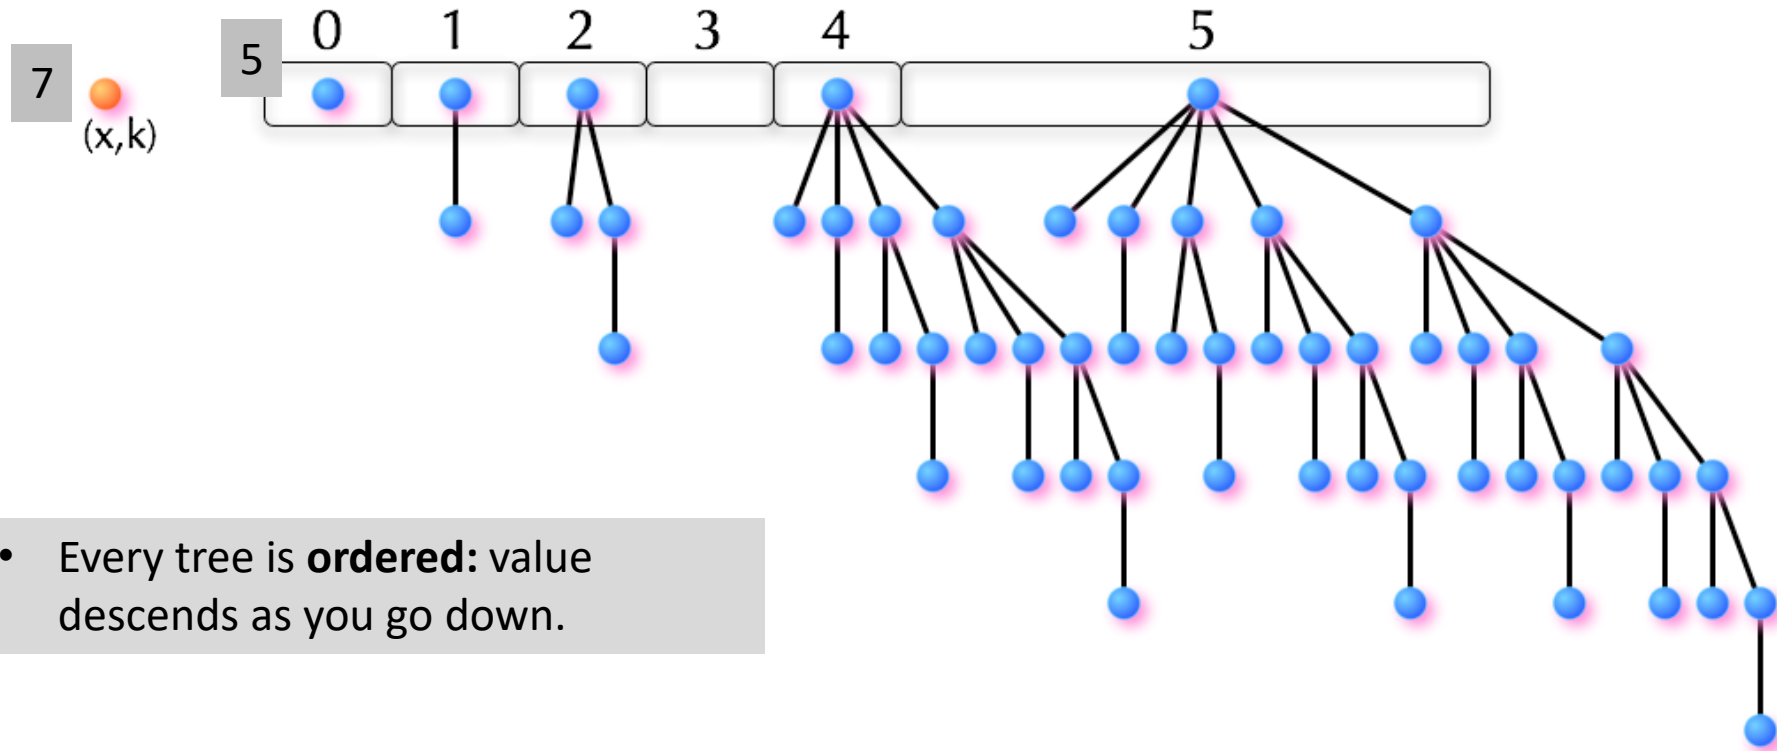
Link B_2 trees. At most 1 tree of each rank so we're done.



Insertion: Which Goes Under?

Insert(x, k): create new B_0 tree (one node) for x

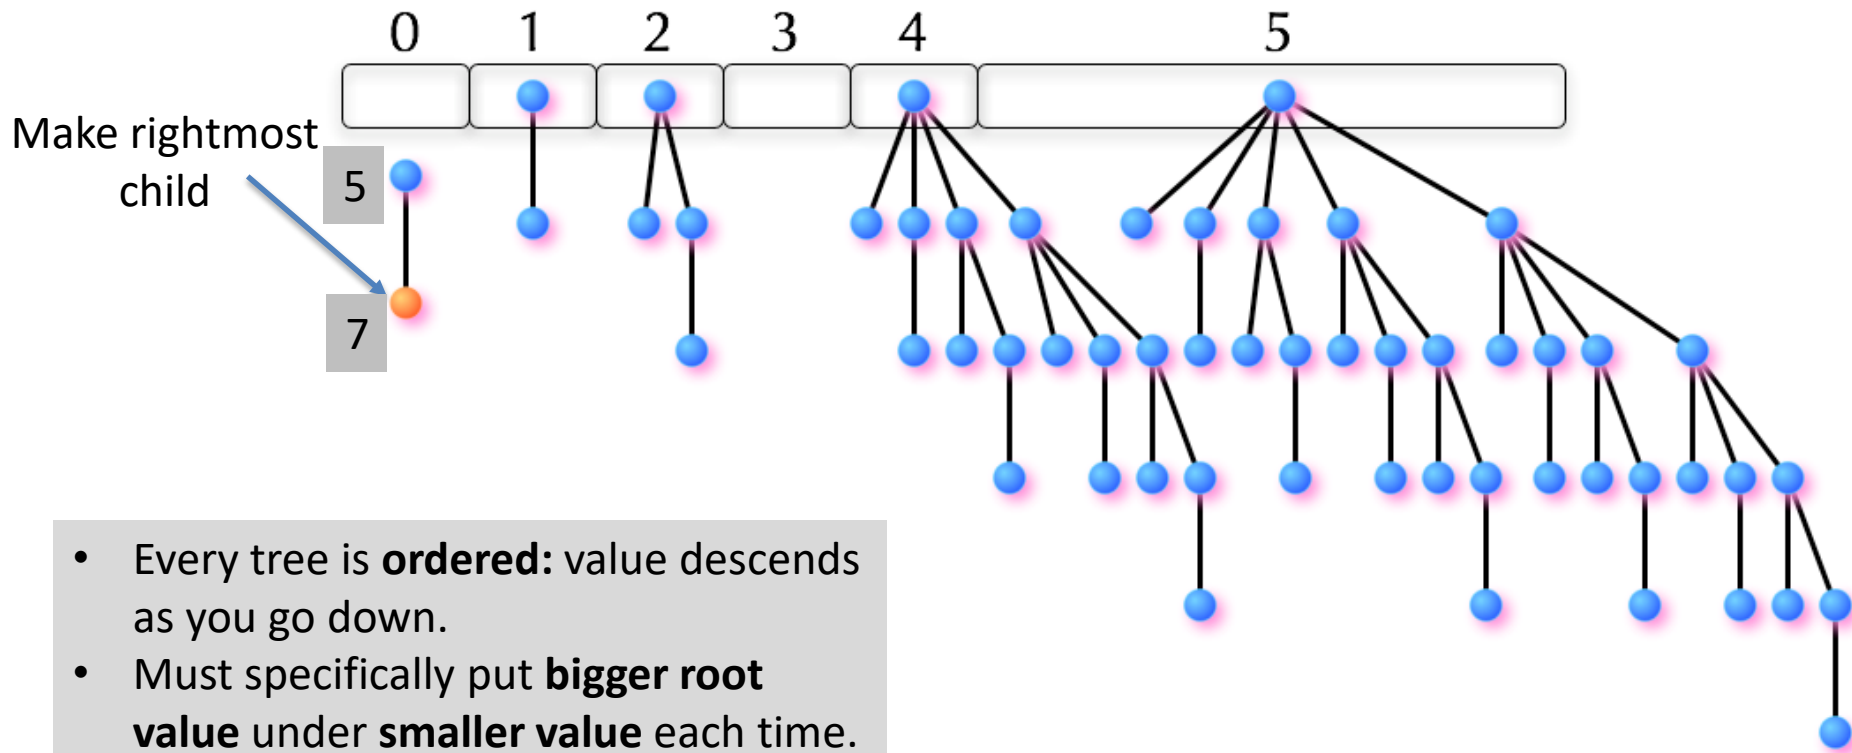
Now more than one B_0 tree



Insertion: Which Goes Under?

Insert(x, k):

Link B_0 trees. Now there's too many B_1 trees.



Binomial *Heaps*

- Structure of a binomial heap:
 - (1) A list of heap-ordered binomial trees
 - (2) At most one tree of each rank
- Insert(X, k) :
 - Create a new node X (a B_0 tree)
 - Repeatedly **link** trees until property is restored

Claim:

Insert takes $O(\log n)$ operations in the worst case.


- Initially: n elements and $O(\log n) + 1$ inserted tree in the collection
- One operation links two trees \rightarrow one fewer in the collection
- Can only perform $O(\log n)$ linkings per insert

Binomial *Heaps*

- Structure of a binomial heap:
 - (1) A list of heap-ordered binomial trees
 - (2) At most one tree of each rank
- Insert(X, k): $O(\log n)$ operations (# of trees)
- Deletemin():
 - Scan **roots** of trees in collection to find the one with min key
 - Remove and return min root **r**
 - Insert children of **r** to the collection

Deletion in a Binomial Heap


Deletemin() : Search all tree roots for **minimum key**



Deletion in a Binomial Heap

Deletemin() :


Delete this root and insert its children one at a time



Deletion in a Binomial Heap

Deletemin() :


...after inserting children B_0, B_1, B_2



Deletion in a Binomial Heap

Deletemin() :


...after inserting children B_0, B_1, B_2



Deletion in a Binomial Heap

Deletemin() :


...after linking two B_3 trees



Deletion in a Binomial Heap

Deletemin() :


...after linking two B_3 trees...and two B_4 trees



Deletion in a Binomial Heap

Deletemin() :

No more than one tree per rank, so we're done.



Binomial *Heaps*

- Structure of a binomial heap:
 - (1) A list of heap-ordered binomial trees
 - (2) At most one tree of each rank
- Insert(X, k): $O(\log n)$ operations (# of trees)
- Deletemin():
 - Scan **roots** of trees in collection to find the one with min key
 - Remove and return min root **r**, insert children of **r**

Claim:

Deletemin takes $O(\log n)$ operations in the worst case.

- Up to $O(\log n)$ insertions, and one insertion can take $O(\log n)$ time! What?

Binomial *Heaps*

- Structure of a binomial heap:
 - (1) A list of heap-ordered binomial trees
 - (2) At most one tree of each rank
- Insert(X, k): $O(\log n)$ operations (# of trees)
- Deletemin():
 - Scan **roots** of trees in collection to find the one with min key
 - Remove and return min root **r**, insert children of **r**

Claim:

Deletemin takes $O(\log n)$ operations in the worst case.


- Initially: n elements and $O(\log n) + 1$ trees in the collection
- One operation links two trees \rightarrow one fewer in the collection
- Can only perform $O(\log n)$ links **across all the insertions**.

Binomial *Heaps*

- Structure of a binomial heap:
 - (1) A list of heap-ordered binomial trees
 - (2) At most one tree of each rank
- Insert(X, k): $O(\log n)$ operations (# of trees)
- Deletemin(): $O(\log n)$ operations (max # children)
- DecreaseKey(X, k):
 - (Can jump straight to the right key)
 - Decrease the value
 - Swap up the tree to fix the ordering


Decreasekey in a Binomial Heap

Decreasekey(x,4) :



Decreasekey in a Binomial Heap


Decreasekey(**x**,4) :



(Find the node in $O(1)$ operations by maintaining array of pointers)


Decreasekey in a Binomial Heap

Decreasekey(**x**,4) :



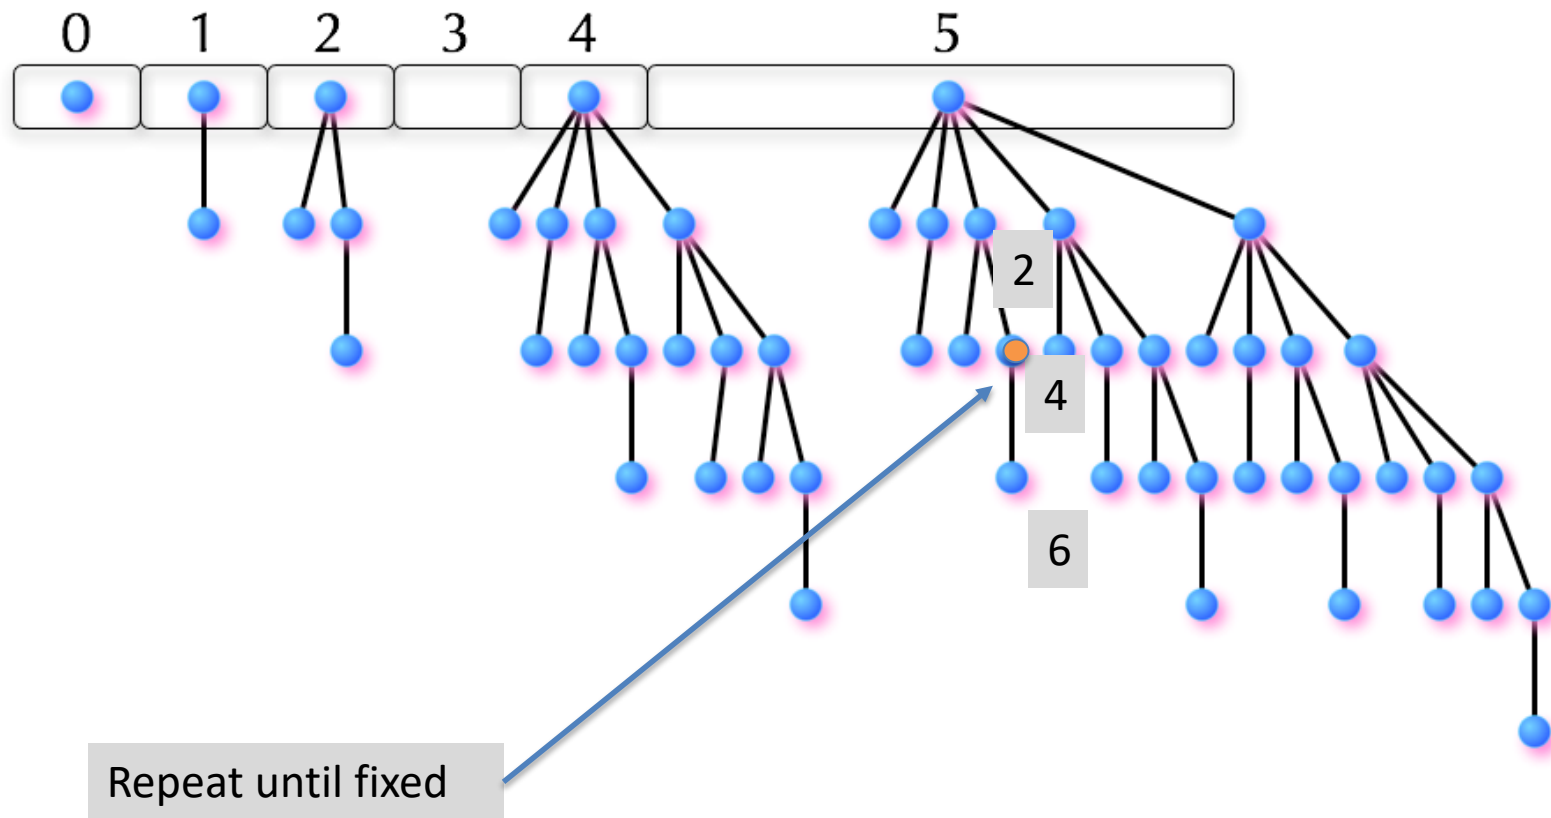
Decreasekey in a Binomial Heap

Decreasekey(**x**,4) :



Decreasekey in a Binomial Heap

Decreasekey(**x**,4) :




Binomial *Heaps*

- Structure of a binomial heap:
 - (1) A list of heap-ordered binomial trees
 - (2) At most one tree of each rank
- Insert(X, k): $O(\log n)$ operations (# of trees)
- Deletemin(): $O(\log n)$ operations (max # children)
- DecreaseKey(X, k): $O(\log n)$ operations (max tree height)

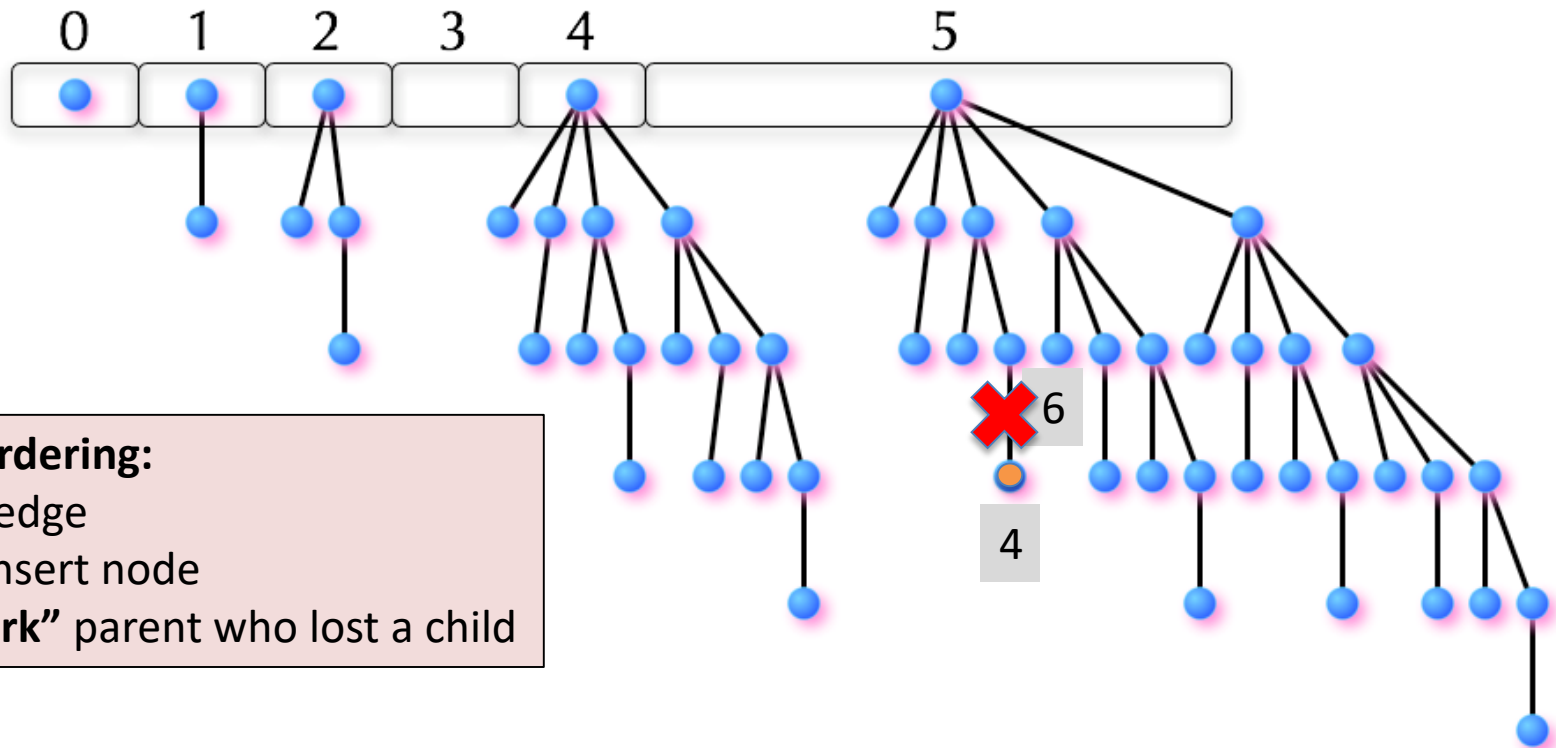
The Fibonacci Heap

- Like binomial heap, but not so picky about the tree structure
- Collection of heap ordered trees (not necessarily binomial)
- At most **one tree per “rank”** (definition will change)
- Insert & Delete same as in binomial heap
- **New DecreaseKey**



The Fibonacci Heap

- Like binomial heap, but not so picky about the tree structure
- Collection of heap ordered trees (not necessarily binomial)
- At most **one tree per “rank”** (definition will change)
- Insert & Delete same as in binomial heap
- **New DecreaseKey**




To fix ordering:


- Cut edge
- Re-insert node
- **“Mark”** parent who lost a child

Slightly different definition of “rank”

- In binomial tree “rank” \equiv “number of children”
- In Fibonacci heap a rank k node EITHER has k children



- OR: It has $k - 1$ children and is **marked** (lost one)




The Fibonacci Heap

- Like binomial heap, but not so picky about the tree structure
 - Collection of heap ordered trees (not necessarily binomial)
 - At most **one tree per “rank”** (definition will change)
 - Insert & Deletemin same as in binomial heap
- Decreasekey(x,k) : **cuts** link from x to parent(x)
 - A node that has lost one child is **marked**
 - A node that has lost two children is **unmarked** and **cuts the link to its parent and reinserts**


Decreasekey in action

- Cut link from node to its parent, mark the parent




Decreasekey in action

- Cut link from node to its parent, mark the parent




Decreasekey in action

- Cut link from node to its parent, mark the parent




Decreasekey in action

- Too many rank 0 nodes, so link them




Decreasekey in action

- Too many rank 0 nodes, so link them




Decreasekey in action

- New decreasekey: cut link from node to its parent, mark the parent



Decreasekey in action

- Parent lost two children
- Cut it and set it to be *unmarked*




To-do list:

- There are two rank-1 trees, fix that
- A node is double-marked, reinsert it


Decreasekey in action

- *Its parent* lost two children
- Cut it as well and set it to be unmarked




Decreasekey in action

- Too many rank 1 nodes, so link them




Decreasekey in action

- Too many rank 1 nodes, so link them



Decreasekey in action

- At most one tree per rank, so we're done.



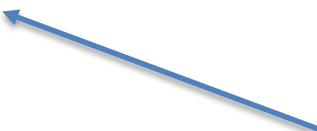
The Challenge

- **Binomial Heaps:**
 - Insert, DeleteMin, DecreaseKey all cost $O(\log n)$ time in the worst case.
- **Fibonacci Heaps:**
 - No changes in Insert, DeleteMin
 - Aiming for $O(1)$ time DecreaseKey
 - DecreaseKey calls Insert
 - **How are we going to get $O(1)$ Insert and DecreaseKey? .**
Are we doomed?

The Challenge

- **Binomial Heaps:**

- Insert, DeleteMin, DecreaseKey all cost $O(\log n)$ time **in the worst case.**



Need “average-case” analysis

- Repeated insertions/cuts can happen
- But very often?

- **Fibonacci Heaps:**


- No changes in Insert, DeleteMin
- Aiming for $O(1)$ time DecreaseKey
- DecreaseKey calls Insert
- **How are we going to get $O(1)$ Insert and DecreaseKey? .**
Are we doomed?

Amortized Analysis

- **Statement:** Total running time is at most
 - $O((\# \text{ insert}) + (\# \text{ decreasekey}) + \log n * (\# \text{ deletemin}))$.
 - *On average*, insert, decreasekey run in time $O(1)$ and deletemin runs in time $O(\log n)$.
- **Strategy:** Maintain a “potential” ϕ in the analysis.
 - Some function of the current state of the heap.
 - Let c_i = actual running time of i -th operation.
 - Let D_i = Fibonacci heap after the i -th operation.
 - Let $\hat{c}_i = c_i + \phi(D_i) - \phi(D_{i-1})$.
 - Total running time $\sum_{i=1}^n c_i = (\sum_{i=1}^n \hat{c}_i) - \phi(D_n) + \phi(D_0)$
 - So, if $\phi(D_0) = 0$ and $\phi(D_n) \geq 0$, $\sum_{i=1}^n c_i \leq \sum_{i=1}^n \hat{c}_i$
 - If $\hat{c}_i \leq O(1)$ whenever i is insert/decreasekey and $\hat{c}_i \leq O(\log n)$ when i is deletemin, we prove what we want!


Our Potential Function

- **Insert:**
 - $O(1)$ work to link two trees, reducing # of trees by 1
 - Might do this repeatedly...
 - But that's okay if **potential depends on the # of trees**




Our Potential Function

- **DecreaseKey:**
 - $O(1)$ work to cut
 - Might have to do this repeatedly...
 - Last cut in the chain:
 - **Increases** the number of trees by 1
 - **Increases** the number of marked nodes by 1



Our Potential Function

- **DecreaseKey:**
 - $O(1)$ work to cut
 - Might have to do this repeatedly...
 - All other cuts in the chain:
 - **Increase** the number of trees by 1
 - **Decrease** the number of marked nodes by 1




Our Potential Function

$$\phi = 2(\# \text{ marked nodes}) + (\# \text{ trees})$$

Amortized Cost

- Insert

- $c_i = 1 + \#links$
- $\Delta\phi = 1 - \#links$
- $\hat{c}_i = 2$



- If you want to more precise bound $c_i \leq a(1 + \#links)$ for some constant a , define
 - $\phi = a \cdot (2(\# \text{ marked nodes}) + (\# \text{ trees}))$
 - Everything will be multiplied by a , but fine.

Amortized Cost


- Deletemin

- $c_i = \#trees + \#children + \#links$

- $\Delta\phi = \#children - 1 - \#links$

- $\hat{c}_i \leq \#trees + 2 \cdot \#children \leq O(\log n)$

Lemma] Tree of rank i has at least i -th Fibonacci number = $\Omega(1.682^i)$.



Amortized Cost

- Decreasekey

- $c_i = 1 + \#cuts + \#links$

- Last cut in the chain


- Increase #trees by 1
 - Increase #markednodes by 1

- Every other cut in chain

- Increase #trees by 1
 - Decrease #markednodes by 1

- $\Delta\phi = 3 - (\#cuts - 1) - \#links$

- $\hat{c}_i = c_i + \Delta\phi = 5.$



Amortized Analysis

- **Statement:** Total running time of Fibonacci Heaps is at most
 - $O((\# \text{ insert}) + (\# \text{ decreasekey}) + \log n * (\# \text{ deletemin}))$.
 - *On average*, insert, decreasekey run in time $O(1)$ and deletemin runs in time $O(\log n)$.

Dijkstra with Fibonacci heaps

- Input: directed graph G with positive lengths, source s
- $\text{dist}[s] = 0$, $\text{dist}[v] = \infty$ for all $v \neq s$, $\text{prev}[s] = \emptyset$. $A = \emptyset$.
- For every $v \in V$
 - **Insert**($Q, v, \text{dist}[v]$)
- While $A \neq V$
 - $v = \text{Deletemin}(Q)$
 - $A \leftarrow A \cup \{v\}$.
 - For every $(v, w) \in E$
 - If $\text{dist}[w] > \text{dist}[v] + \ell(v, w)$
 - $\text{dist}[w] = \text{dist}[w] + \ell(v, w)$, $\text{prev}[w] = v$.
 - **Decreasekey**($Q, w, \text{dist}[w]$).

Insert $O(n)$ times
 $O(n)$ operations **total**

Deletemin $O(n)$ times
 $O(n \log n)$ operations

Decreasekey $O(m)$ times
 $O(m)$ operations **total**

Total: $O(m + n \log n)$ operations

Wrapup

- Upshot: Dijkstra takes $O(m + n \log n)$ time.
- Choosing a good potential function can be tricky!